Who Should Get Money? Estimating Welfare Weights in the U.S.*

Francesco Capozza†
Krishna Srinivasan‡

June 22, 2026

[Most recent version]

Abstract

Evaluating the desirability of a reform typically involves weighing the gains of winners against the losses of losers using welfare weights, which measure the value that society places on a $1 increase in an individual's consumption. These weights can capture various normative ideals such as utilitarianism and equality of opportunity. We develop a portable experimental method to elicit welfare weights and apply it to samples of the U.S. general population. We assess the robustness of the elicited weights to response quality and design features, validate them using measures of support for redistribution, and document their temporal stability. The aggregate welfare weights are progressive: under a constant-elasticity specification, the income elasticity of welfare weights lies between -0.85 and -0.71. These weights are roughly 8–9 times as progressive as the weights implied by the U.S. tax schedule and 4–5 times as progressive as those implied by U.S. tax and transfer policies.

Keywords: Welfare Weights, Policy Views, Income Taxation
JEL Classification: D63, H21, I31

*This work was supported by Grant 2104-31756 from the Russell Sage Foundation, by the University Research Priority Program "URPP Equality of Opportunity" of the University of Zurich, and by WZB. Any opinions expressed are those of the principal investigators alone and should not be construed as representing the opinions of the Russell Sage Foundation. We thank Sandro Ambühl, Aurelien Baillon, Björn Bartling, Katy Bergstrom, Alexander Cappelen, William Dodds, Ernst Fehr, Manuel Grieder, Nathaniel Hendren, Luca Henkel, Michael Kurschilgen, Michel Maréchal, Morten Støstad, Nina Serdarevic, Florian Scheuer, Itai Sher, Stefanie Stantcheva, Dmitry Taubinsky, Bertil Tungodden, and Roberto Weber for their helpful comments. We also thank participants at various conferences and seminars for their helpful comments. This study was pre-registered: https://doi.org/10.1257/rct.16001-1.1

†UB, IEB, CESifo. Email: francesco.capozza@ub.edu
‡UniDistance Suisse. Email: krishna.srinivasan@unidistance.ch.

1 Introduction

Most policy reforms create winners and losers. Evaluating the desirability of such reforms commonly involves weighing the gains of the winners against the losses of the losers using welfare weights. Welfare weights measure the value that society places on providing an additional dollar of consumption to any given individual. While welfare weights are a crucial input for policy evaluation (Saez, 2001; Hendren and Sprung-Keyser, 2020), direct empirical evidence on societal welfare weights is scarce.[1] We develop a portable method to elicit welfare weights using experiments and apply this to samples of the U.S. general population ($N \approx 2000$).

In the experiment, participants in the role of "Social Architects" face pairs of participants in the role of "Recipients." A Social Architect makes several real-stakes decisions that redistribute specified monetary amounts between the pairs of Recipients. Recipients differ in real-world disposable income, with the set of incomes broadly spanning the U.S. income distribution. Social Architects' decisions identify the welfare weights they implicitly assign to the Recipients.

In the canonical utilitarian framework, welfare weights are proportional to Recipients' marginal utility of consumption. More generally, these welfare weights can also capture other normative ideals, such as equality of opportunity, redistribution based on the source of income, and poverty alleviation (Saez and Stantcheva, 2016). For example, a Social Architect guided by equality of opportunity might assign higher welfare weights to Recipients from disadvantaged backgrounds. A key advantage of our approach is that it elicits welfare weights from Social Architects' redistributive choices, without requiring us to specify or uncover the underlying normative ideals. The elicited weights, conditioned on income, can be directly used to evaluate policy reforms conditioned on income.

Pooling responses of all Social Architects, we find that the median welfare weights assigned to Recipients are progressive—i.e., they decrease with income. We summarize the degree of progressivity using a parametric constant-elasticity specification: the median income elasticity of welfare weights is $\nu = -0.71$. This estimate implies that if a Social Architect assigns a welfare weight of 1 to a Recipient, they would assign a weight of 0.61 to a Recipient earning twice as much. In monetary terms, the Social Architect is indifferent between taking $1 from a Recipient and giving $0.61 to a Recipient earning half as much.

[1] Previous studies generally either impose welfare weights (e.g., Saez, 2001) or indirectly infer them from existing policies (e.g., Lockwood and Weinzierl, 2016; Jacobs et al., 2017; Hendren, 2020). Imposed weights, however, are not empirically grounded, so it is unclear whether they reflect societal values. Weights implied by policies, though empirically grounded, may conflate citizens' normative judgments with the political-economy forces that shape policy (Stantcheva, 2016; Lockwood and Weinzierl, 2016).

Since these welfare weights are elicited based on Recipients' disposable incomes resulting from the current tax and transfer system, progressive welfare weights indicate, abstracting from the efficiency costs of taxes, a desire for redistribution beyond that achieved by the current system.

While the aggregate welfare weights are progressive, there is considerable heterogeneity. Most Social Architects assign progressive weights, but the degree of progressivity varies across Social Architects. This heterogeneity is predicted by the Social Architects' background characteristics. Republicans assign less progressive welfare weights relative to Democrats and Independents, consistent with the partisan gap in support for government redistribution documented in the literature (e.g., Kuziemko et al., 2015; Stantcheva, 2021). Social Architects with above-median age assign less progressive welfare weights. By contrast, the income gradient in support for redistribution documented in the literature (e.g., Singhal, 2008; Cohn et al., 2023) is muted in our data: income is not a statistically significant predictor of overall progressivity. However, lower-income Social Architects assign relatively higher weight to Recipients with incomes close to their own.

We conduct several tests to assess the robustness of the elicited welfare weights. To assess robustness to response quality, we flag low-quality responses based on several indicators, such as failing comprehension checks. Responses flagged as low-quality are less progressive; consequently, adjusting for the quality of responses using the procedure of Luttmer and Samwick (2018) leads to slightly more progressive estimates. To assess robustness to features of the experimental design, we vary three features of the decision environment, including the visual placement of the two Recipients on the screen, the order in which Recipient pairs are presented, and the income of the Recipient common across the pairs of Recipients. We also vary two features of our elicitation method—a "staircase method" that presents participants with decisions in an adaptive manner. We vary the first decision to test for anchoring to the first decision, and vary whether participants verify the consequences of their decisions to test for mistakes.

We find that welfare weights are influenced by some of these features—the progressivity estimates vary across treatments. We reconcile the progressivity estimates across treatments using several strategies proposed in the literature. The resulting income-elasticity of welfare weights is more progressive than the pooled benchmark (-0.71) and lies in a relatively narrow range: $\nu \in [-0.85, -0.71]$. These welfare weights can be implemented in optimal policy formulas using the constant-elasticity function c^{ν} with $\nu \in [-0.85, -0.71]$, where c is consumption, or with CRRA utilities with coefficient of relative risk aversion $\gamma = -\nu \in [0.71, 0.85]$. As an application, we calibrate the optimal non-linear labor income tax formula of Saez (2001) with these welfare weights and mid-range estimates of the elasticity of taxable income. The resulting average optimal marginal tax rates are given by 56%–60%, suggesting that the bounds on welfare weights translate into relatively tight bounds on optimal income taxes.

To validate welfare weights, we compare them with two survey items measuring (i) general support for government redistribution and (ii) support for government redistribution at the margin. Social Architects with more progressive welfare weights express stronger support for government redistribution on both measures. When comparing levels, we find that 78% of Social Architects assign progressive welfare weights—implying a preference to redistribute at the margin—closely matching the 74% who support additional progressive government redistribution at the margin. The close alignment between welfare weights and support for government redistribution, both in individual-level correlations and in aggregate shares, supports the validity of our measure of welfare weights.

To gauge temporal stability, we re-elicit Social Architects' welfare weights in a second survey wave conducted four weeks after the first. Welfare weights are temporally stable, both at the individual and aggregate levels. The cross-wave correlation in individual progressivity estimates is comparable to correlations documented in the broader preference-elicitation literature (e.g., Wölbert and Riedl, 2013; Dean and Sautmann, 2021) and to the cross-wave correlations of survey measures of support for government redistribution that we fielded in both waves. At the aggregate level, the median progressivity estimate is very similar across waves. The individual-level correlation in progressivity estimates and the similarity of the aggregate progressivity estimates across waves indicate considerable temporal stability of our measure of welfare weights.

Finally, we compare the welfare weights obtained from our experiment to those implied by the U.S. income tax schedule (obtained from Hendren (2020)) and tax and transfer policies (obtained from Hendren and Sprung-Keyser (2020)). These weights—referred to as "inverse optimum weights"—reflect the welfare weights that can rationalize current policies. Under the constant-elasticity specification, the welfare weights obtained from our experiment are about 8–9 times more progressive than the weights implied by the tax schedule, and about 4–5 times more progressive than those implied by tax and transfer policies. We discuss several possible explanations for this gap in Section 5, though we cannot adjudicate among them.

Our paper is related to a large literature on preferences for redistribution, which studies (i) the normative ideals that guide such preferences (e.g., Drenik and Perez-Truglia, 2018; Almås et al., 2020), (ii) social preferences or distributional preferences using controlled experiments (e.g., Fisman et al., 2007, 2023; Fehr et al., 2024), and (iii) preferences for government redistribution measured in surveys (e.g., Alesina and Angeletos, 2005; Luttmer and Singhal, 2011; Cruces et al., 2013; Durante et al., 2014; Karadja et al., 2017; Kuziemko et al., 2015; Stantcheva, 2021). We contribute to this literature by eliciting welfare weights, which can be directly used to evaluate policies.[2]

Most closely related is a small literature that directly elicits welfare weights. Saez and Stantcheva (2016) elicit welfare weights as a function of disposable income and taxes and use them to calibrate optimal linear income taxes with no behavioral responses. Our paper estimates weights as a function of disposable income alone, which makes them portable across policy domains, and uses them to calibrate optimal non-linear income taxes that incorporate behavioral responses. We also document robustness, validity, and temporal stability of welfare weights, supporting their use in optimal policy formulas.[3] In a recent working paper, Rodemeier and Sun (2025) builds on our methodology and shows that aversion to government as an institution accounts for a substantial share of the gap between welfare weights and people's stated tax preferences.

A related approach recovers "inverse-optimum" welfare weights from the existing tax schedule (e.g., Hendren, 2020; Lockwood and Weinzierl, 2016; Jacobs et al., 2017; Bourguignon and Spadaro, 2012) or from tax and transfer policies (Hendren and Sprung-Keyser, 2020). A key limitation of inverse-optimum weights is that they may not be normatively appealing if they reflect political-economy considerations such as lobbying (Stantcheva, 2016; Lockwood and Weinzierl, 2016). In contrast, our paper directly elicits the welfare weights of citizens using experiments

Finally, a strand of literature incorporates a parsimonious set of normative ideals into optimal policy formulas by modifying individuals' utilities or the objective function (e.g., Weinzierl, 2014; Fleurbaey and Maniquet, 2006). Our approach calibrates "standard" optimal policy formulas using welfare weights elicited from the general population, which can capture a variety of underlying ideals.

The paper is organized as follows. Section 2 outlines the theoretical framework. Section 3 describes the experimental design. Section 4 reports the experimental results. Section 5 compares the elicited welfare weights to those implied by tax and transfer policies. Section 6 presents a discussion.

[2]Some studies, such as Fisman et al. (2007) and Fisman et al. (2023), elicit participants' choices over equity-efficiency tradeoffs—similar to the decisions in our study—between themselves and a random stranger. However, the income of the stranger is left unspecified, so welfare weights assigned to specific income levels cannot be recovered.

[3]There are also two methodological differences. First, Saez and Stantcheva (2016) rely on a non-representative sample recruited from Amazon Mechanical Turks, while our paper recruits samples broadly representative on age, gender, income, education, and region. Second, their elicitation uses hypothetical decisions, whereas ours involves real-stakes decisions. Our paper shows the importance of incentivizing choices.

2 Theoretical Framework

We present a simple theoretical framework to conceptualize welfare weights and to formalize their identification from redistributive choices.

2.1 Setup

Consider a population of N *Recipients* indexed by j. Each Recipient earns income z_j and pays taxes $T(z_j)$ as a function of income. Their consumption (disposable income) c_j is given by their after-tax income $c_j = z_j - T(z_j)$. We assume that earnings are completely inelastic to taxes and transfers: a reform that changes Recipients' tax liabilities generates only mechanical changes in consumption. This assumption allows us to focus on redistributive rather than efficiency concerns.

2.2 Welfare Weights

A *Social Architect* evaluates policy reforms that affect the Recipient based on their welfare weights. The welfare weight $g_j = g(c_j, \theta_j)$ measures the Social Architect's assessment of the value of increasing the consumption of Recipient j by \$1.[4] These welfare weights are a function of the Recipients' consumption c_j and characteristics contained in θ_j (e.g., disability status or parental income).[5] Welfare weights are defined up to a multiplicative constant, as they measure the relative value of consumption of Recipient j.

Welfare weights are a reduced-form representation of underlying normative ideals. In the canonical utilitarian approach, welfare weights are proportional to Recipients' marginal utility of consumption (a function of c_j). However, as shown by Saez and Stantcheva (2016), they can also capture other normative ideals, like equality of opportunity, poverty alleviation, or redistribution based on the source of income. For example, a Social Architect guided by equality of opportunity would assign higher welfare weights to Recipients from disadvantaged backgrounds (captured by θ_j).

The reforms we study are conditioned only on the Recipients' consumption. Thus, even if welfare weights vary with other characteristics contained in θ_j (e.g., parental income) at a given consumption level c, only their average welfare weight at each consumption level is relevant for evaluating these reforms. Denoting $h(c)$ as the number of Recipients with

[4]Formally, a Social Architect maximizes an objective function (not necessarily welfarist), and their welfare weight on Recipient j measures the marginal increase in their objective function from a \$1 increase in Recipient j's consumption.

[5]As shown by Saez and Stantcheva (2016), some characteristics in θ_j, such as disability status, may enter Recipients' utility functions, giving rise to welfare weights that reflect standard utilitarian considerations; however, other characteristics, such as parental income, may not enter Recipients' utility functions but still affect a Social Architect's welfare weights.

consumption c, this average welfare weight is given by:

$$\bar{g}(c) = \frac{\sum_{j:c_j=c} g_j}{h(c)}. \tag{1}$$

A schedule of welfare weights across consumption levels can be consistent with multiple underlying ideals. For example, welfare weights decreasing with consumption can arise from utilitarianism or equality of opportunity. However, only the average welfare weights $\bar{g}(c)$ are needed to evaluate reforms conditioned on income.

2.3 Identifying Welfare Weights

To illustrate the identification, consider a setting with two Recipients with consumption levels c_ℓ and c_h, where $c_h > c_\ell$. Consider a "small" (not necessarily budget-neutral) reform $(r_\ell, -r_h)$, which increases the consumption of the low-consumption Recipient by r_ℓ and decreases the consumption of the high-consumption Recipient by r_h. That is, the transfer amounts are small enough that the Social Architect's valuation of a marginal dollar at each consumption level is approximately unchanged by the reform.

If a Social Architect assigns average welfare weights $\bar{g}(c_\ell)$ and $\bar{g}(c_h)$ to the low-consumption and high-consumption Recipients, respectively, they value the reform at $\bar{g}(c_\ell)\cdot r_\ell - \bar{g}(c_h)\cdot r_h$—the value they place on the gain to the low-consumption Recipient net of the value they place on the loss to the high-consumption Recipient. At the reform that leaves the Social Architect indifferent between implementing the reform and maintaining the status quo (0, 0), the gain just offsets the loss, $\bar{g}(c_\ell)\cdot r_\ell = \bar{g}(c_h)\cdot r_h$, implying

$$\tilde{g}(c_l, c_h) = \frac{\bar{g}(c_h)}{\bar{g}(c_\ell)} = \frac{r_\ell}{r_h}. \tag{2}$$

Equation (2) shows that the ratio of welfare weights at different consumption levels is identified by the ratio of reform amounts that makes the Social Architect indifferent between implementing the reform and maintaining the status quo. If $\tilde{g} < 1$, welfare weights decrease with Recipients' consumption ($\bar{g}(c_\ell) > \bar{g}(c_h)$), which we define as *progressive* welfare weights. Analogously, $\tilde{g} > 1\,(\bar{g}(c_\ell) < \bar{g}(c_h))$ is defined as *regressive* welfare weights and $\tilde{g} = 1\,(\bar{g}(c_h) = \bar{g}(c_\ell))$ is defined as equal weights.

Thus, under the maintained assumptions, choices over redistributive reforms identify relative welfare weights without requiring the Social Architect to articulate—or the researcher to specify—the underlying normative ideals. The next section describes the experimental design that implements this elicitation.

3 Experimental Design and Sample

This section describes the experimental design and sample used to elicit the welfare weights of the U.S. general population and to test their robustness, validity, and temporal stability. The experimental design, sample restriction, and analyses were pre-registered.[6] Deviations from the pre-registration are discussed in Appendix Section B. The complete set of instructions can be found in Appendix Section G.

3.1 Eliciting Welfare Weights

3.1.1 Approach

Participants in our experiment are assigned to one of two roles: *Social Architect* or *Recipient*. Social Architects decide whether to implement various "reforms" that redistribute specified monetary amounts between pairs of Recipients. Recipients differ in their real-world disposable incomes, which we use as the empirical counterpart of consumption c_j in the theoretical framework. Under the framework in Section 2, these decisions identify the welfare weights they implicitly assign to Recipients. Social Architects' decisions may have real consequences: one randomly selected decision made by one randomly selected Social Architect is implemented. Recipients are passive participants who receive payments based on the Social Architects' choices.

Our approach has five key features. First, to capture Social Architects' assessments of the value of transfers given the current tax and transfer system, we present them with Recipients' disposable incomes under the current system.[7] Social Architects also learn that the Recipients are U.S. citizens, at least 18 years old, and randomly drawn from the U.S. general population. Second, the reforms constitute "small" one-time transfers; with large transfers, a Social Architect's valuation of the first dollar transferred may differ from their valuation of the last dollar, preventing clean elicitation of welfare weights at Recipients' pre-reform income levels. To evaluate larger reforms, we use a parametric function for the welfare weight (described in Section 4.2). Third, to help ensure that welfare weights are unconfounded by views about real-world policies or government, such as beliefs about the behavioral responses to taxation, the reforms in the experiment make no references to the government or real-world policies. Fourth, the reforms are presented in a gain-loss frame—one Recipient loses while the other gains—so that the elicited welfare weights can be applied to evaluate real-world policies that typically also involve winners and losers (Saez, 2001;

[6] See https://doi.org/10.1257/rct.16001-1.1.

[7] Eliciting welfare weights based on pre-tax incomes is challenging because it requires Social Architects to ignore Recipients' disposable incomes, despite their incentivized decisions affecting these incomes.

Hendren and Sprung-Keyser, 2020).[8] Finally, following the tradition in welfare economics (e.g., Cappelen et al., 2013), Social Architects are not directly affected by their own decisions, which helps minimize the role of self-interest motives.

3.1.2 Decisions

To elicit the welfare weights assigned by a Social Architect to a pair of Recipients with different disposable incomes, we present the Social Architect with several binary choices between implementing a reform and maintaining the status quo. Implementing a reform $(r_\ell, -r_h)$ would increase the disposable income of the low-income Recipient by r_ℓ and decrease the disposable income of the high-income Recipient by r_h, while maintaining the status quo leaves both incomes unchanged. To allow for negative transfers, we endow both Recipients with an initial endowment of \$1,000, which is incorporated in their disposable incomes, and explicitly disclose this to the Social Architects. The endowment is less salient by design: Social Architects only see disposable incomes that includes the endowment when making a decision (e.g., a Recipient with \$59,000 in real-world income appears as \$60,000). By Equation (2), the reform that makes a Social Architect indifferent between implementing it and maintaining the status quo identifies their welfare weights $\tilde{g}(c_\ell, c_h) = \bar{g}(c_h)/\bar{g}(c_\ell) = r_\ell/r_h$.

Figure 1 presents a screenshot of one of the decisions presented to the Social Architects. At the top of the screen, we remind Social Architects that their choices may have real consequences (see Section 3.1.4 for details on incentives). We randomize the placement of the Recipients on the screen (left vs. right) across Social Architects. We refer to a reform as a "proposed change" rather than "reform" or "policy" to avoid references to real-world policies or government. At the bottom of the screen, we present the final incomes of Recipients implied by the two options.

Table 1 presents the full set of reforms $(r_\ell, -r_h)$ used to elicit welfare weights and the implied $\tilde{g}$ for a Social Architect indifferent at each reform. To cover a wide range of $\tilde{g}$ without requiring extremely large transfers, we simultaneously vary r_ℓ and r_h across the reforms.[9] The reform amounts range from \$50 to \$950 in increments of \$50, with the average amount (\$500) comparable to several tax and transfer policies in the U.S. (see Appendix Table A1

[8] A gain-gain frame of the form $(0, 0)$ vs. (r_ℓ, r_h) with $r_\ell, r_h > 0$ cannot identify welfare weights, since the reform would always be preferred. A gain-gain frame that involves giving r_ℓ to the low-income Recipient or giving r_h to the high-income Recipient can, in principle, identify welfare weights; we avoided it because participants may be reluctant to favor the high-income Recipient when the low-income Recipient receives nothing, which would bias the elicited weights toward progressivity. We conducted an additional study (Study 2) with 1,965 participants to compare the two framings (Details in Appendix Section F). The evidence is mixed: the gain-gain frame yields more progressive weights when the Recipient common across the pairs is lower-income—a difference in income elasticity of 0.12, comparable to the 0.14-wide bounds reported in Section 4.6—and no difference when that Recipient is higher-income.

[9] For instance, to obtain $\tilde{g} = 19$ (the largest $\tilde{g}$ in the table) with a fixed $r_h =$ \$500 requires $r_\ell =$ \$9,500.

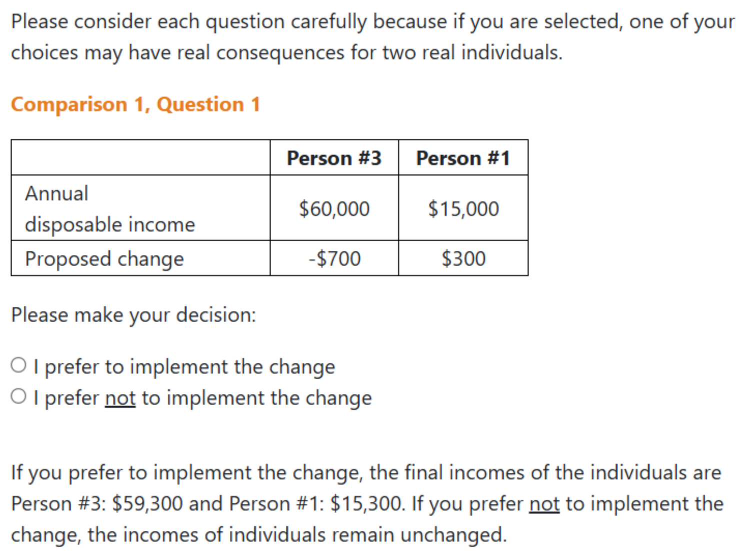
Please consider each question carefully because if you are selected, one of your choices may have real consequences for two real individuals.

Comparison 1, Question 1

	Person #3	Person #1
Annual disposable income	\$60,000	\$15,000
Proposed change	-\$700	\$300

Please make your decision:

I prefer to implement the change
I prefer not to implement the change

If you prefer to implement the change, the final incomes of the individuals are Person #3: \$59,300 and Person #1: \$15,300. If you prefer not to implement the change, the incomes of individuals remain unchanged.

Figure 1: Screenshot of a Decision Presented to Social Architects

for examples). The budget-neutral reform in Row 10 corresponds to $\tilde{g} = 1$, reforms above Row 10 correspond to $\tilde{g} < 1$, and reforms below Row 10 correspond to $\tilde{g} > 1$.[10] The reforms are symmetric around Row 10: $\tilde{g}$ in Row $10 - n$ is equal to $\frac{1}{\tilde{g}}$ in Row $10 + n$, $\forall n \in \{1, \ldots, 9\}$.

We aim to identify a Social Architect's switch-point—the row where they switch from preferring the status quo (where no reform is implemented) to preferring the reform listed in that row. We approximate the indifference point as the mid-point of the reforms in the switch-point row and the previous row.[11] To identify a Social Architect's switch-point, we use a "staircase method," which presents them with three to five decisions that are adaptively selected based on their previous choices. This method allows us to obtain the point of indifference with few decisions, reducing the cognitive burden of the task. [12] Because this

[10]The reforms are generally not budget-balanced, so the experimenter keeps any surplus and funds any deficit. A Social Architect who places positive welfare weight g_E on the experimenter has an elicited ratio $(\bar{g}(c_h) - g_E)/(\bar{g}(c_\ell) - g_E)$, which amplifies departures from the true ratio of 1; since a majority of Social Architects have progressive weights, the elicited weights would overstate progressivity. We did not make the source of the unallocated funds salient, to keep the focus on the Recipients. Rodemeier and Sun (2025) test for precisely this concern: informing subjects that any funds they leave unallocated simply disappear—and thus have no value to the experimenter—leaves choices unchanged relative to a control condition in which nothing was specified; furthermore, when that study makes a productive alternative use of unspent funds salient, behavior does shift. Together, these results indicate that control subjects are likely inattentive to the experimenter's budget.

[11]If a Social Architect always prefers a reform (i.e., switches in the first row), we take the mid-point of (\$50, $-$\$950) and (\$0, $-$\$1000), with the latter (corresponding to $\tilde{g} = 0$) giving the minimum possible amount to the low-income Recipient. If a Social Architect always prefers the status quo (i.e., never switches), we take the mid-point of (\$950, $-$\$50) and (\$1000, $-$\$0), with the latter reform (corresponding to $\tilde{g} = \infty$) taking the minimum possible amount from the high-income Recipient.

[12]Compared to a Multiple Price List (MPL), which presents all decisions at once, the staircase method is simpler for the general population to understand; Falk et al. (2018) uses the staircase method with nationally

Table 1: Set of Reforms Used to Elicit Welfare Weights

Row	Reform $(r_\ell, -r_h)$	$\tilde{g}$
1	(\$50, −\$950)	0.05
⋮	⋮	⋮
6	(\$300, −\$700)	0.43
⋮	⋮	⋮
8	(\$400, −\$600)	0.67
9	(\$450, −\$550)	0.82
10	(\$500, −\$500)	1
11	(\$550, −\$450)	1.22
12	(\$600, −\$400)	1.5
⋮	⋮	⋮
14	(\$700, −\$300)	2.33
⋮	⋮	⋮
18	(\$900, −\$100)	9
19	(\$950, −\$50)	19

Notes: The table presents the set of reforms used to elicit welfare weights. A reform $(r_\ell, -r_h)$ increases the income of the low-income Recipient by r_ℓ and decreases the income of the high-income Recipient by r_h. $\tilde{g} = \frac{r_\ell}{r_h}$ is the ratio of the reform amounts.

adaptive elicitation can be sensitive to the starting point, we randomize the first decision in the staircase, following the approach of Luttmer and Samwick (2018) and Bursztyn et al. (2025). For half the Social Architects, the reform in the first decision is (\$300, −\$700), while for the other half, it is (\$700, −\$300). The first decision is identical across all comparisons. Appendix Figures A2 and A3 present a graphical representation of the staircase method with the two different starting points. We discuss this and other randomizations in more detail in the next section.

3.1.3 Recipients' Incomes and Comparisons

A Social Architect makes redistributive decisions for five different pairs of Recipients, involving six Recipients in total (see Table 2); a sixth comparison, identical to the third, serves as a consistency check. A common Recipient appears in every comparison, allowing us to combine the five pairwise welfare-weight ratios into a single schedule of relative welfare weights

representative samples worldwide. Furthermore, it avoids the bias associated with the order of presentation of the decisions on the screen observed in the MPL method (e.g., Jack et al., 2022). Evidence suggests that the staircase method can outperform the MPL method in out-of-sample prediction (Toubia et al., 2013). Potential limitations of the staircase method are addressed in Section 3.2.2.

across all six income levels. To check robustness, we randomize the income of the common Recipient to be either $60,000 (see Panel A) or $120,000 (see Panel B). We also randomize the order of the five comparisons: half the Social Architects see them in the order shown in the table, and the other half in reverse order.

The incomes of the six Recipients are chosen to span a broad range of the U.S. disposable income distribution while satisfying three design constraints: (i) the highest income is capped at $500,000 to keep the recruitment of high-earners feasible, (ii) the ratios of neighboring incomes are constant to minimize cognitive burden, and (iii) incomes are rounded to reduce errors due to numeracy. The resulting incomes range from $15,000 to $480,000, with each subsequent income double the previous one (see Appendix Figure A4 for a graphical representation).

Table 2: Recipients' Incomes Across Comparisons

Panel A: Common Recipient earns $60,000					
	Comparison				
Recipient	1	2	3	4	5
Low-Income	$15,000	$30,000	$60,000	$60,000	$60,000
High-Income	$60,000	$60,000	$120,000	$240,000	$480,000
Panel B: Common Recipient earns $120,000					
	Comparison				
Recipient	1	2	3	4	5
Low-Income	$15,000	$30,000	$60,000	$120,000	$120,000
High-Income	$120,000	$120,000	$120,000	$240,000	$480,000

3.1.4 Incentives

To incentivize thoughtful decisions, Social Architects are informed that one participant in the study would be randomly selected at the end of the study, and one of their decisions would be randomly selected and implemented.[13] The decision could be selected from either of the two waves of data collection (data collection is discussed in the next section). Each decision saliently reminds Social Architects of the incentives (see Figure 1).

3.2 Assessing Robustness, Validity, and Temporal Stability

To assess the robustness, validity, and temporal stability of the elicited welfare weights, we conduct various tests.

[13]An additional study (Study 3) conducted on Prolific ($N = 1992$) compares welfare weights elicited under real stakes (as in this study) with those elicited under hypothetical stakes. Welfare weights are more progressive under hypothetical stakes (with a difference in income elasticity of 0.30), consistent with real-stakes decisions inducing more careful considerations of tradeoffs. Design details, other differences between the two studies, and full results are reported in Appendix Section F.

3.2.1 Assessing Robustness to Quality of Responses

We assess the robustness of welfare weights to individual-level response quality using several proxies. First, we flag Social Architects failing one or more comprehension questions on the first attempt as having low-quality responses. We included three comprehension questions that test whether Social Architects have understood (i) that they will make decisions regarding six Recipients, (ii) that they will be presented with disposable incomes (rather than pre-tax incomes), and (iii) that their decisions may have real stakes. Second, we elicit Social Architects' confidence in their decisions on a five-point scale ("Not at all"; "Very little"; "Little"; "Somewhat"; "Very much") and classify those not choosing the highest category as having low response quality. Third, we check decision consistency by including a sixth comparison identical to the third comparison (Recipients earning \$60,000 and \$120,000) and flag those whose welfare weights are inconsistent in sign (progressive in one and regressive in the other) across comparisons as having low-quality responses. Finally, we flag completion times beyond two standard deviations from the mean within each treatment as an indicator of low response quality. Similar proxies have been used in the literature to assess response quality (e.g., Luttmer and Samwick, 2018; Enke et al., 2023; Stantcheva, 2023).

3.2.2 Assessing Robustness to Design Features

We assess whether the estimated welfare weights are robust to features of the experimental design by introducing several treatments that vary these features across Social Architects.

The first three treatment dimensions vary features of the decision environment, including (i) the placement of the Recipients on the screen (left vs. right), (ii) the order of presentation of the comparisons, and (iii) the Recipient common across the comparisons (\$60,000 vs. \$120,000). The set of high-to-low income ratios of Recipients is identical across treatments that vary the common Recipient, which allows us to isolate the role of the common Recipient.

The remaining two treatments vary the features of the staircase method to address two potential limitations: anchoring to the first decision and mistakes that may propagate through the adaptive elicitation. First, to test and correct for anchoring, we randomize the reform in the first decision to be either (\$300, −\$700) or (\$700, −\$300), following Luttmer and Samwick (2018) and Bursztyn et al. (2025). Second, to test for mistakes, we vary whether Social Architects are prompted to verify the consequences of their choice before proceeding. In the verification treatment, Social Architects are shown the income distributions implied by each option—implementing the reform or maintaining the status quo—and asked which income distribution they prefer; they can proceed only if their preferred income distribution is consistent with their original choice (see Appendix Figure A1). For example, if a Social Architect chose to implement the reform (\$300, −\$700) when comparing Recipients earning

$15,000 and $60,000, they are shown the implied income distributions—$15,300 and $59,300 (reform) versus $15,000 and $60,000 (status quo); they can only proceed if they indicate that they prefer $15,300 and $59,300. To minimize burden, this verification appears only in the first decision of each comparison, where mistakes generate the largest distortions in elicited welfare weights.[14] Social Architects are informed in advance that the verification will take place. Our approach is similar to that of Bursztyn et al. (2025), who present participants with their point of indifference implied by their decisions and let them revise; related approaches prompt participants to reconsider their choices (e.g., Burchardi et al., 2021; Berry et al., 2020; Abdellaoui et al., 2019).

3.2.3 Correlation with Support for Redistribution

We validate our measure of welfare weights by comparing them with measures of support for government redistribution. We measure general support for government redistribution with a question used in the General Social Survey (GSS) and by Alesina et al. (2018) and support for government redistribution at the margin (i.e., beyond the current tax and transfer system). The order of the questions is randomized across Social Architects.

3.2.4 Do Welfare Weights Capture Fairness Concerns?

To test whether welfare weights capture fairness concerns (as intended) or other concerns (e.g., beliefs about labor distortions from taxes), we include survey questions capturing various concerns, as well as misperceptions about the tax system and the income distribution. Detailed descriptions of the measures are presented in Appendix Section A.

3.2.5 Temporal Stability of Welfare Weights

To gauge the temporal stability of welfare weights, we re-elicited Social Architects' welfare weights in a second survey wave conducted four weeks after the first. Participants with complete responses in Wave 1 were invited to participate in Wave 2. The welfare weights elicitation protocol (including treatment assignments) was identical across waves. The four-week interval limits the scope for external shocks or changes in personal circumstances to alter preferences, while still allowing us to test whether the elicited weights are stable. We benchmark the stability of welfare weights against the stability of our two measures of support for government redistribution, which were included in both waves.

[14] Early errors in the staircase distort the elicited weights more than later ones, because the first decision fixes the coarsest split and later decisions only refine it. To formalize this, suppose a Social Architect errs in decision k but answers all other decisions correctly, so the elicited welfare-weight ratio is $x_k(\tilde{g}) \cdot \tilde{g}$ rather than the true ratio $\tilde{g}$. Across the true ratios $\tilde{g}$ on our grid, the variance of the log distortion $\ln x_k(\tilde{g})$ is 3.39, 1.79, 0.85, 0.50, and 0.31 for mistakes in the first through fifth decisions, respectively; this implies a typical multiplicative distortion of roughly 6×, 3.8×, 2.5×, 2×, and 1.7×.

3.3 Data Collection

We conducted the study across two waves. In Wave 1 of data collection, we recruited participants in the role of Social Architects from the data collection provider Prolific.[15] Recruitment was based on three demographic quotas available on Prolific: gender, political affiliation, and age. We included one attention check and dropped participants failing the check. We implemented the experiment using oTree (Chen et al., 2016). The data collection for Wave 1 began on 12 May 2025 and lasted three days. The median completion time was 15.7 minutes. Our final sample includes 1,996 participants.[16]

We conducted the second wave of data collection roughly four weeks after the completion of the first wave. We invited 1,986 participants to Wave 2, including those who completed the survey in the first wave and had a valid Prolific ID prior to data collection. The data collection for Wave 2 began on 9 June 2025 and lasted three days. The median completion time was 9 minutes. Of those invited, 1,447 (73%) returned for Wave 2, and after our pre-registered sample restrictions, our final sample includes 1,397 participants.[17]

3.4 Summary Statistics

Table 3 presents the sample averages of Social Architects' background characteristics and the population averages. The Wave 1 sample broadly resembles the population but under-represents very low-income earners (income <$30,000) and individuals with a high-school education or less, while over-representing middle-income earners ($30,000–$150,000) and the college-educated. Our main progressivity estimates are estimated using sampling weights that match the sample averages of background characteristics (age, gender, individual income, education, and region of residence) with the population averages specified in Table 3. The weights are defined for 1,977 Wave 1 participants with complete responses on all five demographic variables; all weighted estimates are computed on this sample. The table shows that the Wave 1 and Wave 2 samples are also very similar.[18]

We also test representativeness using two survey items fielded in our experiment with wording identical to that used in large national polls. First, we find that 17% of the weighted Wave 1 sample report that the government can be trusted "most of the time" or "just about

[15] Prolific has been used in several recent studies (e.g., Bursztyn et al., 2025; Enke et al., 2023).

[16] We recruited 2,194 participants in Wave 1. Following our pre-registration, we drop those with multiple survey responses (0.18%), who failed the attention check (0.42%), and who did not complete the study.

[17] Following our pre-registered sample selection, we drop participants with multiple survey responses (0.35%), who failed the attention check (0.14%), and who did not complete the survey.

[18] Tables A2 report balance tests across treatment arms in Wave 1. For each background characteristic in Table 3, we regress the characteristic on indicators for the five treatment dimensions. F-tests of the null that the treatment indicators are jointly zero do not reject at the 5% level for any of the 21 characteristics, indicating no evidence of imbalance across treatments.

Table 3: Summary Statistics of Sample and Population

Variable	Wave 1	Wave 2	Population
Male	0.49	0.49	0.49
Individual income < 30,000	0.26	0.28	0.41
Individual income 30–59,999	0.29	0.28	0.26
Individual income 60–99,999	0.26	0.26	0.18
Individual income 100–149,999	0.12	0.12	0.08
Individual income ≥ 150,000	0.07	0.07	0.07
Education: High school or less	0.11	0.11	0.37
Education: Some college	0.18	0.19	0.20
Education: Bachelor or Associate	0.45	0.47	0.30
Education: Masters or above	0.26	0.24	0.13
Age: 18–24	0.12	0.09	0.11
Age: 25–34	0.18	0.16	0.17
Age: 35–44	0.17	0.17	0.17
Age: 45–54	0.15	0.16	0.16
Age: 55–64	0.25	0.27	0.16
Age: ≥ 65	0.13	0.14	0.23
Region: Northeast	0.16	0.17	0.17
Region: Midwest	0.18	0.18	0.20
Region: South	0.44	0.43	0.39
Region: West	0.22	0.22	0.24
Republican	0.32	0.30	0.28

Notes: The table presents the average background characteristics of our sample and the U.S. population. Columns (2) and (3) report the sample demographics based on the 1996 participants in Wave 1 and 1397 participants in Wave 2, respectively. Column (4) reports U.S. population demographics, computed using the American Community Survey (ACS) 1-year estimates from 2024 (most recent year available) for individuals aged 18 and older (Ruggles et al., 2025). The population share of Republicans is the average share of individuals identifying as Republican, based on the Gallup poll in 2024.

always," closely matching the share (22%) in the Pew Research Center's 2024 national poll. Second, 65% of the weighted Wave 1 sample supports reducing income differences between the rich and the poor, which is slightly higher than the corresponding share (53%) in the 2024 General Social Survey (GSS).[19] These parallels to large-scale national surveys provide additional reassurance regarding the representativeness of our sample.

[19]Responses range on a scale from 1 to 7, which we reverse-code such that higher values indicate stronger preferences for reducing income differences between the rich and the poor. We code values above 4 as supporting reductions in income differences. Unlike the GSS question, our question does not include a "Don't know" option.

4 Welfare Weights of the General Population

4.1 Data Description

Figure 2 presents the cumulative distribution function (CDF) of $\tilde{g}(c_\ell, c_h) \equiv \bar{g}(c_h)/\bar{g}(c_\ell)$—the ratio of welfare weights assigned to the high-income and low-income Recipients—across the five pairs of Recipients, using data from Wave 1. For readability, we refer to this ratio simply as $\tilde{g}$. We infer $\tilde{g}$ from each Social Architect's switching behavior using Equation (2). Figure 2a corresponds to the treatment in which the common Recipient earns $60,000, and Figure 2b to the treatment in which they earn $120,000. Appendix Figures A6a and A6b present the same distributions as heat maps. In each figure, Panel A includes comparisons in which the common Recipient is the higher-income earner, and Panel B includes comparisons in which the common Recipient is the lower-income earner. The dashed horizontal line at 50 marks the 50^{th} percentile; the median $\tilde{g}$ for each comparison can be read as the value at which the corresponding CDF crosses this line.

There are several patterns in the data. First, a majority of Social Architects assign higher welfare weights to lower-income Recipients in each comparison: the distribution of $\tilde{g}$ is concentrated at low values in every comparison, and the median $\tilde{g}$ lies below 1 in all cases. Second, Social Architects assign more progressive welfare weights to pairs of Recipients whose incomes differ by a larger ratio: in every panel, the CDF for the comparison with the smaller income ratio first-order stochastically dominates those with larger ratios. Third, there is some bunching in the distributions: 14–34% of Social Architects bunch at the progressive extreme ($\tilde{g} = 0.03$), while 8–15% bunch at the regressive extreme ($\tilde{g} = 39$); relatively few bunch at values close to equal weights ($\tilde{g} = 0.90$ or $\tilde{g} = 1.11$). We discuss bunching in more detail in Section 4.4. Fourth, these aggregate progressive patterns mask substantial heterogeneity: within each comparison, $\tilde{g}$ ranges from 0.03 (most progressive) to 39 (most regressive). Finally, the $60,000 vs. $120,000 comparison, which is common across treatments, yields very similar distributions of $\tilde{g}$ in both treatments (see Panel B of Figure 2a and Panel A of Figure 2b). This overlap is easier to see in the heat maps (Appendix Figures A6a and A6b).

4.2 Progressivity of Welfare Weights

4.2.1 Non-Parametric Approach

Because the common Recipient appears in each of the five pairwise comparisons, for each Social Architect, we can combine the elicited ratios $\tilde{g}$ into a non-parametric schedule of welfare weights over the six income levels, normalizing the weight on the common Recipient to 1. Figure 3 plots the median schedule (solid line). We focus on the median welfare weights, since they correspond to the aggregate preference under the median voter theorem (Downs,

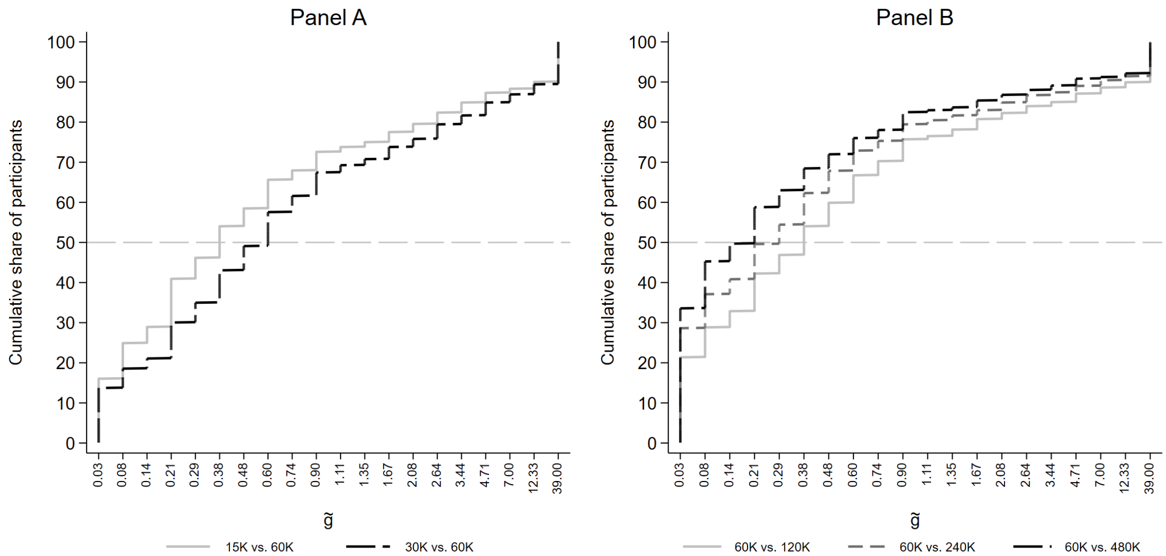

(a) Common Recipient earns $60,000

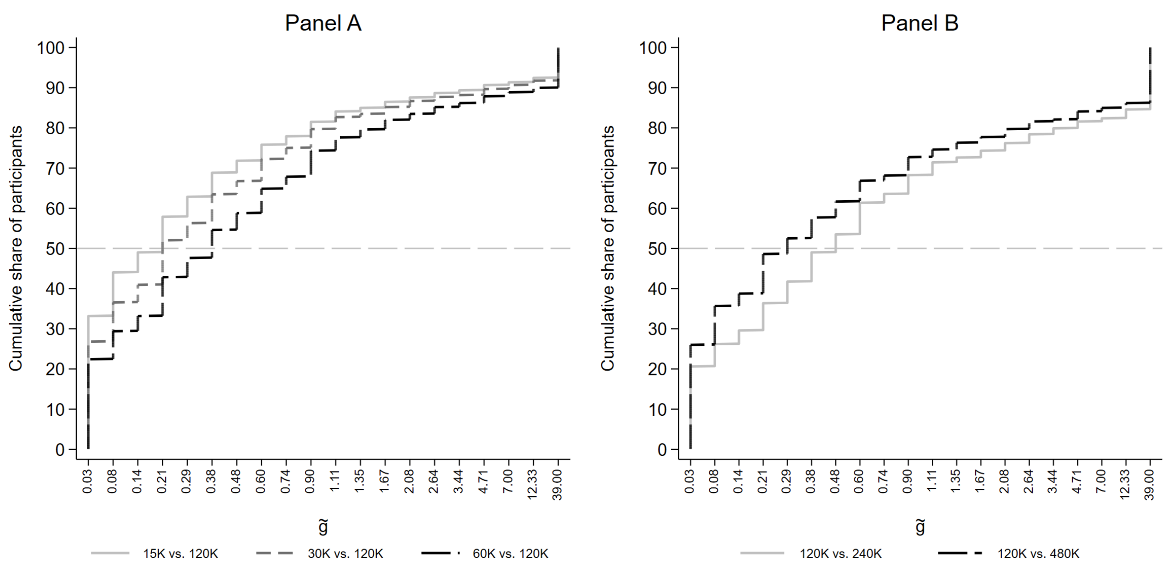

(b) Common Recipient earns $120,000

Figure 2: Empirical CDFs of Welfare Weight Ratios $\tilde{g}$

Notes: This figure presents the cumulative distribution function (CDF) of $\tilde{g}$—the ratio of welfare weights assigned to the high-income and low-income Recipients. The CDFs use Wave 1 data and exclude the repeated sixth comparison used for the consistency check. Figure (a) includes treatments in which the common Recipient earns $60,000, and Figure (b) includes treatments in which the common Recipient earns $120,000. Panel A groups comparisons in which the common Recipient is the higher-income earner; Panel B groups comparisons in which the common Recipient is the lower-income earner. The dashed horizontal line at 50 marks the 50$^{\text{th}}$ percentile. Values of $\tilde{g} < 1$ indicate progressive welfare weights; values of $\tilde{g} > 1$ indicate regressive weights.

1957). In Panel (a), the common Recipient earns \$60,000, and in Panel (b), the common Recipient earns \$120,000.

In both panels, the median welfare weights schedule is weakly decreasing in income, with steeper declines at lower incomes. In Panel (a), the weight at \$15,000 is approximately 2.6, declining to 1.00 at \$60,000 (by normalization) and falling further to 0.21 at \$480,000. In Panel (b), the weights at \$15,000 and \$30,000 are nearly indistinguishable—both approximately 4.7—before falling to 1 at \$120,000 (by normalization) and 0.29 at \$480,000.

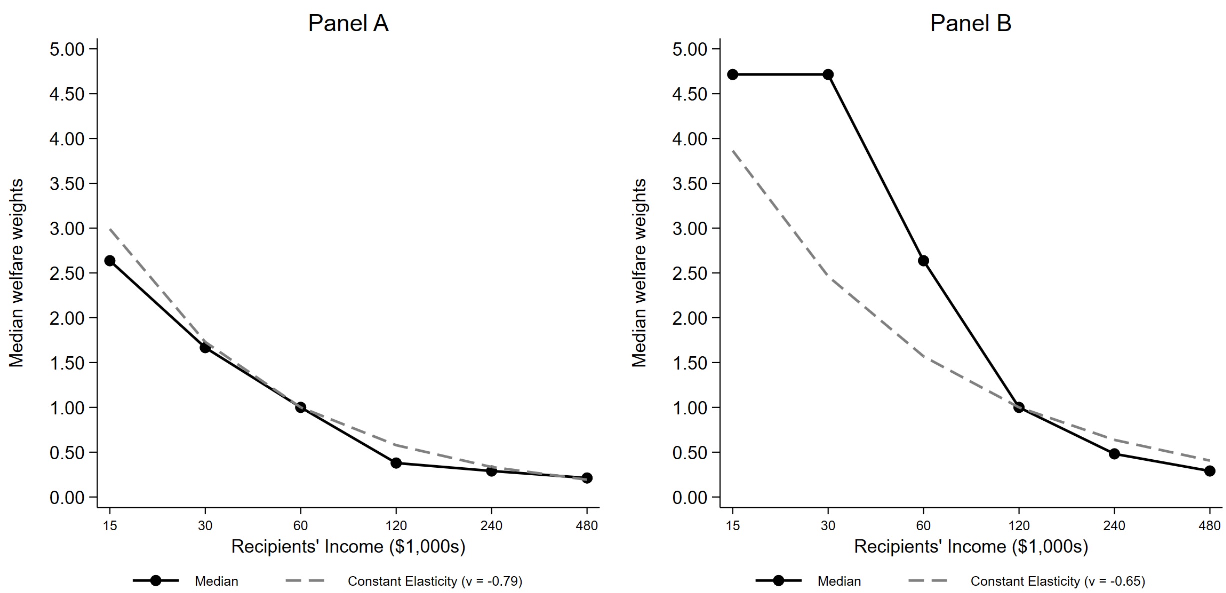

Figure 3: Median Welfare Weight Schedule Across Income

Notes: The solid line presents the median welfare weight assigned by Social Architects at each Recipient income level, with the common Recipient's weight normalized to 1. Panel (a) includes treatments in which the common Recipient earns \$60,000, and Panel (b) includes treatments in which the common Recipient earns \$120,000. The gray dashed lines show the fitted constant-elasticity schedule $g(c) \propto c^{\nu}$, where c denotes Recipient income and ν is the median of the individual elasticity estimates ν_i obtained from Equation (3). The unweighted median estimates are $\nu = -0.79$ in Panel (a) and $\nu = -0.65$ in Panel (b).

While the aggregate schedule is progressive, there is considerable heterogeneity at the individual level: 40% of Social Architects assign welfare weights that are weakly decreasing in Recipients' incomes (i.e., weakly progressive), while 8% assign weights that are weakly increasing (i.e., weakly regressive). The remaining Social Architects have non-monotonic weights.[20]

[20]Non-monotonic weights are not necessarily due to measurement error. They could also be due to various normative ideals. For example, some Social Architects may be broadly progressive but assign a lower weight to very low-income individuals, whom they perceive as "lazy" (Drenik and Perez-Truglia, 2018).

4.2.2 Parametric Approach

The non-parametric schedules in Figure 3 show how aggregate welfare weights vary over incomes without imposing a functional form. However, for comparisons across demographic subgroups, treatments, and robustness specifications, it is useful to summarize each Social Architect's welfare-weight schedule with a single progressivity parameter.

Specifically, for Social Architect i, let g_{ir} denote the welfare weight assigned to Recipient r, whose disposable income is c_r. We approximate the welfare weights schedule using a constant-elasticity specification $g_{ir} \propto c_r^{\nu_i}$, which is commonly used in the optimal taxation literature (e.g., Saez, 2002; Allcott et al., 2019). The parameter ν_i is the income elasticity of welfare weights: $\nu_i < 0$ indicates progressive weights (weights decline with income), while $\nu_i > 0$ indicates regressive weights (weights increase with income). Under a utilitarian social welfare function with an additively separable utility function that exhibits constant relative risk aversion (CRRA) in consumption, $-\nu_i$ corresponds to the coefficient of relative risk aversion. To estimate ν_i for each Social Architect, we model the conditional mean of assigned welfare weights as

$$E[g_{ir} \mid c_r] = \exp(\alpha_i + \nu_i \log(c_r)), \tag{3}$$

where α_i is an intercept that absorbs the normalization of welfare weights. We estimate Equation (3) using Poisson pseudo-maximum likelihood (PPML).[21] We aggregate estimates across Social Architects by reporting the median ν_i.

To assess whether the parametric constant-elasticity specification provides a useful summary of the non-parametric schedules, we overlay the fitted parametric schedule implied by the median estimate of ν on the median non-parametric schedule. The fitted schedules are normalized to 1 at the common Recipient's income, matching the normalization of the non-parametric schedules. The unweighted median estimate is $\nu = -0.79$ in the \$60,000 common-Recipient treatment and $\nu = -0.65$ in the \$120,000 common-Recipient treatment. As Figure 3 shows, the fitted schedule (dashed line) closely follows the non-parametric schedule (solid line) in the \$60,000 treatment. In the \$120,000 treatment, the fitted schedule matches the non-parametric schedule well at and above \$120,000, but understates the welfare weights assigned to lower-income Recipients.

Appendix Section E.1 reports two additional diagnostics that support using the constant-elasticity specification as a parsimonious summary of the welfare-weight schedules. First, relative to linear and constant-semi-elasticity alternatives, the constant-elasticity specifica-

[21] We use PPML rather than OLS on the log-linearized version of Equation (3), since log-linear OLS can be severely biased in the presence of heteroskedasticity (Silva and Tenreyro, 2006).

tion achieves the best in-sample fit based on the root-mean-squared error (RMSE). Second, re-estimating ν after dropping each Recipient income in turn yields median estimates close to the full-sample medians, though the estimates vary somewhat more in the \$120,000 common-Recipient treatment.

For tractability, we pool data across the 60K and 120K common Recipient treatments and apply sampling weights to obtain our benchmark median estimate of ν; we assess the sensitivity of our estimates to the common-Recipient income in Section 4.5. We also report 95% confidence intervals of the median obtained via percentile bootstrap; hereafter we refer to these as 95% CIs.[22] The pooled weighted median progressivity estimate is $\nu = -0.71\,(95\%\ \text{CI} = [-0.78, -0.64])$.[23] Under the constant elasticity specification, a doubling of income multiplies the welfare weight by $2^{-0.71} = 0.61$. This estimate implies that a Social Architect who assigns a welfare weight of 1 to a Recipient assigns a welfare weight of 0.61 to a Recipient earning twice as much. In monetary terms, the Social Architect is indifferent between taking \$1 from a Recipient (decreasing welfare by $0.61 \cdot \$1$) and giving \$0.61 to a Recipient earning half as much (increasing welfare by $1 \cdot \$0.61$). Since these welfare weights were elicited over disposable incomes under the current tax-and-transfer system, our finding of progressive welfare weights implies that, abstracting from efficiency costs, Social Architects favor additional redistribution at the margin.

While the median estimate is progressive, the distribution of individual estimates reveals both broad agreement on the direction of welfare weights and meaningful heterogeneity in the magnitude of the estimates. Figure 4 presents the cumulative distribution of individual progressivity estimates ν_i. Most Social Architects have progressive estimates: 78% assign progressive weights ($\nu_i < 0$), while 22% assign regressive weights ($\nu_i > 0$). These estimates are unchanged when we apply sampling weights. However, the distribution is dispersed: individual estimates range from -4.18 to 4.00. These patterns point to meaningful heterogeneity across Social Architects, but the extent of dispersion should be interpreted cautiously because each ν_i is estimated from six data points.

4.3 Welfare Weights and Background Characteristics

In this section, we explore the role of Social Architects' background characteristics in predicting the observed heterogeneity in the progressivity of welfare weights. Table 4 presents median regression estimates of progressivity (ν_i) on background characteristics.

Republicans assign less progressive welfare weights relative to Democrats and Indepen-

[22] The bootstrap resamples respondents with replacement and recomputes the weighted median using the original sampling weights. Thus, the intervals are conditional on the weights and do not account for uncertainty in their construction.

[23] The unweighted median is identical: $\nu = -0.71\,(95\%\ \text{CI} = [-0.75, -0.66])$.

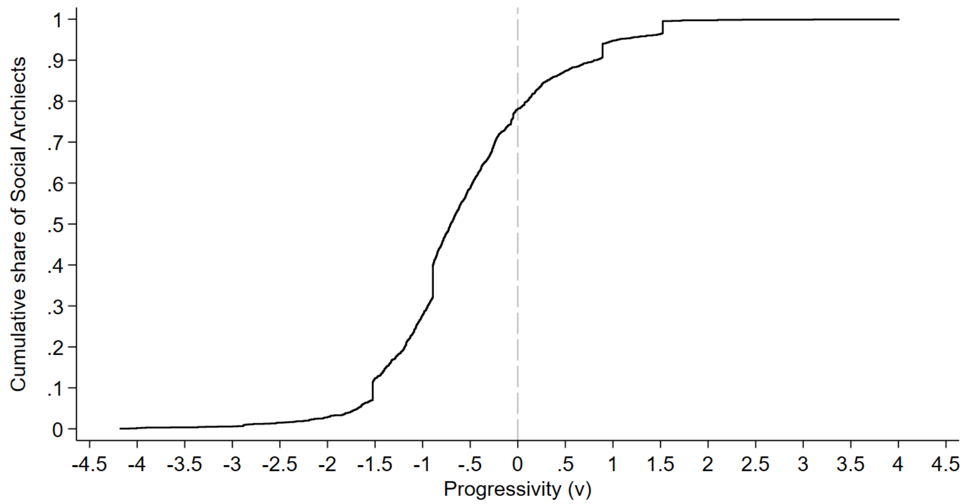

Figure 4: Distribution of Progressivity of Welfare Weights

Notes: The figure presents the cumulative distribution function (CDF) of the progressivity of welfare weights (ν_i) using data from Wave 1. Values below zero indicate progressive welfare weights; values above zero indicate regressive welfare weights.

dents: the conditional median progressivity estimate among Republicans is 0.12 points higher than among non-Republicans. This result is consistent with the partisan gap in support for government redistribution observed in the literature (e.g., Alesina et al., 2018; Stantcheva, 2021; Kuziemko et al., 2015). While Republicans assign less progressive weights on aggregate, the median Republican estimate is progressive.

Social Architects with above median income (above $54,000) assign less progressive weights, consistent with the income gap in support for redistribution observed in the literature (e.g., Singhal, 2008; Cohn et al., 2023). However, the observed effect of income is not statistically significant. While income is a weak predictor of overall progressivity, Social Architects tend to assign a higher weight to Recipients with incomes similar to their own relative to other Recipients—a pattern that is most robust among lower-income Social Architects (see Appendix Section E.2).[24]

Among the remaining demographic characteristics, above-median aged (55 years or older) Social Architects assign less progressive welfare weights. We do not find a statistically significant gender or education difference.

[24]This is consistent with an explanation of self-interest motives or alternative explanations, such as Social Architects having different normative views regarding Recipients with similar incomes.

Table 4: Welfare Weights and Background Characteristics

	(1)
Republican	0.117**
	(0.052)
High Income	0.067
	(0.050)
High Age	0.137***
	(0.049)
High Education	-0.090
	(0.056)
Male	0.019
	(0.047)
Constant	-0.790***
	(0.045)
Observations	1996

Notes: The table presents coefficient estimates from a median regression. The dependent variable is the progressivity of the welfare weights (ν_i). *Republican* equals 1 for Republicans and 0 otherwise. *High Income* equals 1 for Social Architects with above-median income and 0 otherwise. *High Age* equals 1 for Social Architects with above-median age and 0 otherwise. *High Education* equals 1 for Social Architects with above-median education and 0 otherwise. *Male* equals 1 if a Social Architect is male and 0 otherwise. The regression uses data from Wave 1. Standard errors in parentheses.
$^{*}p < 0.10, ^{**}p < 0.05, ^{***}p < 0.01$.

4.4 Assessing Robustness to Quality of Responses

4.4.1 Indicators of Response-Quality

We assess the robustness of the estimated welfare weights to individual-level response quality using four pre-registered indicators. We flag responses as low-quality when participants (i) fail one or more comprehension questions on the first attempt, (ii) exhibit extreme completion times, (iii) provide internally inconsistent welfare weights (type classification—i.e., progressive or regressive—differs across the identical third and sixth comparisons), or (iv) report low confidence in their decisions. About half the Social Architects are flagged by one or more indicators, with most of these being flagged by only one indicator (see Appendix Figure A7).[25] On our indicator of consistency—which provides a direct measure of response quality—we find that 89% of Social Architects retain their type classification (progressive or

[25] Overall, about 37% of Social Architects are flagged by only one response-quality indicator, 11% by two indicators, 2% by three indicators, and 0.2% by all four indicators. Overlap across pairs of indicators is limited (see Appendix Figure A7).

regressive) across the identical third and sixth comparisons. Using a quantitative measure of consistency, the correlation in $\tilde{g}$ across the two comparisons is 0.84 ($p < 0.01$).[26]

Table 5 presents median regression estimates of progressivity (ν_i) on the four pre-registered indicators of response quality, controlling for treatments and demographics. The median value of ν is 0.14 points higher among those who fail at least one comprehension question on their first try and 0.50 points higher among those with inconsistent welfare weights; both differences are statistically significant. It is also 0.07 points higher among Social Architects reporting low confidence, though this difference is only marginally significant. The coefficient on extreme completion time is not significant. These results suggest that Social Architects flagged as having low-quality responses tend to assign less progressive weights.

Table 5: Welfare Weights and Response Quality

	(1)
Fail comprehension	0.145**
	(0.066)
Inconsistent (sign)	0.501***
	(0.067)
Extreme time	-0.057
	(0.096)
Low confidence	0.074*
	(0.043)
Observations	1996
Controls?	Yes

Notes: The table presents coefficient estimates from a median regression. The dependent variable is the progressivity of the welfare weights (ν_i). *Fail comprehension* equals 1 if a Social Architect failed one or more comprehension questions on the first try and 0 otherwise. *Inconsistent (sign)* equals 1 if, in the identical third and sixth comparisons, a Social Architect assigns progressive welfare weights ($\tilde{g} < 1$) in one and regressive welfare weights ($\tilde{g} > 1$) in the other. *Extreme time* equals 1 if a Social Architect's time spent on the survey lies beyond two standard deviations of the mean. *Low confidence* equals 1 if a Social Architect reports confidence levels lower than the highest category of "Very Much." Controls include treatment indicators (those specified in Table 6) and demographic controls (those specified in Table 4). The regression uses data from Wave 1. Standard errors in parentheses.
$^{*}p < 0.10, ^{**}p < 0.05, ^{***}p < 0.01$.

[26] A simple summary of the inconsistency is the step distance between the third and sixth comparisons on the staircase grid of $\tilde{g}$: 53% of Social Architects have identical welfare weights across comparisons, 67% are within 1 step, 77% are within two steps, and 82% are within three. Appendix Figure A5 presents the full distribution. We use this measure as a robustness check for our quality adjustment (see Footnote 27).

4.4.2 Response-Quality Adjustment

We adjust the progressivity estimates for response quality following Luttmer and Samwick (2018). Each coefficient in Table 5 represents the median bias introduced by having a low-quality response. For each Social Architect whose answer is flagged as low-quality by a given indicator (indicator takes a value of 1), we subtract that indicator's estimated bias (the coefficient estimate) from the estimated ν_i. This correction yields unbiased estimates under two conditions: (1) differences in true underlying progressivity between those with high- and low-quality responses are fully captured by demographic controls (rather than by the quality indicators), and (2) the median ν among those with high-quality responses is measured without error. Applying this adjustment yields a weighted median progressivity estimate of $\nu = -0.83$, which is 0.12 points lower than our unadjusted benchmark estimate of -0.71 (see Figure 5).[27]

4.4.3 Corner Solutions

Another concern is the share of Social Architects at corner solutions. They are particularly problematic when Social Architects are consistently at a corner solution in all five comparisons. In our data, 12% always prefer the reform and 6% always prefer the status quo across all five comparisons. Some are likely genuine extreme preferences; others may be mistakes due to confusion or trembling-hand errors. Our response-quality indicators flag 14% of those who always select the reform and 19% of those who always select the status-quo as low-quality.

4.5 Assessing Robustness to Design Features

To assess the robustness of welfare weights to the features of the experimental design, we implemented several treatments that vary the features across Social Architects. Table 6 presents median regression estimates of progressivity (ν_i) on treatment indicators. Column (1) includes only treatment indicators as explanatory variables, while Column (2) additionally includes background characteristics as controls.

The first three treatment dimensions vary the features of the decision environment. The placement of the Recipients on the screen (left vs. right) does not have a significant effect on the progressivity of weights. The order of the comparisons has a small effect on the progressivity of the weights, but it is only marginally significant without controls and insignificant with controls. In contrast, the common Recipient across comparisons matters:

[27] As a robustness check, we replace the qualitative inconsistency indicator with a stricter quantitative measure that flags any non-identical response across the third and sixth comparisons, comprising 47% of Social Architects. Using this indicator, the quality-adjusted weighted progressivity estimate is $\nu = -0.83$, which is identical to the estimate obtained using the qualitative inconsistency measure.

welfare weights are more progressive when the common Recipient earns \$60,000 rather than \$120,000 (ν is 0.13 points lower, or 0.15 with controls). Because the two treatments contain the same set of high-to-low income ratios, this difference reflects sensitivity to features of the comparison other than relative income ratios, such as the relative position of the common Recipient.[28]

The remaining two treatment dimensions vary the features of the staircase method, addressing two potential limitations of the method—anchoring and mistakes. We find evidence of anchoring: welfare weights are more progressive when the first decision in the staircase is more progressive (\$300, −\$700) than when it is less progressive (\$700, −\$300), with ν 0.18–0.19 points lower in both columns. We also find evidence of mistakes: welfare weights are more progressive when Social Architects are asked to verify the consequences of their choices compared to when they are not (ν is 0.07 points lower and marginally significant, and 0.09 lower with controls and significant). This result indicates that some less progressive choices are likely mistakes.[29] We find very similar results when considering univariate regression in which the treatment indicators enter the regressions separately (see Appendix Table A5).

4.5.1 Reconciling Estimates Across Frames

The treatment effects documented above show that elicited welfare weights vary systematically across some treatments, raising the question of which estimate should be used for welfare analysis. We reconcile estimates following the framework of Bernheim and Rangel (2009), which provides choice-theoretic foundations for welfare analysis when choices depend on features of the decision environment (i.e., "frames"). Drawing on this framework and subsequent empirical work, we use three strategies: (i) when no frame is welfare-relevant, we embrace the normative ambiguity and report bounds across frames (e.g., Goldin and Reck, 2022); (ii) when a frame reduces mistakes, the choices in that frame are welfare-relevant (e.g., Chetty et al., 2009; Allcott and Taubinsky, 2015; Ambuehl et al., 2022); (iii) when variation across frames identifies a bias under a model, we estimate and correct the bias (e.g., Luttmer and Samwick, 2018; Chetty et al., 2009; Allcott et al., 2022). Figure 5 presents the resulting progressivity estimates.

Among the frames that vary the income of the common Recipient across comparisons,

[28]In the \$60K common-Recipient treatments, the common Recipient is the lower-earner in 3 of 5 comparisons, compared with 2 of 5 in the \$120K common-Recipient treatments. If respondents are more responsive to comparisons in which the common Recipient is at a disadvantage—consistent with inequity aversion (Fehr and Schmidt, 1999)—this asymmetry can pull weights toward progressivity in the \$60K treatment.

[29]An alternative explanation is that the verification induces experimental demand. To mitigate demand effects, we explicitly informed Social Architects before they began the task that they would be prompted in the first decision of each comparison. Moreover, we did not explicitly ask them to change their choice when there was an inconsistency (which might induce demand effects). Instead, we asked them to make their choice and indicate their preferred final distribution once again.

Table 6: Treatment Effects

	(1)	(2)
Low-income recipient on Left	-0.044	-0.026
	(0.045)	(0.047)
Common recipient 60K	-0.131***	-0.150***
	(0.045)	(0.047)
Order of comparisons	-0.076*	-0.045
	(0.045)	(0.047)
First decision (300, -700)	-0.188***	-0.182***
	(0.045)	(0.047)
Verification	-0.075*	-0.093**
	(0.045)	(0.047)
Constant	-0.448***	-0.545***
	(0.055)	(0.071)
Observations	1996	1996
Controls?	No	Yes

Notes: The table presents coefficient estimates from a median regression. The dependent variable is the progressivity of the welfare weights (ν_i). The explanatory variables include treatment indicators and demographic controls (those specified in Table 4). The regressions use data from Wave 1. Standard errors in parentheses.
$^{*}p < 0.10,^{**} p < 0.05,^{***} p < 0.01$.

there is no obvious welfare-relevant frame. This is supported by results showing that indicators of response quality and temporal stability are not different across the frames (see Appendix Table A7). We embrace the normative ambiguity about the welfare-relevant frame and report bounds on the progressivity estimates. We estimate the weighted median value of ν separately in the \$60K and \$120K common Recipient treatments.[30]

When frames differ in whether Social Architects are asked to verify the consequences of their choice, we treat the frame without the consequence verification as one in which some choices reflect mistakes and the frame with the verification—because verification plausibly reduces such mistakes—as the "welfare-relevant" frame. Choices in the verification frame are more progressive: the weighted median of ν is -0.78, about 0.07 points below the benchmark (-0.71). Estimating within this frame separately by treatments with varying common Recipients gives our bounds: the weighted median is -0.85 when the common Recipient earns \$60,000 and -0.71 when it earns \$120,000, i.e., $\nu \in [-0.85, -0.71]$.

Finally, frames that randomly vary the first decision of the staircase provide evidence of anchoring, which we interpret as bias rather than welfare-relevant variation in preferences.

[30]The weighted median value of ν when the common Recipient earns \$60,000 is -0.83 (95% CI = $[-0.89, -0.72]$), and when the common Recipient earns \$120,000 is -0.64 (95% CI = $[-0.72, -0.55]$).

We correct the progressivity estimates for anchoring following Luttmer and Samwick (2018) and Bursztyn et al. (2025). In particular, we assume that the reported reform at indifference is a weighted average of the starting value and the true indifference reform, and use the randomized variation in the starting value to back out the true reform (details in Appendix Section E.3). The estimated anchoring is small in each comparison (Appendix Table A11); in the pooled sample, correcting for it moves the median progressivity only from $\nu = -0.71$ to $\nu = -0.76$. Re-estimating the corrected median separately in the \$60K and \$120K common-Recipient treatments yields $\nu \in [-0.87, -0.69]$, essentially identical to the bounds obtained from frames with consequence verifications of $[-0.85, -0.71]$.

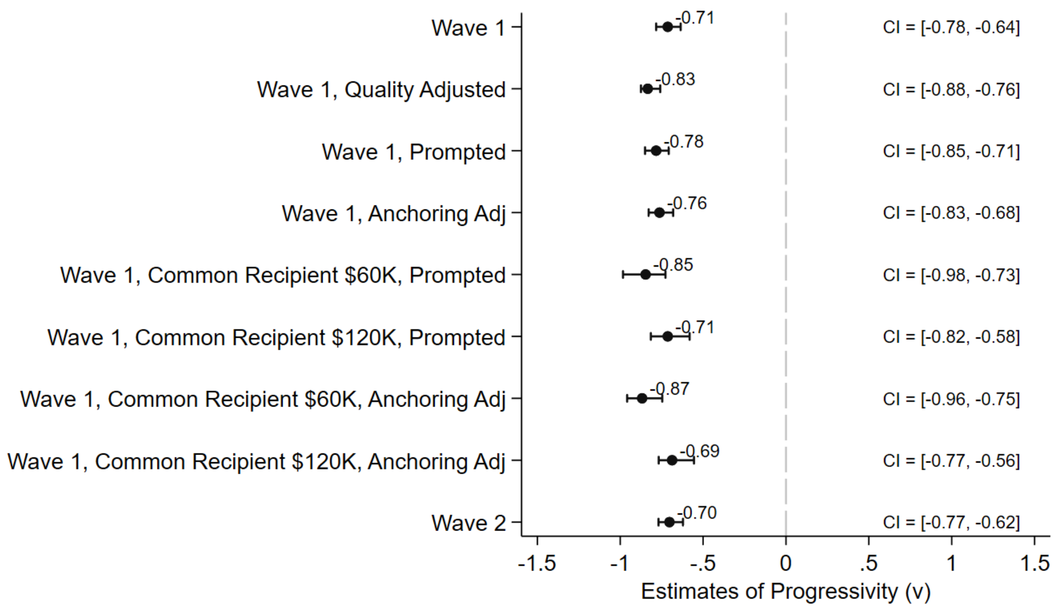

Figure 5: Estimates of Median Progressivity (ν)

Notes: The figure presents median progressivity estimates ν from various specifications. All estimates are weighted using sampling weights to match the averages of the specified sample to the population averages. Confidence intervals are 95% confidence intervals obtained via percentile bootstrap with 1,000 weighted replications.

4.6 Summary and Implications of Progressivity Estimates

The bounds on progressivity estimates obtained across robustness to design features (discussed in Section 4.5.1) are given by $\nu \in [-0.85, -0.71]$. The response-quality adjusted estimate given by $\nu = -0.83$ (discussed in Section 4.4) also lies within these bounds. We compare these estimates to common benchmarks: $|\nu| = 0.25$ is "weak," $|\nu| = 1$ is "fairly strong," and $|\nu| = 4$ is "extremely strong" redistributive tastes (Saez, 2002; Allcott et al., 2019). Our estimates ($|\nu| \in [0.71, 0.85]$) fall below the "fairly strong" threshold, which corresponds to logarithmic utility and is commonly used in the optimal-tax literature. In Section 5, we also compare them to the weights implied by current U.S. tax and transfer policies, and find the elicited weights to be substantially more progressive.

These welfare weights can be implemented in optimal policy formulas using the constant-elasticity function c^{ν} with $\nu \in [-0.85, -0.71]$, or CRRA utility with a coefficient of relative risk aversion of $\gamma = -\nu \in [0.71, 0.85]$. As an application, we calibrate the optimal non-linear labor income tax formula of Saez (2001) (details on methodology and results are in Appendix Section E.6). The less progressive endpoint of the bounds ($\nu = -0.71$) yields an average optimal MTR of 56%, while the more progressive endpoint ($\nu = -0.85$) yields 60%. Thus, the bounds of progressivity estimates produce relatively tight bounds on optimal marginal tax rates.

The preceding sections established the progressivity of welfare weights and their robustness. In the next sections, we turn to understanding the validity (Section 4.7) and temporal stability (Section 4.8) of welfare weights, before comparing them to the weights implied by current U.S. tax and transfer policies (Section 5).

4.7 Welfare Weights and Support for Redistribution

We validate our measure of welfare weights by comparing it with two measures of support for government redistribution. The first—general support for redistribution—measures support for reducing income differences between the rich and the poor (ranging from 1 to 7), which we reverse-code so that higher values indicate stronger support. The second—support for redistribution at the margin—measures support for redistribution beyond that achieved by the current tax and transfer system (ranging from -2 to 2), with positive (negative) values indicating a desire for redistribution from high-income (low/middle-income) individuals to low/middle-income (high-income) individuals and a value of zero indicating a desire for no additional redistribution. Social Architects support increasing government redistribution on both measures (See Appendix Figures A8 and A9).

We test the individual-level correlation using linear regressions presented in Table 7. In Column (1), the outcome variable is standardized general support for redistribution, and in Column (3), the outcome variable is standardized support for redistribution at the margin. The explanatory variable is the standardized progressivity estimate, ν_i. Social Architects with more progressive welfare weights (lower values of ν_i) exhibit stronger support for redistribution on both measures, and the associated coefficients are highly statistically significant. The relationship is approximately linear and not driven by extreme values (Appendix Figures A10). Based on the R^2, welfare weights account for about 11–16% of the variation in support across the two measures presented in Columns (1) and (3). For comparison, stated political

Table 7: Correlation with Support for Redistribution

	General redistribution		Redistribute at margin	
	Continuous	Binary	Continuous	Binary
	(1)	(2)	(3)	(4)
$\nu - std$	-0.398***		-0.335***	
	(0.023)		(0.020)	
$1(\nu < 0)$		0.369***		0.443***
		(0.026)		(0.025)
Constant	0.000	0.371***	0.000	0.399***
	(0.021)	(0.023)	(0.021)	(0.023)
Observations	1996	1996	1996	1996
R^2	0.158	0.104	0.112	0.177
Outcome var	Continuous	Binary	Continuous	Binary
Explanatory var	Continuous	Binary	Continuous	Binary

Notes: The table presents coefficient estimates from linear regressions. Columns (1) and (2) use general support for redistribution—support for the government to do something to reduce income differences between the rich and poor; this ranges from 1–7, recoded so that higher values indicate stronger support. Columns (3) and (4) use support for redistribution at the margin—support for redistribution beyond that achieved by the current tax and transfer system; this ranges from -2 to $+2$, where positive (negative) values indicate a desire for redistribution from high-income (low/middle-income) to low/middle-income (high-income) individuals, and zero indicates no desired additional redistribution. Columns (1) and (3) regress standardized measure of support on standardized progressivity of welfare weights; Columns (2) and (4) regress binary indicators of support—values above 0 for the redistribution at the margin measure and values above 4 for the general support for redistribution measure—on an indicator equal to 1 if $\nu_i < 0$ (progressive weights) and 0 otherwise. All standardized variables have mean 0 and standard deviation 1. The regression uses data from Wave 1. HC3 standard errors in parentheses.
$^{*}p < 0.10, ^{**}p < 0.05, ^{***}p < 0.01$.

affiliation—a strong predictor of redistributive support (e.g., Stantcheva, 2021; Kuziemko et al., 2015)—explains only about 3–8% of the variation in support for redistribution (Appendix Table A6). We find similarly strong relationships when we use binary versions of the outcomes (split at each outcome's midpoint) and ν_i (split at 0), presented in Columns (2) and (4).

Having examined individual-level correlations, we now compare aggregate (unweighted) shares. The appropriate comparison for welfare weights is support for redistribution at the margin rather than the general support for redistribution measure: the former, like welfare weights, is measured at the margin and allows for both progressive and regressive redistribution, whereas the latter captures broader attitudes toward redistribution. About 78% of Social Architects assign progressive welfare weights—implying a preference to redistribute at the margin—closely matching the 74% of Social Architects supporting additional progressive government redistribution at the margin.[31]

Overall, the close alignment between welfare weights and support for government redistribution, both in individual-level correlations and in aggregate shares, helps validate our measure of welfare weights.[32]

4.8 Temporal Stability

To examine the temporal stability of welfare weights, we collected data from the same Social Architects in two waves, fielded four weeks apart. We find a high degree of temporal stability in type classifications: 81% of Social Architects retain the same type classification (progressive or regressive weights) across waves. Among the 19% who do not, 56% transition from regressive to progressive weights and 44% in the opposite direction, consistent with switching reflecting noise rather than systematic changes in preferences. To explore the continuous measure of progressivity, Figure 6 presents a bin scatter plot of ν_i across the two waves, using the sample of Social Architects with valid responses in both waves. The points line up along a clear upward gradient, and the Pearson correlation of ν_i across waves is 0.55 ($p < 0.01$), indicating substantial temporal stability.

The cross-wave correlation in welfare weights of 0.55 is similar to the estimates found in the literature.[33] It is also similar to the stability of two survey measures of redistribution that we fielded in both waves. Specifically, the cross-wave correlation for welfare weights exceeds the correlation for support for redistribution at the margin ($0.44, p < 0.01$)—the more directly comparable item—but is below the correlation for general support for redistribution ($0.72, p < 0.01$), which captures a broader attitudinal measure.

Next, we examine whether welfare weights are stable at the aggregate level. Aggregate stability is more informative for policy than individual stability, because stable aggregate measures can be used in policy evaluation even if individual measures are unstable due to

[31]For the general redistribution measure, 66% report support above the mid-point of 4 on the 7-point scale; this share rises to 78% if respondents at the midpoint are included as supporters.

[32]In Appendix Section E.4, we conduct an additional validation exercise, which tests whether welfare weights capture fairness concerns—as intended. We find that welfare weights do capture fairness concerns. However, they also capture other concerns, including beliefs about the economic costs of taxing high incomes and broader views about the appropriate scope of government. We also test whether welfare weights capture misperceptions measured in the survey. We find that welfare weights do not capture misperceptions about tax levels and the prevalence of low incomes individuals.

[33]Most evidence on the stability of social preferences comes from studies that track individuals over relatively long horizons. These studies find modest temporal correlations—0.28 in rural Paraguay (Chuang and Schechter, 2015), 0.12–0.28 in Vietnam (Carlsson et al., 2014), and 0.39–0.46 in the United States (Fisman et al., 2023). By contrast, work that looks at much shorter intervals has focused on risk and time preferences, which appear considerably more stable, with correlations of 0.55 (Schoemaker and Hershey, 1992) and 0.50 (Wölbert and Riedl, 2013) observed for risk preferences and 0.64 for time preferences (Dean and Sautmann, 2021).

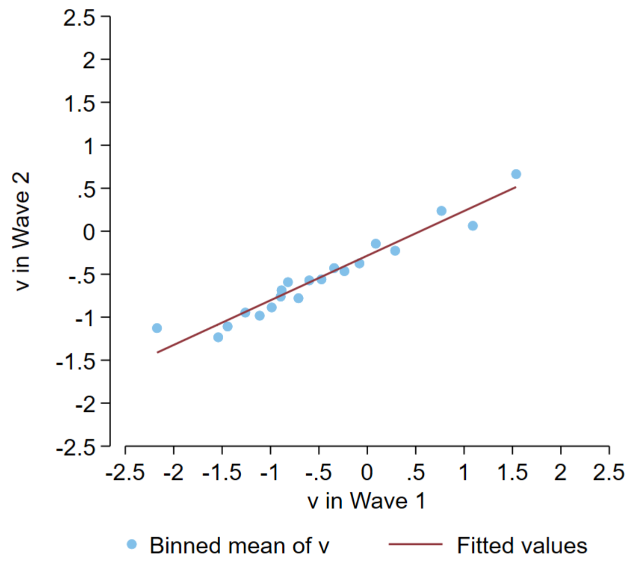

Figure 6: Temporal Stability of Welfare Weights

Notes: The figure presents a binned scatter plot of progressivity estimates ν_i in Wave 1 and Wave 2 using 20 equally-sized bins. The figure uses the sample of 1397 Social Architects with valid responses in both waves.

measurement error. The weighted median progressivity estimate in Wave 2 is $\nu = -0.70$, which is close to the Wave 1 benchmark estimate of $\nu = -0.71$ (Figure 5).

Overall, the correlation in progressivity estimates across waves and the similarity of the aggregate progressivity estimates across waves indicate considerable temporal stability of our measure of welfare weights.

5 Comparing Welfare Weights

In the previous sections, we estimated the progressivity of the elicited welfare weights, bounded their progressivity across our robustness checks, and showed them to be valid and temporally stable. In this section, we compare the welfare weights elicited in our experiment to those implied by the U.S. income tax schedule and by broader tax and transfer policies. The welfare weights implied by policies—referred to as "inverse optimum weights"—represent the weights that can rationalize current policies.

We obtain the inverse optimum weights implied by the U.S. income tax schedule from Hendren (2020). This paper reveals the implicit welfare weights from the observed marginal tax rates (MTRs), the income distribution, and estimates of the elasticity of taxable income (ETI). The baseline estimate of inverse-optimum weights uses MTRs in 2012 and mid-range ETI estimates of about 0.3 (Details in Section E.5). The resulting weights are decreasing with incomes but increase at the top percentile (see Appendix Figure A11).

The inverse optimum weights implied by U.S. tax and transfer policies are derived using the framework of Hendren and Sprung-Keyser (2020). In this framework, the welfare weight $g(z)$ assigned to beneficiaries of a policy earning z is the inverse of the policy's Marginal Value of Public Funds (MVPF)—the beneficiaries' welfare gain from \$1 of government spending. Intuitively, a planner is willing to incur efficiency losses by implementing a policy with a low MVPF if the beneficiaries are assigned a high welfare weight. We use the MVPF estimates of tax and transfer policies in Hendren and Sprung-Keyser (2020) to infer the welfare weights implied by these policies (Details in Section E.5). The resulting weights are roughly decreasing with the beneficiaries' incomes (see Appendix Figure A12).

The inverse-optimum weights and the weights we elicit are defined over different income concepts: pre-tax income z for the former, disposable income c for the latter. To compare them, we express both over disposable income. Let h denote the tax-and-transfer schedule, so $c = h(z)$. Since h is strictly increasing, it is invertible, and each c corresponds to a unique $z = h^{-1}(c)$. We can thus express the inverse-optimum weights over disposable income as $g\big(h^{-1}(c)\big)$. With both sets now expressed over disposable income, we summarize each with a constant-elasticity specification, $g(c) \propto c^{\nu}$.

We find that the estimated progressivity is $\nu = -0.09$ (SE $= 0.01, 95\%$ CI $= [-0.12, -0.07]$) for the tax schedule and $\nu = -0.16$ (SE $= 0.08, 95\%$ CI $= [-0.31, 0.00]$) for tax and transfer policies. The welfare weights elicited using our experiment ($\nu \in [-0.85, -0.71]$) are about 8–9 times as progressive as the inverse-optimum weights implied by the income tax schedule and about 4–5 times as progressive as those implied by tax and transfer policies. Figure 7 presents the distribution of welfare weights implied by these estimates against the disposable income distribution. The discrepancy between the elicited weights and those implied by current tax and transfer policy is concentrated at the bottom of the income distribution, where the elicited weights are considerably larger.

There are at least four classes of explanations for why the weights elicited in our experiment are more progressive than those implied by policies. First, citizens' welfare weights elicited in the experiment may differ from those considered by politicians. For example, politicians may have misperceptions about citizens' views (e.g., Broockman and Skovron, 2018; Lucas et al., 2025; Sheffer et al., 2024). Second, even if the welfare weights elicited in our experiment align with those considered by politicians, the political process may aggregate citizens' weights differently from the population aggregation used in our analysis. For example, politicians may weigh only certain subgroups (e.g., voters) or overweight high-income groups. Third, the inverse-optimum weights might reflect inputs and considerations beyond the redistributive preferences we elicit. For example, if politicians use higher estimates of elasticity of taxable income than those in Hendren (2020), the inverse optimum

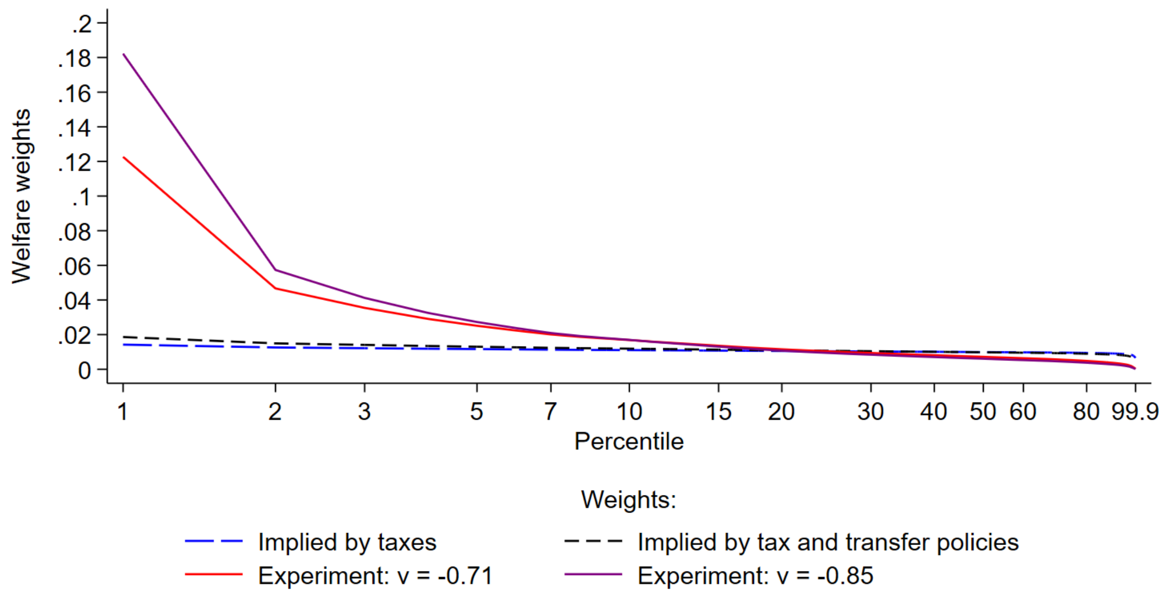

Figure 7: Comparison of Welfare Weights

Notes: The figure plots welfare weights against percentiles of the disposable income distribution. The x-axis has a natural log spacing. Welfare weights are interpolated using c^{ν}, where c is disposable income, and re-normalized to sum to 1. The figure plots the inverse-optimum weights implied by the income tax schedule ($\nu = -0.09$), inverse-optimum weights implied by tax and transfer policies ($\nu = -0.16$), and the two endpoints of the progressivity bounds obtained from the experiment ($\nu = -0.71$ and $\nu = -0.85$). Data on disposable incomes (which includes in-cash and in-kind transfers) is obtained from Piketty et al. (2018); individuals with negative disposable income are excluded, and each individual is treated as a single filer.

weights would be less progressive. As another example, considerations such as aversion to government, which are not accounted for in the elicited welfare weights by design, might lead the inverse-optimum weights to be less progressive Rodemeier and Sun (2025). Finally, other political economy frictions may distort the translation of preferences into policy.

6 Discussion

We develop a portable method to elicit welfare weights from real-stakes online experiments and apply it to a sample of the U.S. general population. The aggregate welfare weights are progressive, with bounds on the income elasticity of $\nu \in [-0.85, -0.71]$ across our robustness specifications. The elicited weights are also valid and temporally stable. Comparing these estimates to the weights implied by policies, we find that the weights elicited in the experiment are about 8–9 times as progressive as the inverse-optimum weights implied by the U.S. income tax schedule and about 4–5 times as progressive as those implied by tax and transfer policies.

Several challenges arise when using the elicited welfare weights for policy evaluation.

First, the welfare weights estimated in our experiment cannot be directly used to evaluate non-marginal reforms.[34] For non-marginal reforms, the marginal value of the first dollar may differ from the marginal value of the last dollar. We assume a parametric form for the weights—in which welfare weights are a function of consumption—which allows evaluating large reforms, under the assumption that shifts in consumption induced by the reform capture the relevant shifts in Recipient characteristics that Social Architects value.

Second, since our elicited welfare weights depend only on income, they apply directly to policies conditioned on income. Evaluating policies conditioned on other characteristics—such as age, gender, or ethnicity—requires the additional assumption that welfare weights are invariant to those characteristics. Future research can explore how welfare weights differ across non-income characteristics.

Finally, the elicited welfare weights may not be applicable across time, countries, and policy domains. There is evidence in the literature suggesting that people's support for redistribution may differ across time (Fisman et al., 2015), countries (e.g., Almås et al., 2020; Falk et al., 2018), and policy domains (Srinivasan et al., 2025). Future work can test whether welfare weights differ across time, countries, and policy domains.

References

Abdellaoui, M., Kemel, E., Panin, A. and Vieider, F. M. (2019), 'Measuring time and risk preferences in an integrated framework', *Games and Economic Behavior* **115**, 459–469.

Alesina, A. and Angeletos, G.-M. (2005), 'Fairness and redistribution', *American Economic Review* **95**(4), 960–980.

Alesina, A., Stantcheva, S. and Teso, E. (2018), 'Intergenerational mobility and preferences for redistribution', *American Economic Review* **108**(2), 521–554.

Allcott, H., Kim, J., Taubinsky, D. and Zinman, J. (2022), 'Are high-interest loans predatory? Theory and evidence from payday lending', *The Review of Economic Studies* **89**(3), 1041–1084.

Allcott, H., Lockwood, B. B. and Taubinsky, D. (2019), 'Regressive sin taxes, with an application to the optimal soda tax', *The Quarterly Journal of Economics* **134**(3), 1557–1626.

Allcott, H. and Taubinsky, D. (2015), 'Evaluating behaviorally motivated policy: Experimental evidence from the lightbulb market', *American Economic Review* **105**(8), 2501–2538.

Almås, I., Cappelen, A. W. and Tungodden, B. (2020), 'Cutthroat capitalism versus cuddly socialism: Are americans more meritocratic and efficiency-seeking than scandinavians?', *Journal of Political Economy* **128**(5), 1753–1788.

[34] It is worth noting that alternative approaches, such as using the inverse-optimum weights implied by policies (Hendren, 2020), also cannot be used to evaluate non-marginal reforms.

Ambuehl, S., Bernheim, B. D. and Lusardi, A. (2022), 'Evaluating deliberative competence: A simple method with an application to financial choice', *American Economic Review* **112**(11), 3584–3626.

Bernheim, B. D. and Rangel, A. (2009), 'Beyond revealed preference: choice-theoretic foundations for behavioral welfare economics', *The Quarterly Journal of Economics* **124**(1), 51–104.

Berry, J., Fischer, G. and Guiteras, R. (2020), 'Eliciting and utilizing willingness to pay: Evidence from field trials in northern Ghana', *Journal of Political Economy* **128**(4), 1436–1473.

Bourguignon, F. and Spadaro, A. (2012), 'Tax–benefit revealed social preferences', *The Journal of Economic Inequality* **10**(1), 75–108.

Broockman, D. E. and Skovron, C. (2018), 'Bias in perceptions of public opinion among political elites', *American Political Science Review* **112**(3), 542–563.

Burchardi, K. B., De Quidt, J., Gulesci, S., Lerva, B. and Tripodi, S. (2021), 'Testing willingness to pay elicitation mechanisms in the field: Evidence from Uganda', *Journal of Development Economics* **152**, 102701.

Bursztyn, L., Handel, B., Jiménez-Durán, R. and Roth, C. (2025), 'When product markets become collective traps: The case of social media', *American Economic Review* **115**(12), 4105–4136.

Cappelen, A. W., Konow, J., Sørensen, E. Ø. and Tungodden, B. (2013), 'Just luck: An experimental study of risk-taking and fairness', *American Economic Review* **103**(4), 1398–1413.

Carlsson, F., Johansson-Stenman, O. and Nam, P. K. (2014), 'Social preferences are stable over long periods of time', *Journal of Public Economics* **117**, 104–114.

Chen, D. L., Schonger, M. and Wickens, C. (2016), 'oTree—An open-source platform for laboratory, online, and field experiments', *Journal of Behavioral and Experimental Finance* **9**, 88–97.

Chetty, R., Looney, A. and Kroft, K. (2009), 'Salience and taxation: Theory and evidence', *American Economic Review* **99**(4), 1145–1177.

Chuang, Y. and Schechter, L. (2015), 'Stability of experimental and survey measures of risk, time, and social preferences: A review and some new results', *Journal of Development Economics* **117**, 151–170.

Cohn, A., Jessen, L. J., Klašnja, M. and Smeets, P. (2023), 'Wealthy Americans and redistribution: The role of fairness preferences', *Journal of Public Economics* **225**, 104977.

Cruces, G., Perez-Truglia, R. and Tetaz, M. (2013), 'Biased perceptions of income distribution and preferences for redistribution: Evidence from a survey experiment', *Journal of Public Economics* **98**, 100–112.

Dean, M. and Sautmann, A. (2021), 'Credit constraints and the measurement of time preferences', *Review of Economics and Statistics* **103**(1), 119–135.

Downs, A. (1957), 'An economic theory of political action in a democracy', *Journal of Political Economy* **65**(2), 135–150.

Drenik, A. and Perez-Truglia, R. (2018), 'Sympathy for the diligent and the demand for workfare', *Journal of Economic Behavior & Organization* **153**, 77–102.

Durante, R., Putterman, L. and Van der Weele, J. (2014), 'Preferences for redistribution and perception of fairness: An experimental study', *Journal of the European Economic Association* **12**(4), 1059–1086.

Enke, B., Graeber, T. and Oprea, R. (2023), 'Confidence, self-selection, and bias in the aggregate', *American Economic Review* **113**(7), 1933–1966.

Falk, A., Becker, A., Dohmen, T., Enke, B., Huffman, D. and Sunde, U. (2018), 'Global evidence on economic preferences', *The Quarterly Journal of Economics* **133**(4), 1645–1692.

Fehr, E., Epper, T. and Senn, J. (2024), 'Social preferences and redistributive politics', *Review of Economics and Statistics* pp. 1–45.

Fehr, E. and Schmidt, K. M. (1999), 'A theory of fairness, competition, and cooperation', *The Quarterly Journal of Economics* **114**(3), 817–868.

Fisman, R., Jakiela, P. and Kariv, S. (2015), 'How did distributional preferences change during the great recession?', *Journal of Public Economics* **128**, 84–95.

Fisman, R., Jakiela, P., Kariv, S. and Vannutelli, S. (2023), 'The distributional preferences of Americans, 2013–2016', *Experimental Economics* **26**(4), 727–748.

Fisman, R., Kariv, S. and Markovits, D. (2007), 'Individual preferences for giving', *American Economic Review* **97**(5), 1858–1876.

Fleurbaey, M. and Maniquet, F. (2006), 'Fair income tax', *The Review of Economic Studies* **73**(1), 55–83.

Goldin, J. and Reck, D. (2022), 'Optimal defaults with normative ambiguity', *Review of Economics and Statistics* **104**(1), 17–33.

Hendren, N. (2020), 'Measuring economic efficiency using inverse-optimum weights', *Journal of Public Economics* **187**, 104198.

Hendren, N. and Sprung-Keyser, B. (2020), 'A unified welfare analysis of government policies', *The Quarterly Journal of Economics* **135**(3), 1209–1318.

Jack, B. K., McDermott, K. and Sautmann, A. (2022), 'Multiple price lists for willingness to pay elicitation', *Journal of Development Economics* **159**, 102977.

Jacobs, B., Jongen, E. L. and Zoutman, F. T. (2017), 'Revealed social preferences of Dutch political parties', *Journal of Public Economics* **156**, 81–100.

Karadja, M., Mollerstrom, J. and Seim, D. (2017), 'Richer (and holier) than thou? The effect of relative income improvements on demand for redistribution', *Review of Economics and Statistics* **99**(2), 201–212.

Kuziemko, I., Norton, M. I., Saez, E. and Stantcheva, S. (2015), 'How elastic are preferences for redistribution? Evidence from randomized survey experiments', *American Economic Review* **105**(4), 1478–1508.

Lockwood, B. B. and Weinzierl, M. (2016), 'Positive and normative judgments implicit in US tax policy, and the costs of unequal growth and recessions', *Journal of Monetary Economics* **77**, 30–47.

Lucas, J., Sheffer, L., Loewen, P. J., Walgrave, S., Soontjens, K., Amsalem, E., Bailer, S., Brack, N., Breunig, C., Bundi, P. et al. (2025), 'Politicians' theories of voting behavior', *American Political Science Review* **119**(3), 1304–1321.

Luttmer, E. F. P. and Samwick, A. A. (2018), 'The welfare cost of perceived policy uncertainty: Evidence from social security', *American Economic Review* **108**(2), 275–307.

Luttmer, E. F. P. and Singhal, M. (2011), 'Culture, context, and the taste for redistribution', *American Economic Journal: Economic Policy* **3**(1), 157–179.

Rodemeier, M. and Sun, G. (2025), 'positive versus normative disagreement in optimal income taxation'. SSRN Working Paper 5687044.

Ruggles, S., Flood, S., Sobek, M., Backman, D., Cooper, G., Rivera Drew, J. A., Richards, S., Rodgers, R., Schroeder, J. and Williams, K. C. (2025), 'IPUMS USA: Version 16.0 [dataset]'.
URL: *https://doi.org/10.18128/D010.V16.0*

Saez, E. (2001), 'Using elasticities to derive optimal income tax rates', *The Review of Economic Studies* **68**(1), 205–229.

Saez, E. (2002), 'Optimal income transfer programs: intensive versus extensive labor supply responses', *The Quarterly Journal of Economics* **117**(3), 1039–1073.

Saez, E. and Stantcheva, S. (2016), 'Generalized social marginal welfare weights for optimal tax theory', *American Economic Review* **106**(1), 24–45.

Schoemaker, P. J. and Hershey, J. C. (1992), 'Utility measurement: Signal, noise, and bias', *Organizational Behavior and Human Decision Processes* **52**(3), 397–424.

Sheffer, L., Loewen, P. J. and Lucas, J. (2024), 'Long-term policymaking and politicians' beliefs about voters: Evidence from a 3-year panel study of politicians', *Governance* **37**(2), 395–410.

Silva, J. M. C. S. and Tenreyro, S. (2006), 'The log of gravity', *The Review of Economics and Statistics* **88**(4), 641–658.

Singhal, M. (2008), 'Quantifying preferences for redistribution'. Working Paper.

Srinivasan, K., Valasek, J. and Wang, W. (2025), 'Fair (p)redistribution'. SSRN Working Paper 6464018.

Stantcheva, S. (2016), 'Comment on "Positive and normative judgments implicit in US tax policy and the costs of unequal growth and recessions" by Benjamin Lockwood and Matthew Weinzierl', *Journal of Monetary Economics* **77**, 48–52.

Stantcheva, S. (2021), 'Understanding tax policy: How do people reason?', *The Quarterly Journal of Economics* **136**(4), 2309–2369.

Stantcheva, S. (2023), 'How to run surveys: A guide to creating your own identifying variation and revealing the invisible', *Annual Review of Economics* **15**(1), 205–234.

Toubia, O., Johnson, E., Evgeniou, T. and Delquié, P. (2013), 'Dynamic experiments for estimating preferences: An adaptive method of eliciting time and risk parameters', *Management Science* **59**(3), 613–640.

Weinzierl, M. (2014), 'The promise of positive optimal taxation: normative diversity and a role for equal sacrifice', *Journal of Public Economics* **118**, 128–142.

Wölbert, E. and Riedl, A. (2013), 'Measuring time and risk preferences: Reliability, stability, domain specificity'. Working Paper.

Appendix

Who Should Get Money? Estimating Welfare Weights in the U.S.

Francesco Capozza

Krishna Srinivasan

Table of Contents

A Variable Definitions

Misperceptions

We elicit Social Architects' perceptions about the level of taxes paid by individuals, and the share of individuals with incomes below $15,000.

Overestimate taxes: We ask Social Architects five questions designed to elicit their perceptions about the level of taxes paid by individuals in society. In particular, Social Architects are asked about their beliefs regarding (i) the share of households in the top tax bracket, (ii) the top marginal tax rate, (iii) the average tax rate of those in the top tax bracket, (iv) the share of households who pay no taxes, and (v) the average tax rate of households with the median income. We focus on perceptions along these five dimensions because they were the most predictive of people's support for redistribution in Stantcheva (2021). Social Architects can select a number from 0 to 100 using a slider for each of the five questions. We identify misperceptions in each of the five variables as follows.

- *Gap in top-taxes* = *Beliefs about top-taxes* - 32.7
- *Gap in top-MTR* = Beliefs about top MTR - 37
- *Gap in top-share* = *Beliefs about top-share* - 0.73
- *Gap in non-filers* = 44 - *Beliefs about non-filers*
- *Gap in median-taxes* = *Beliefs about median-taxes* - 13

We take the actual values from Stantcheva (2021). We orient the gap in non-filers such that a higher gap in non-filers corresponds to an overestimation in the level of taxes paid. We standardize each of the five misperception variables such that they have a mean of 0 and a standard deviation of 1. Then, we create an index by taking the equally weighted average of the five standardized variables and standardizing the resulting variable.

Overestimate low-income share: We elicit Social Architects' beliefs about the share of individuals earning less than $15,000. This income level corresponds to the income of the Recipient with the lowest income in our experiment. Social Architects can select a number from 0 to 100 using a slider. We identify Architects' misperceptions by subtracting the actual value (11) from their responses. We obtain the actual value by looking at the share of individuals whose disposable income is below $15,000 in the data obtained from Piketty et al. (2018) (variable *diinc*). Finally, we standardize the misperceptions.

Views about Taxes and Government

We ask Social Architects several questions that elicit their views about the tax system and their trust in government, each capturing a unique mechanism that may explain people's support for redistribution. Many of these questions are drawn from Stantcheva (2021).

Behavioral responses high earners: Takes a value of 1 if a Social Architect indicates that the extent to which taxing high-income earners would encourage them to work less is "A moderate amount," "A lot," or "A great deal," and a value of 0 if the Social Architect indicates "A little," or "None at all."

Higher taxes high-incomes hurt economy: Takes a value of 1 if a Social Architect indicates that taxing high-income earners would "Hurt economic activity in the U.S.," and a value of 0 if the Social Architect indicates "Not have an effect on economic activity in the U.S." or "Help economic activity in the U.S."

No belief trickle down: Takes a value of 1 if a Social Architect indicates that the lower class and working class would "Mostly lose" or "Neither lose nor win" if taxes on high-income earners were cut and a value of 0 if the Social Architect indicates that they would "Mostly win."

Fair distribution of income: Takes a value of 1 if a Social Architect indicates that the current income distribution is "Somewhat fair" or "Very fair" and a value of 0 if the Social Architect indicates "Neither unfair nor fair," "Somewhat unfair," or "Very unfair."

Low trust in government: Takes a value of 1 if a Social Architect trusts the government "Only some of the time" or "Never" and a value of 0 for responses "Most of the time" or "Just about always."

Govt should do less: Takes a value of 1 if a Social Architect indicates that "Government is doing too much" or "Government is doing just the right amount" to solve the country's problems and a value of 0 for "Government should do more."

B Pre-registration

The experimental design, sample restrictions, and analyses were pre-registered. We list the following deviations from the pre-analysis plan:

1. We pre-registered the inclusion of participants with an approval rating between 95 and 100 on Prolific. However, the "representative sample" option on Prolific does not permit screening by approval rate. Matching our data with Prolific's data on participants shows that our sample in Wave 1 includes participants with approval ratings between 93 and 100.

2. For our response-quality adjustment, we pre-registered three separate proxies indicating failure on each of the three comprehension checks. The proportions failing each check are 8%, 3%, and 2%, respectively. To improve power, we instead use a single indicator that flags failure on any of the three checks, which flags 11% of responses.

3. We pre-registered flagging respondents with unusual completion times (± 2 standard deviations from the mean) separately by treatment. In Wave 1, we additionally construct this indicator separately within the soft launch (293 participants) and the full launch, because server congestion slowed the soft launch—the median completion time was 24.4 minutes in the soft launch, versus 14 minutes in the full launch.

4. The following exploratory analyses were not pre-registered:
   - Table A6 on the effect of political affiliation on support for redistribution.
   - Table A7 on the effect of the common Recipient on response quality indicators

5. Following Stantcheva (2021), in the decomposition in Appendix Section E.4, we code all mediators in the same direction relative to the outcome. This is a sign convention for readability and leaves each variable's contribution unchanged.

6. We pre-registered recruitment quotas based on gender, ethnicity, and age. We instead used quotas based on gender, political affiliation, and age, replacing ethnicity with political affiliation.

7. We pre-registered the consistency check as examining whether welfare weights are similar in sign across the identical third and sixth comparisons. In addition to this, we construct a quantitative consistency measure based on the step distance between the two comparisons on the staircase grid of $\tilde{g}$, which we use as a robustness check for the response-quality adjustment.

C Additional Figures

Your choice on whether to implement the change (above) and your preferred distribution of final incomes (below) contradict each other. Please make your choice and indicate your preferred distribution of final incomes once again.

Please consider each question carefully because if you are selected, one of your choices may have real consequences for two real individuals.

Comparison 1, Question 1

	Person #1	Person #3
Annual disposable income	$15,000	$60,000
Proposed change	$300	-$700

Please make your decision:

- ○ I prefer to implement the change
- ○ I prefer not to implement the change

If you prefer to implement the change, the final incomes of the individuals are Person #1: $15,300 and Person #3: $59,300. If you prefer not to implement the change, the incomes of individuals remain unchanged.

Which of the following final income distributions do you prefer?

- ○ Person #1: $15,300 and Person #3: $59,300
- ○ Person #1: $15,000 and Person #3: $60,000

Figure A1: Screenshot of a Decision in Treatments with Consequence Verifications

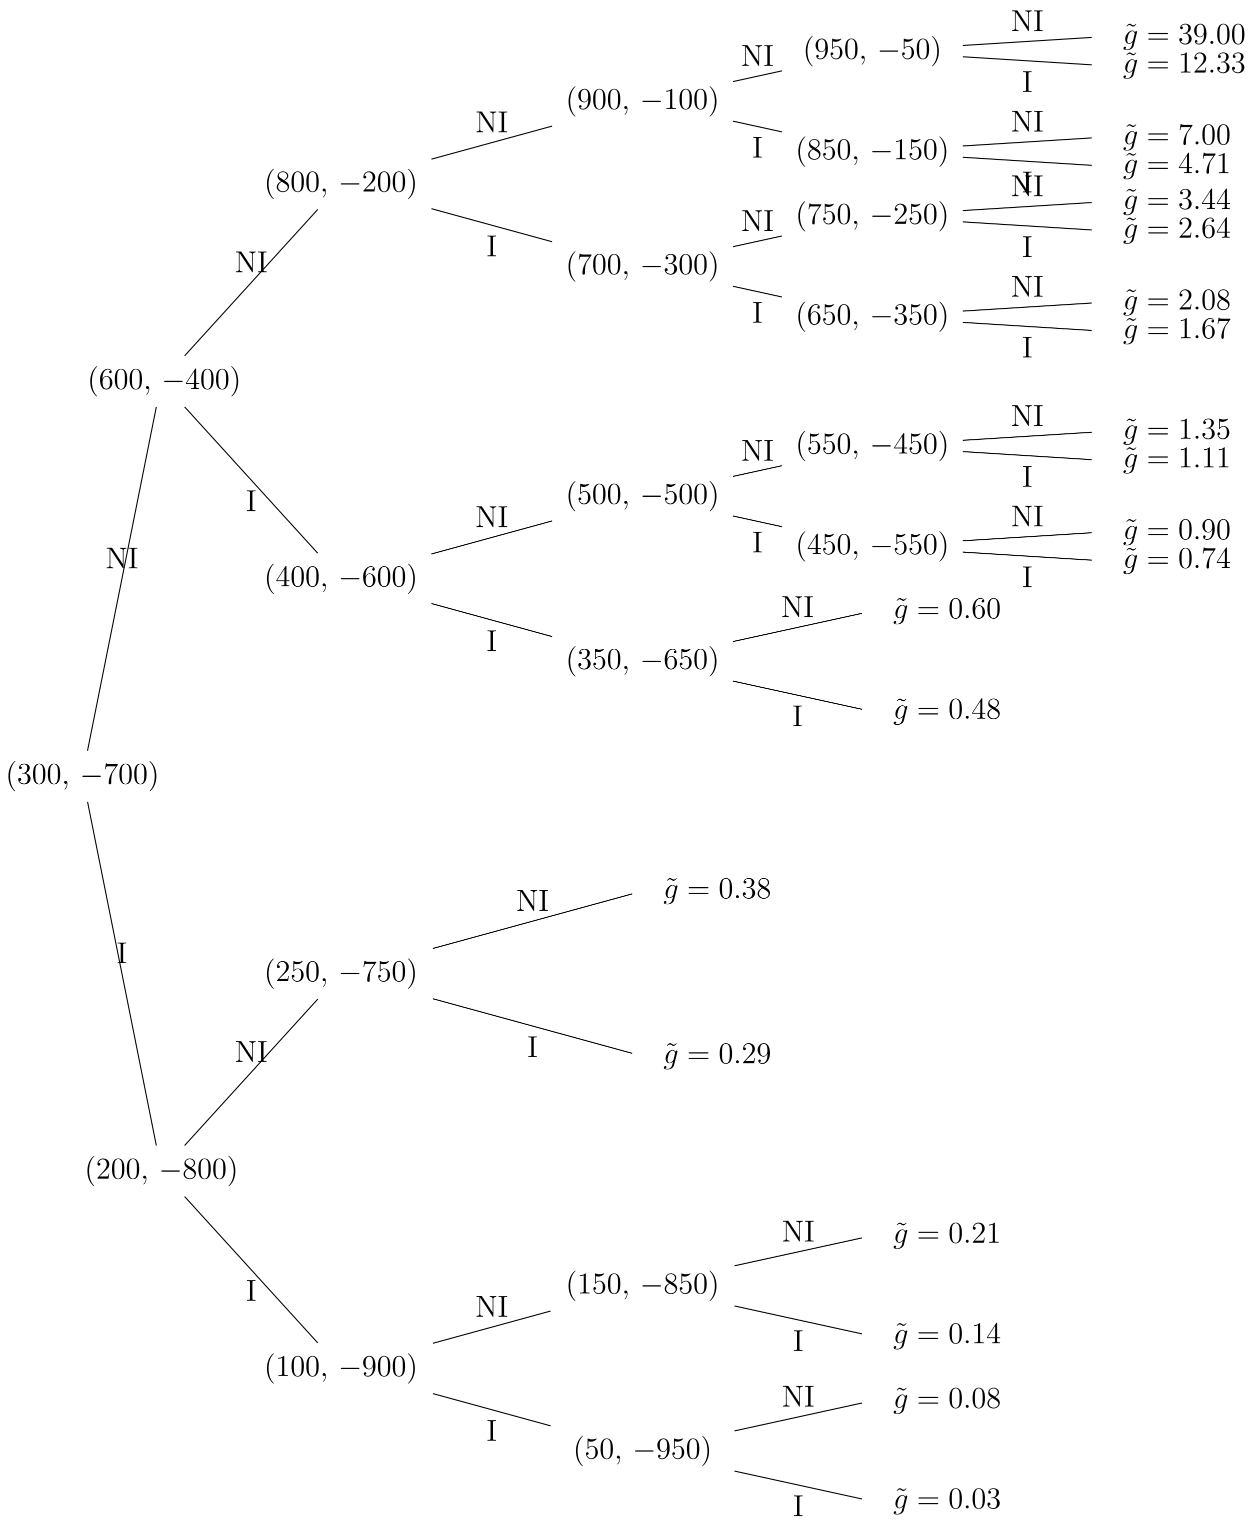

Figure A2: Reforms Selected by the Staircase Method, Starting at (300, −700)

Notes: The figure shows the reform selected by the staircase method at each node, depending on whether the reform at the previous node was implemented ("I") or not ("NI"); the first decision is (300, −700). The parameter $\tilde{g}$ at each terminal node is the implied ratio of the welfare weight on the higher-income Recipient to that on the lower-income Recipient.

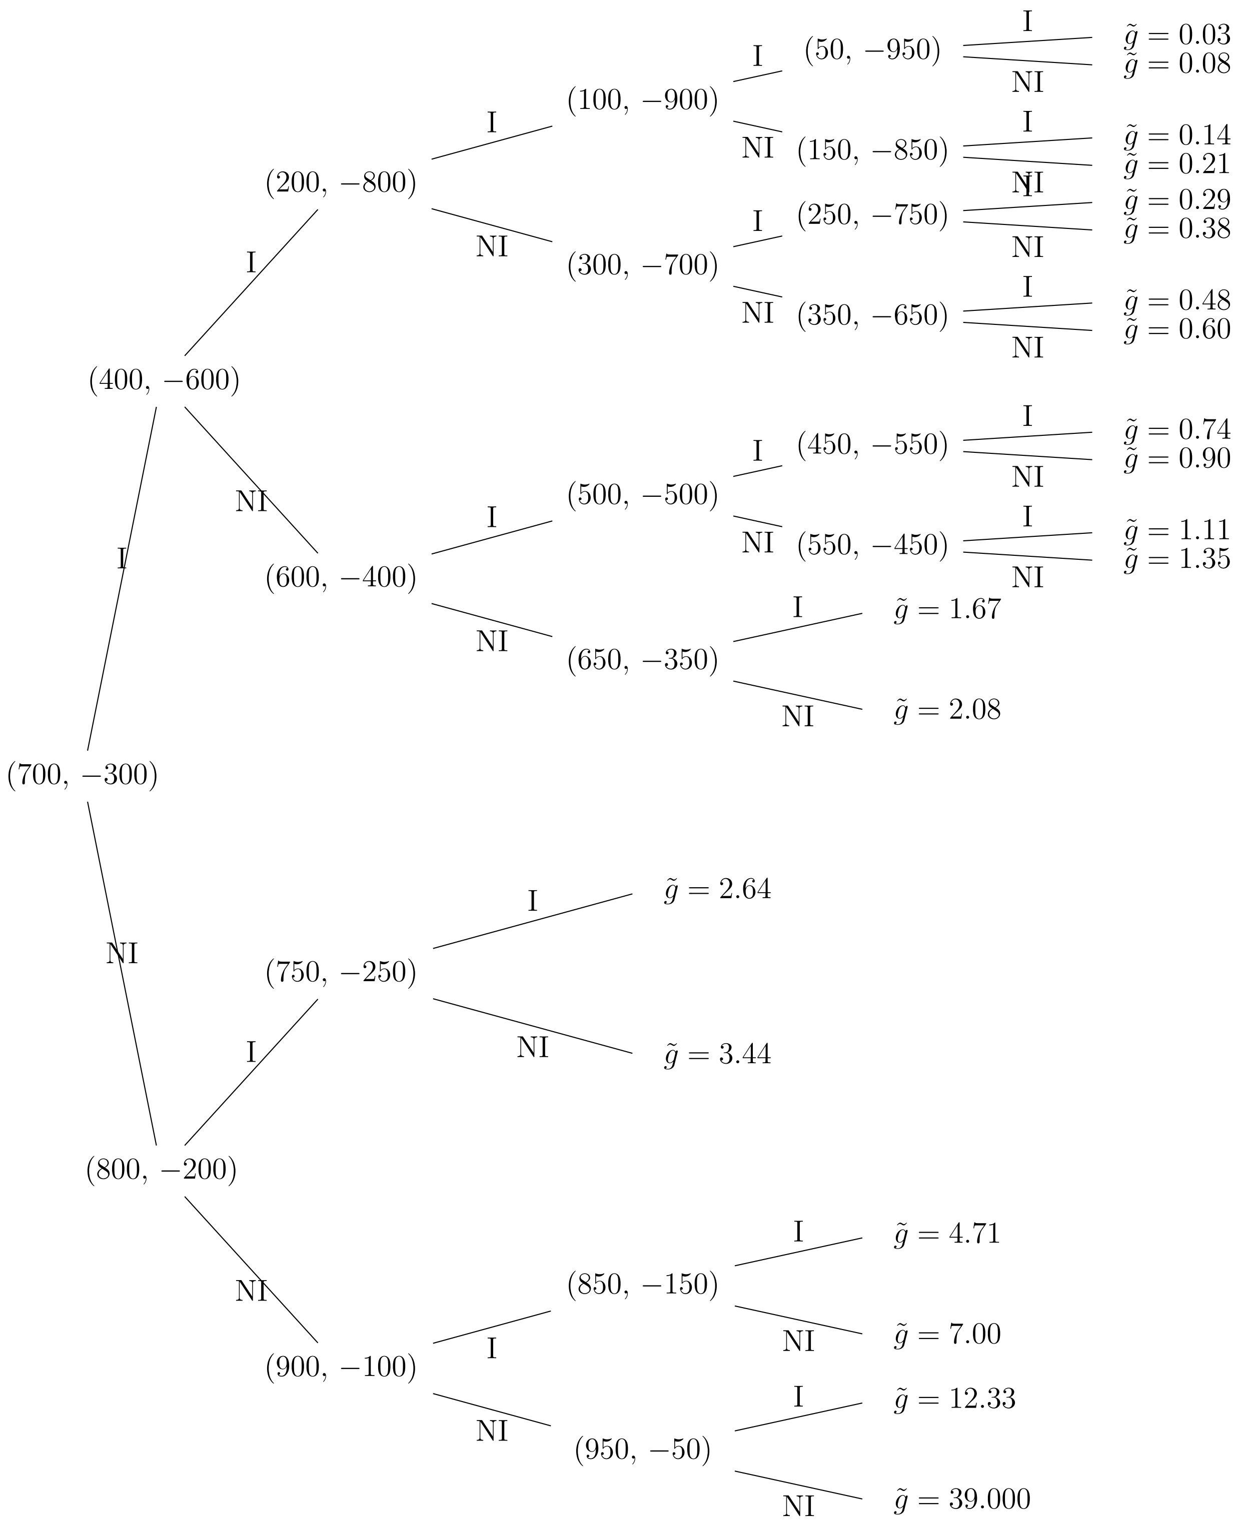

Figure A3: Reforms Selected by the Staircase Method, Starting at (700, −300)

Notes: The figure shows the reform selected by the staircase method at each node, depending on whether the reform at the previous node was implemented ("I") or not ("NI"); the first decision is (700, −300). The parameter $\tilde{g}$ at each terminal node is the implied ratio of the welfare weight on the higher-income Recipient to that on the lower-income Recipient.

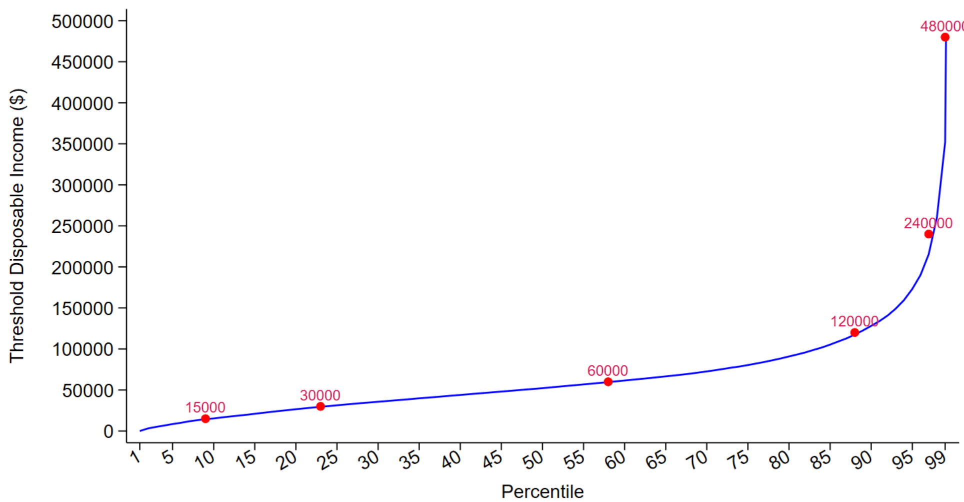

Figure A4: Disposable Incomes of the Recipients

Notes: The figure plots the incomes of the six Recipients (dots) against the disposable income distribution (line) in the U.S. in 2019. The horizontal axis indicates the percentiles, and the vertical axis indicates the threshold annual disposable incomes corresponding to the percentiles. Data on disposable incomes (which includes in-cash and in-kind transfers) is obtained from Piketty et al. (2018); individuals with negative disposable income are excluded, and each individual is treated as a single filer.

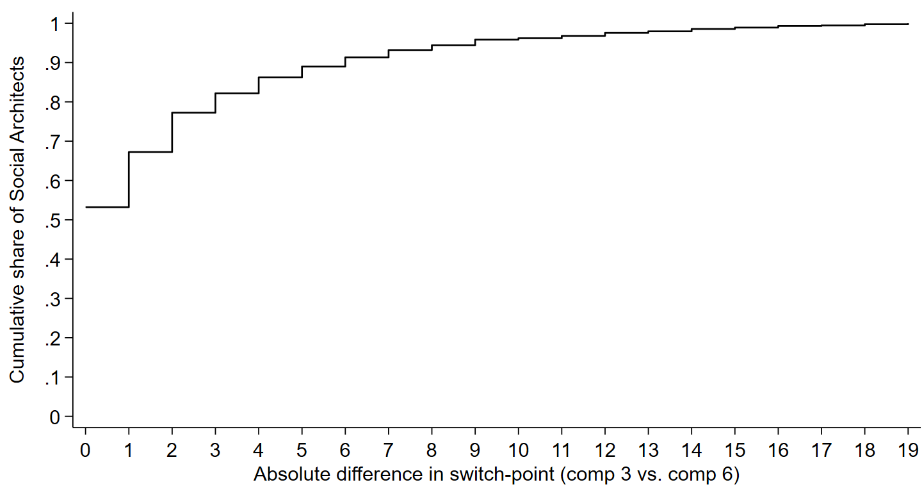

Figure A5: Consistency in Welfare Weights

Notes: The figure presents the distribution of the absolute difference in switch-points between the identical third and sixth comparisons. The 19 reforms in Table 1 define 20 possible switch-points, so the difference ranges from 0 (perfect consistency) to a maximum of 19.

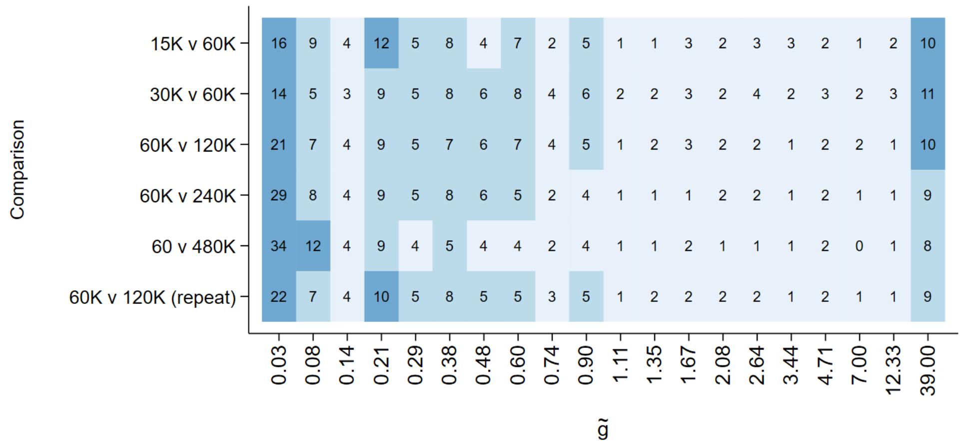

(a) Common Recipient earns $60K

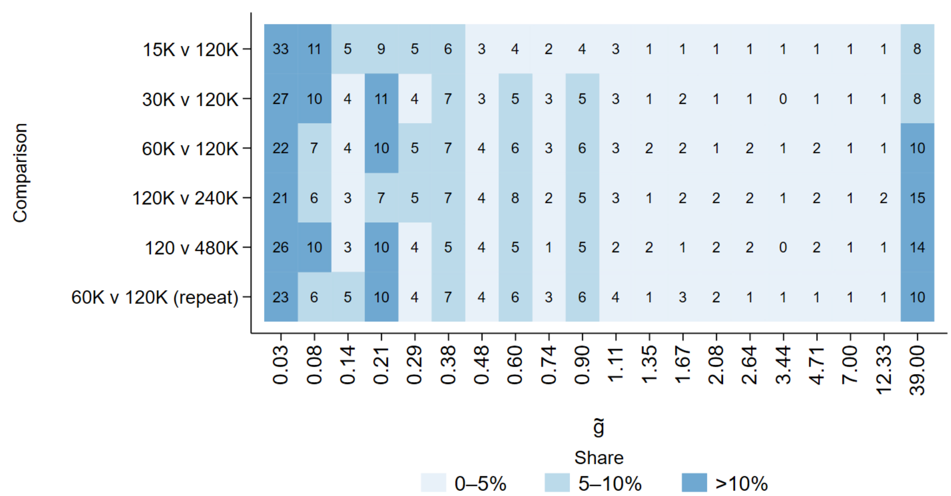

(b) Common Recipient earns $120K

Figure A6: Distribution of $\tilde{g}$ across Comparisons and Treatments

Notes: The figures present the distribution of $\tilde{g}$—the ratio of welfare weights assigned to the high-income and low-income Recipients—across the six comparisons. Each cell presents the share of participants with a given $\tilde{g}$ in a given comparison, and the shares sum to 100% in each row. The figures in both panels use data from Wave 1; Figure (a) uses data from treatments in which the common Recipient across pairs earns $60,000, while Figure (b) uses data from treatments in which the common Recipient earns $120,000.

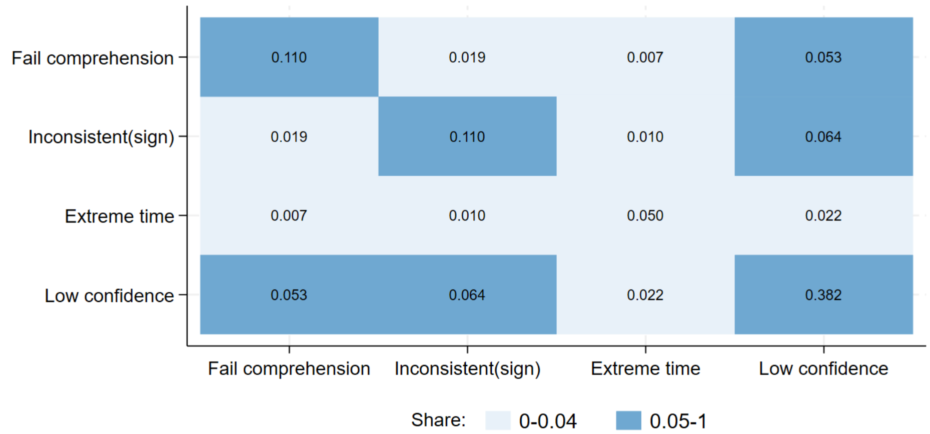

(a) Wave 1

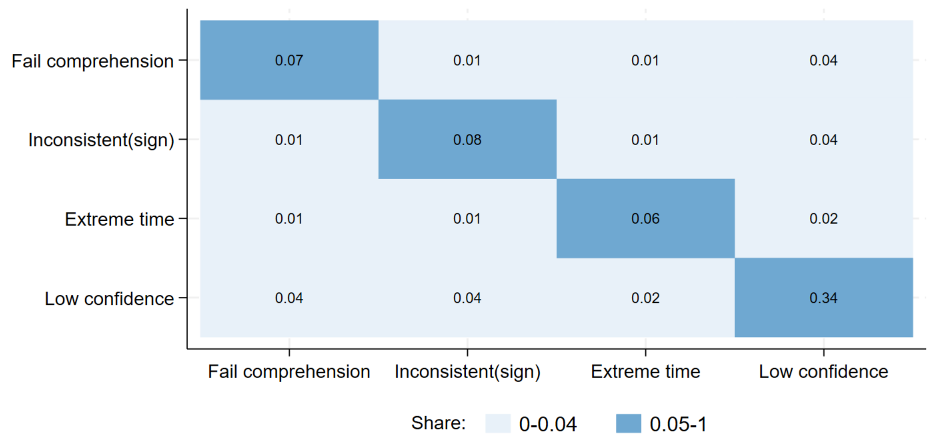

(b) Wave 2

Figure A7: Share Flagged as Low-Quality Based on Indicators

Notes: Each figure presents the share of observations where each indicator pair simultaneously equals 1; diagonals show marginal shares. *Fail comprehension* equals 1 if a Social Architect failed one or more comprehension questions on the first try. *Inconsistent (sign)* equals 1 if, in the identical third and sixth comparisons, a Social Architect assigns progressive weights ($\tilde{g} < 1$) in one and regressive ($\tilde{g} > 1$) in the other. *Extreme time* equals 1 if time spent on the survey lies beyond two standard deviations of the mean. *Low confidence* equals 1 if reported confidence is below the highest category ("Very Much"). Figure (a) uses Wave 1 data; Figure (b) uses Wave 2 data.

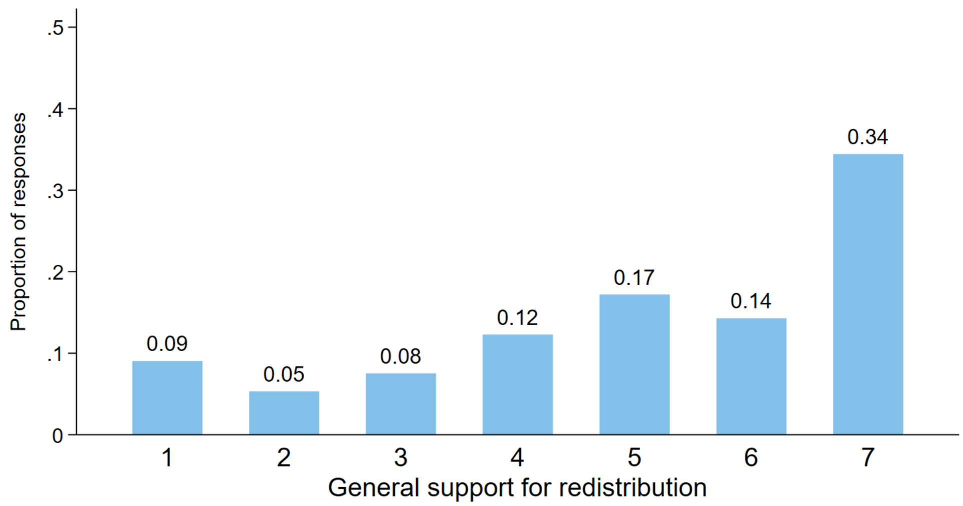

Figure A8: Distribution of General Support for Redistribution

Notes: This figure presents the frequency of responses for the general support for redistribution question. Responses take values from 1 through 7, which are recoded such that higher values indicate stronger support for the government to do something to reduce income differences between the rich and poor. The figure uses data from Wave 1.

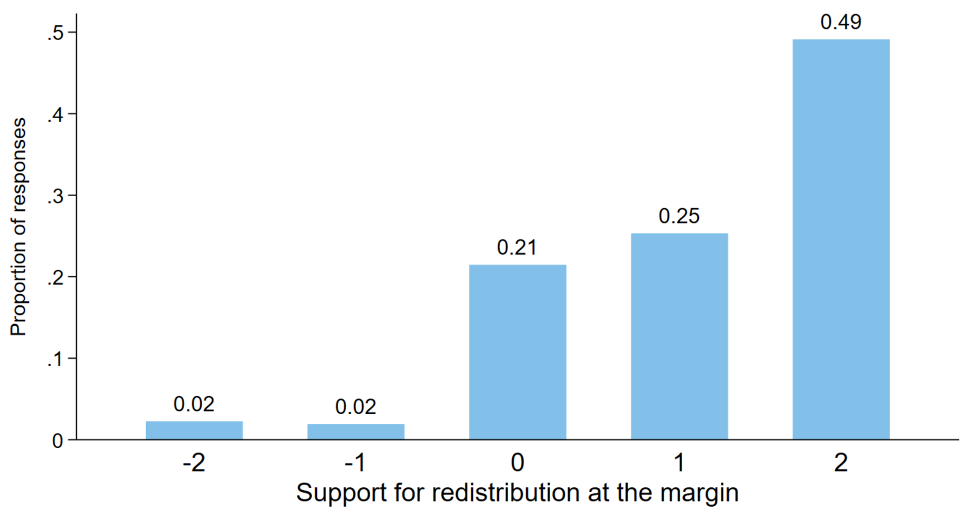

Figure A9: Distribution of Support for Redistribution at the Margin

Notes: This figure presents the frequency of responses for the support for redistribution at the margin question. Responses take values from -2 to $+2$, where positive (negative) values indicate a desire for redistribution from high-income (low/middle-income) individuals to low/middle-income (high-income) individuals. A value of zero indicates a desire for no additional redistribution. The figure uses data from Wave 1.

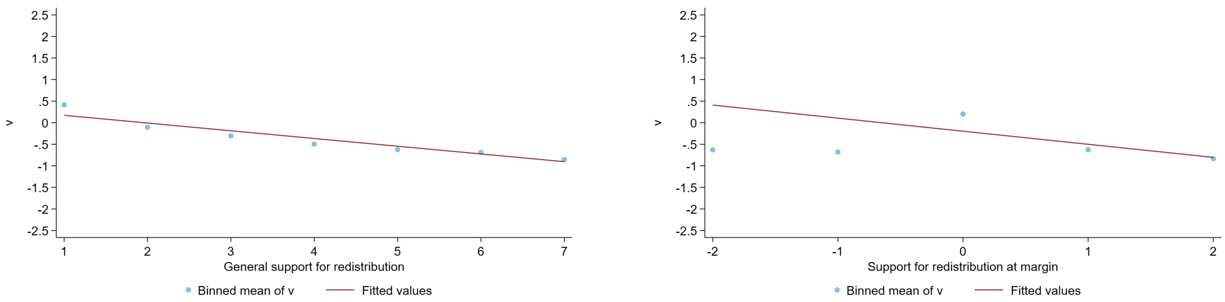

(a) ν and General Redistribution (b) ν and Redistribution at the Margin

Figure A10: Progressivity and Support for Redistribution

Notes: The figure plots the binned means of ν (y-axis) against responses to support for redistribution questions. Figure (a) uses the general support for redistribution question; support takes values from 1 through 7, which are reverse coded such that higher values indicate stronger support for the government to do something to reduce income differences between the rich and poor. Figure (b) uses the support for redistribution at the margin question; support takes values from -2 to $+2$, where positive (negative) values indicate a desire for redistribution from high-income (low/middle-income) individuals to low/middle-income (high-income) individuals and a value of zero indicates a desire for no additional redistribution. Both panels use data from Wave 1.

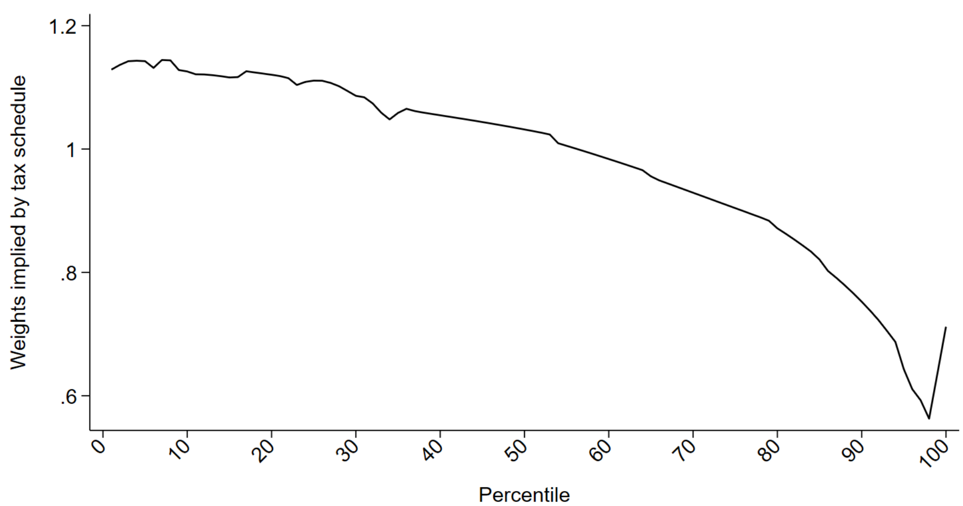

Figure A11: Welfare Weights Implied by the Tax Schedule

Notes: The figure plots the inverse-optimum welfare weights implied by the income tax schedule, obtained from Hendren (2020), for each percentile of the pre-tax income distribution. We exclude individuals with negative incomes. See Appendix Section E.5 for details.

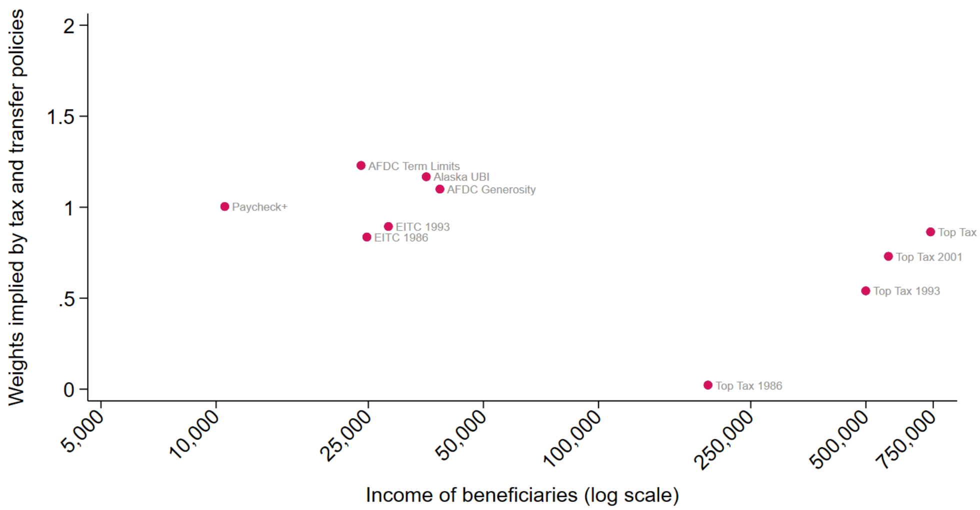

Figure A12: Welfare Weights Implied by Tax and Transfer Policies

Notes: The figure plots the inverse-optimum welfare weights implied by tax and transfer policies against the approximate incomes of the beneficiaries of the policies. The x-axis is on a log scale. The welfare weight of a policy is the inverse of its marginal value of public funds (MVPF). Data on the MVPF of the policies is obtained from Hendren and Sprung-Keyser (2020). See Appendix Section E.5 for details.

D Additional Tables

Table A1: Example U.S. Tax and Transfer Policies

Policy	**Amount**
Massachusetts one-time electric-bill credit	$50 per residential account
Credit for Other Dependents	Maximum $500
Energy-Efficient Home Improvement Credit	Range from $150 to $1,500
Georgia 2025 surplus tax rebate	Maximum range $250–$500
North Dakota Primary Residence Credit	Maximum $1,600
NYC Enhanced Real-Property Tax Credit	Maximum $500
Earned Income Tax Credit (no qualifying children, 2024)	Maximum $632

Notes: The table presents several U.S. tax and transfer policies with low reform amounts.

Table A2: Randomization Check

Panel A: Gender, income, and education

	Male	Individual income < 30,000	Individual income 30-59,999	Individual income 60-99,999	Individual income 100-149,999	Individual income ≥ 150,000	Education: High school or less	Education: Some college	Education: Bachelor or Associate	Education: Masters or above
Low-income recipient on Left	-0.008	0.012	0.019	-0.027	-0.004	-0.001	0.024*	-0.002	0.003	-0.026
	(0.022)	(0.020)	(0.020)	(0.020)	(0.014)	(0.012)	(0.014)	(0.017)	(0.022)	(0.020)
Common recipient 60K	-0.029	-0.001	0.035*	-0.014	-0.016	-0.004	0.017	0.000	-0.045**	0.027
	(0.022)	(0.020)	(0.020)	(0.020)	(0.014)	(0.012)	(0.014)	(0.017)	(0.022)	(0.020)
Order of comparisons	0.003	0.005	0.014	-0.009	-0.014	0.004	0.008	-0.026	0.001	0.016
	(0.022)	(0.020)	(0.020)	(0.020)	(0.014)	(0.012)	(0.014)	(0.017)	(0.022)	(0.020)
First decision (300, -700)	-0.051**	0.011	-0.002	-0.028	0.020	-0.001	0.006	0.007	0.010	-0.023
	(0.022)	(0.020)	(0.020)	(0.020)	(0.014)	(0.012)	(0.014)	(0.017)	(0.022)	(0.020)
Verification	-0.006	-0.018	0.020	0.009	0.001	-0.012	-0.014	0.018	-0.009	0.005
	(0.022)	(0.020)	(0.020)	(0.020)	(0.014)	(0.012)	(0.014)	(0.017)	(0.022)	(0.020)
Constant	0.536***	0.257***	0.244***	0.297***	0.124***	0.078***	0.085***	0.177***	0.472***	0.266***
	(0.027)	(0.024)	(0.024)	(0.025)	(0.018)	(0.014)	(0.016)	(0.021)	(0.027)	(0.024)
F-statistic	1.433	0.340	1.074	0.948	0.787	0.270	1.239	0.688	0.879	1.172
p-value (F-test)	0.209	0.889	0.373	0.449	0.559	0.930	0.288	0.633	0.495	0.321

Table A2: Randomization Check *(continued)*

Panel B: Age, region, and party

	Age: 18-24	Age: 25-34	Age: 35-44	Age: 45-54	Age: 55-64	Age: $\geq$ 65	Region: Northeast	Region: Midwest	Region: South	Region: West	Republican
Low-income recipient on Left	-0.004	-0.002	-0.001	0.022	0.020	-0.035**	-0.028*	0.036**	-0.010	0.002	0.008
	(0.014)	(0.017)	(0.017)	(0.016)	(0.019)	(0.015)	(0.017)	(0.017)	(0.022)	(0.018)	(0.021)
Common recipient 60K	-0.006	0.008	-0.011	-0.003	-0.004	0.015	-0.013	0.001	0.020	-0.008	-0.011
	(0.014)	(0.017)	(0.017)	(0.016)	(0.019)	(0.015)	(0.017)	(0.017)	(0.022)	(0.019)	(0.021)
Order of comparisons	0.036**	-0.012	0.012	0.007	-0.034*	-0.009	0.037**	-0.022	-0.008	-0.007	-0.001
	(0.014)	(0.017)	(0.017)	(0.016)	(0.019)	(0.015)	(0.017)	(0.017)	(0.022)	(0.018)	(0.021)
First decision (300, -700)	0.005	0.017	0.007	0.010	-0.032*	-0.007	-0.017	0.009	0.009	-0.000	0.012
	(0.014)	(0.017)	(0.017)	(0.016)	(0.019)	(0.015)	(0.017)	(0.017)	(0.022)	(0.018)	(0.021)
Verification	-0.020	0.013	0.022	-0.005	-0.006	-0.003	-0.023	0.023	-0.030	0.031*	-0.025
	(0.014)	(0.017)	(0.017)	(0.016)	(0.019)	(0.015)	(0.017)	(0.017)	(0.022)	(0.019)	(0.021)
Constant	0.110***	0.165***	0.158***	0.139***	0.274***	0.153***	0.187***	0.153***	0.452***	0.208***	0.330***
	(0.017)	(0.020)	(0.021)	(0.020)	(0.024)	(0.020)	(0.020)	(0.021)	(0.027)	(0.023)	(0.026)
F-statistic	1.663	0.456	0.564	0.509	1.335	1.391	2.190	1.516	0.634	0.667	0.442
p-value (F-test)	0.140	0.809	0.728	0.769	0.247	0.224	0.053	0.182	0.674	0.648	0.820

Notes: The table reports coefficient estimates from linear regressions of each characteristic on treatment indicators, using Wave 1 data. HC3 standard errors in parentheses. $^{*}\ p < 0.10$, $^{**}\ p < 0.05$, $^{***}\ p < 0.01$.

Table A5: Treatment Effects

	(1)	(2)	(3)	(4)	(5)	(6)
Low-income recipient on Left	-0.024					-0.044
	(0.047)					(0.045)
Common recipient 60K		-0.138***				-0.131***
		(0.046)				(0.045)
Order of comparisons			-0.081*			-0.076*
			(0.047)			(0.045)
First decision (300, -700)				-0.171***		-0.188***
				(0.045)		(0.045)
Verification					-0.066	-0.075*
					(0.047)	(0.045)
Constant	-0.689***	-0.654***	-0.668***	-0.607***	-0.668***	-0.448***
	(0.034)	(0.033)	(0.033)	(0.032)	(0.033)	(0.055)
Observations	1996	1996	1996	1996	1996	1996

Notes: The table presents coefficient estimates from a median regression. The dependent variable is the progressivity of the welfare weights (ν_i). The explanatory variables include treatment indicators. The regressions use data from Wave 1. Standard errors in parentheses.
$^{*}p < 0.10, ^{**}p < 0.05, ^{***}p < 0.01$.

Table A6: Support for Redistribution and Political Affiliation

	General redistribution (std)	Redistribute at margin (std)
	(1)	(2)
Republican	-0.616***	-0.340***
	(0.049)	(0.049)
Constant	0.198***	0.110***
	(0.024)	(0.026)
Observations	1996	1996
R^2	0.083	0.025

Notes: The table presents coefficient estimates from linear regressions. Column (1) uses general support for redistribution—support for the government to do something to reduce income differences between the rich and poor; this ranges from 1–7, recoded so that higher values indicate stronger support. Column (2) uses support for redistribution at the margin—support for redistribution beyond that achieved by the current tax and transfer system; this ranges from -2 to $+2$, where positive (negative) values indicate a desire for redistribution from high-income (low/middle-income) to low/middle-income (high-income) individuals, and zero indicates no desired additional redistribution. Both columns regress the standardized measure of support (mean 0 and standard deviation of 1) on an indicator equal to 1 for Republicans and 0 otherwise. The regression uses data from Wave 1. HC3 standard errors in parentheses.
$^{*}p < 0.10,^{**} p < 0.05,^{***} p < 0.01$.

Table A7: Quality Indicators and Treatments

	Fail comprehension	Inconsistent (sign)	Extreme Time	Low Confidence	Stable
	(1)	(2)	(3)	(4)	(5)
Common recipient 60K	-0.019	0.002	-0.008	0.029	0.033
	(0.014)	(0.014)	(0.010)	(0.022)	(0.021)
Constant	0.120***	0.109***	0.054***	0.368***	0.795***
	(0.010)	(0.010)	(0.007)	(0.015)	(0.015)
Observations	1996	1996	1996	1996	1397

Notes: The table presents coefficient estimates from linear regressions. *Fail comprehension* equals 1 if a Social Architect failed one or more comprehension questions on the first try and 0 otherwise. *Inconsistent (sign)* equals 1 if, in the identical third and sixth comparisons, a Social Architect assigns progressive welfare weights ($\tilde{g} < 1$) in one and regressive welfare weights ($\tilde{g} > 1$) in the other. *Extreme time* equals 1 if a Social Architect's time spent on the survey lies beyond two standard deviations of the mean. *Low confidence* equals 1 if a Social Architect reports confidence levels lower than the highest category of "Very Much." *Stable* equals 1 if a Social Architect's welfare weights across the two waves retain the same type in classification (progressive or regressive) across waves. The explanatory variable is a treatment indicator. The regressions in Columns (1) through (4) use data from Wave 1, while the regression in Column (5) uses data from both waves. HC3 standard errors in parentheses.
*p<0.1, **p<0.05, ***p<0.01

E Additional Analysis

E.1 Assessing the Constant-Elasticity Specification

In this section, we report two diagnostics of the constant-elasticity specification $g(c) \propto c^{\nu}$. The first compares it to alternative one-parameter specifications. The second examines whether the estimated progressivity parameter is sensitive to any single Recipient income.

Alternative specifications. We compare the constant-elasticity specification to two alternative one-parameter specifications: a linear specification and a constant semi-elasticity specification. The linear specification is given by

$$g(c) \propto 1 + \beta c, \tag{A1}$$

and the constant semi-elasticity specification is given by

$$g(c) \propto \exp(\delta c). \tag{A2}$$

Under additively separable utility with constant absolute risk aversion (CARA) in consumption, $-\delta$ is the coefficient of absolute risk aversion. Because welfare weights are identified only up to a multiplicative constant, only the slope parameter (ν, β, or δ) is identified in each specification. We estimate each specification separately for each Social Architect.

To compare fit on a common scale, we compute, for each Social Architect, the in-sample root mean squared error (RMSE) between predicted and observed welfare weights, after normalizing each predicted schedule to equal one at the common Recipient's income.[1] Table A8 reports the median and mean RMSE across Social Architects. Across both common-Recipient treatments, the constant-elasticity specification yields the lowest median and mean RMSE, indicating that it provides a better summary of how welfare weights vary with income.

Leave-one-out estimates. We next examine whether the estimated progressivity parameter is driven by any particular Recipient income. For each Social Architect, we re-estimate ν_i after dropping one Recipient income at a time, and then recompute the median estimate

[1]For each Social Architect i, $RMSE_i = \sqrt{R^{-1} \sum_r (\hat{g}_{ir} - g_{ir})^2}$, where $\hat{g}_{ir}$ is the welfare weight predicted by the fitted schedule at Recipient income c_r, g_{ir} is the weight implied by Social Architect i's switching behavior at c_r, and R (=6) is the number of Recipients.

Table A8: Fit of One-Parameter Welfare-Weight Schedules

	Common Recipient: $60,000		Common Recipient: $120,000	
	Median	Mean	Median	Mean
Constant elasticity (ν)	1.087	7.034	4.870	11.547
Linear (β)	1.785	7.460	7.786	12.219
Constant semi-elasticity (δ)	1.448	7.855	5.828	14.674

Notes: This table reports the median and mean in-sample root mean squared error (RMSE) of predicted welfare weights across Social Architects for three specifications. For each Social Architect, each specification is estimated using the welfare weights assigned to the six Recipients. The welfare weights schedule and the predicted schedule are normalized to equal one at the common Recipient's income before computing the RMSE. The specifications use data from Wave 1.

across Social Architects. If a single constant-elasticity parameter summarizes the schedule well, the median estimate should remain similar when any one income is excluded.

Table A9 presents the leave-one-out medians. In the $60,000 treatment, the medians range from −0.83 to −0.74, close to the full-sample unweighted median of −0.79 observed in this treatment. In the $120,000 treatment, the medians range more widely, from −0.79 to −0.62, around a full-sample median of −0.65. The greater sensitivity in the $120,000 treatment is consistent with the weaker constant-elasticity fit at low incomes shown in Figure 3.

Table A9: Leave-One-Out Estimates of ν_i

Excluded Recipient	Common Recipient	
	$60,000	$120,000
$15,000	−0.83	−0.79
$30,000	−0.74	−0.63
$60,000	−0.78	−0.73
$120,000	−0.74	−0.63
$240,000	−0.76	−0.63
$480,000	−0.79	−0.62
None	−0.79	−0.65

Notes: This table presents the median estimates of progressivity ν after re-estimating each ν_i on the subsample that excludes one Recipient income at a time. For example, the entry −0.83 in the first row is the median ν in the $60,000-common-Recipient treatment when the $15,000 Recipient is excluded from each Social Architect's welfare-weight schedule. The bottom row reports the full-sample median estimate for reference. The specifications use data from Wave 1.

E.2 Welfare Weights and Proximity to Own Income

This section explores whether Social Architects assign higher welfare weights to Recipients with incomes similar to their own. Because Recipients' incomes are disposable incomes whereas Social Architects' measured incomes are pre-tax, we convert Social Architects' pre-tax incomes to disposable incomes using the NBER TAXSIM35 model (Feenberg and Coutts, 1993). We define disposable income as pre-tax wages net of federal and state income taxes and the employee share of payroll taxes, computed under 2023 federal and state tax rules (the most recent year available). For this calculation we use each Social Architect's state of residence and age, and assume each is a single filer with no dependents.

We estimate linear regressions of the following form

$$g_{ij} = \alpha_i + \beta_j \, \text{Near}_{ij} + \varepsilon_{ij}, \tag{A3}$$

where i indexes Social Architects, $j \in \{1, \ldots, 6\}$ indexes Recipients, and g_{ij} is the weight Social Architect i assigns to Recipient j. Near_{ij} is an indicator equal to one if Social Architect i's disposable income is close to Recipient j's income (defined two ways below). The specification includes Social Architect fixed effects α_i, so the coefficients identify within-Architect variation in the weight assigned to Recipients.

We use two definitions of proximity, which differ in how close a Social Architect's income must be to a Recipient's to count as "near." The first matches each Social Architect to the single closest Recipient. We partition the income line into six brackets, bounded by the midpoints between adjacent Recipient incomes, and set $\text{Near}_{ij} = 1$ if Social Architect i's income falls in the bracket for which Recipient j is the closest Recipient. These brackets are [\$0, \$22,500) for the \$15,000 Recipient, [\$22,500, \$45,000) for \$30,000, [\$45,000, \$90,000) for \$60,000, [\$90,000, \$180,000) for \$120,000, [\$180,000, \$360,000) for \$240,000, and [\$360,000, ∞) for \$480,000. Because the brackets partition the income line, every Social Architect falls in exactly one bracket and is matched to exactly one Recipient. The second uses a stricter notion of closeness: $\text{Near}_{ij} = 1$ only if Social Architect i's income is within ±20% of Recipient j's income (e.g., [\$12,000, \$18,000] for the \$15,000 Recipient). Because Recipient incomes roughly double across the distribution, these ±20% bands are disjoint and do not cover the income line—incomes between \$18,000 and \$24,000, for instance, fall in no band. A Social Architect is therefore near at most one Recipient, and those whose income lands in a gap are near none.

Table A10 reports the results, which are broadly similar across the two definitions of proximity. The tendency to weight Recipients near one's own income more heavily is concentrated at the bottom of the income distribution: Social Architects whose income is closest

Table A10: Welfare Weights and Proximity to Own Income

	(1)	(2)
Near Recipient earning 15K	0.195***	0.176***
	(0.012)	(0.022)
Near Recipient earning 30K	0.117***	0.111***
	(0.010)	(0.013)
Near Recipient earning 60K	-0.002	0.003
	(0.007)	(0.008)
Near Recipient earning 120K	-0.041***	-0.034*
	(0.015)	(0.018)
Near Recipient earning 240K	-0.074*	-0.042
	(0.041)	(0.054)
Near Recipient earning 480K	0.161	0.342
	(0.227)	(0.259)
Constant	0.154***	0.162***
	(0.001)	(0.001)
Observations	11976	11976

Notes: The table presents coefficient estimates from linear regressions of Equation (A3). The dependent variable g_{ij} is the weight assigned by Social Architect i to Recipient j. The indicators *Near Recipient earning 15K* through *Near Recipient earning 480K* equal 1 if Social Architect i's income is closest to that Recipient's income (Column 1) or within ±20% of it (Column 2), as defined in the text. The regressions use data from Wave 1 and include Social Architect fixed effects. Robust standard errors (HC1) in parentheses.
$^{*}p < 0.10$, $^{**}p < 0.05$, $^{***}p < 0.01$.

to the \$15,000 or \$30,000 Recipient assign that Recipient a higher weight, with coefficients that are positive and significant at the 1% level under both definitions (0.19 and 0.12 for the bracket definition; 0.18 and 0.11 for the ±20% definition). The pattern does not carry over to higher incomes: for the \$60,000, \$120,000, and \$240,000 Recipients, the coefficients are near zero or negative in both columns. For the \$480,000 Recipient, the coefficient is positive but very imprecisely estimated (0.161 and 0.342, both insignificant), reflecting how few Social Architects have incomes close to \$480,000. Overall, the tendency to assign higher weights to Recipients with similar incomes is robust across the two definitions but confined to the lower end of the income distribution.

E.3 Correcting for Anchoring

We find that the progressivity of the welfare weights is affected by the first decisions of the staircase method, consistent with anchoring to the first decision. To account for this, we adjust our progressivity estimates following Luttmer and Samwick (2018) and Bursztyn et al. (2025).

Let r_c^r denote the reported reform amount accruing to the low-income Recipient at the point of indifference in comparison c, for $c \in \{1, \ldots, 5\}$. We assume r_c^r is a weighted average of the starting value r^s (either \$300 or \$700, depending on the treatment) and the actual value r_c^a:

$$r_c^r \;=\; \beta_c \, r^s \;+\; (1-\beta_c)\, r_c^a. \tag{A4}$$

By rearranging the terms, we obtain the actual value of the reform:

$$r_c^a \;=\; r_c^r \;-\; \frac{\beta_c}{1-\beta_c}\big(r^s - r_c^r\big). \tag{A5}$$

We estimate β_c by regressing the reported value r_c^r on the starting value r^s, including controls for response quality (described in Section 4.4) and demographics (described in Section 4.3). Table A11 reports the regression results. We find only moderate effects of anchoring, with values of β_c ranging from 0.06 to 0.11. The coefficients in only two of the five regressions are statistically significant at the 5% level.

Using the estimates β_c for each comparison $c \in \{1, \ldots, 5\}$, we compute the actual reform amounts using Equation A5. We truncate these reform amounts at \$25 and \$975—the bounds of the unadjusted reform amounts. We then estimate welfare weights with these actual reform amounts using the procedures described in the main text.

Table A11: Adjustment of Starting Point

	r_1^r (1)	r_2^r (2)	r_3^r (3)	r_4^r (4)	r_5^r (5)
Starting value	0.071**	0.115***	0.066*	0.068*	0.062*
	(0.034)	(0.034)	(0.034)	(0.036)	(0.036)
Observations	1996	1996	1996	1996	1996
Controls?	Yes	Yes	Yes	Yes	Yes

Notes: The table reports coefficient estimates from linear regressions. The dependent variables in Columns (1)–(5) are the reported reform amounts accruing to the low-income Recipient at the point of indifference in each of the five comparisons: $r_c^r \; \forall \, c \in \{1, \ldots, 5\}$. *Starting Value* equals 300 if the first decision is randomized to be (\$300, −\$700) and 700 if the first decision is (\$700, −\$300). Controls include response quality indicators (those used in Table 5) and demographic indicators (those used in Table 4). The regressions use data from Wave 1. HC3 standard errors in parentheses.
*p<0.1, **p<0.05, ***p<0.01

E.4 What Do Welfare Weights Capture?

Section 4.7 showed that welfare weights predict support for government redistribution. Here we ask which underlying concerns account for this relationship and, in particular, whether fairness concerns—which welfare weights are intended to capture—are among them. We elicit Social Architects' fairness concerns by asking them to rate the fairness of the current disposable incomes on a five-point scale from "Very unfair" to "Very fair." This measure likely captures several fairness ideals, including redistribution based on the source of income, but it does not capture some important ones, such as equality of opportunity. We also elicit other concerns—(i) beliefs that higher taxes hurt the economy, (ii) trickle-down economics, (iii) beliefs about labor distortions from taxing high-income earners, (iv) trust in government, and (v) the scope of government—and misperceptions about (vi) the level of taxes individuals pay and (vii) the share of individuals with incomes below \$15,000. Definitions of the variables can be found in Appendix Section A. Because fairness is multifaceted and difficult to measure with a single survey item, we view this exercise as a conservative test, as it is likely to understate the role of fairness concerns.

We proceed in two steps. First, we estimate the overall variation in support for redistribution at the margin that can be explained by welfare weights using a linear regression. Some of this overall variation may be explained by various concerns, which "mediate" the effect of welfare weights. In the second step, we decompose this overall variation in support for redistribution explained by welfare weights into the variation explained by each concern using a Gelbach decomposition (Gelbach, 2016). A variable is said to be captured by welfare weights if it is both correlated with welfare weights and predictive of support for redistribution. The decomposition nets out the overlap among the variables and does not depend on the order in which they are added.

Table A12 presents results from a linear regression in which the dependent variable is an indicator equal to 1 if a Social Architect supports additional progressive redistribution at the margin (a positive value on the -2 to $+2$ scale) and 0 otherwise. In Column (1), the only explanatory variable is $1(\nu < 0)$, an indicator equal to 1 if a Social Architect has progressive welfare weights and 0 otherwise. Its coefficient is 0.44: Social Architects with progressive weights are 44 percentage points more likely to support additional progressive redistribution, which we take as the overall relationship between welfare weights and support for redistribution. In Column (2), we add the measured beliefs and misperceptions. The coefficient on $1(\nu < 0)$ falls to 0.32, implying that about 27% of this relationship runs through these other variables.

Table A13 reports the decomposition, following Gelbach (2016). Each value is a factor's contribution to the change in the coefficient on welfare weights as we move from the base

Table A12: Support for Redistribution and Mechanisms

	Redistribute at margin	
	(1)	(2)
$1(\nu < 0)$	0.443***	0.321***
	(0.021)	(0.021)
Fair distribution of income		-0.112***
		(0.024)
Behavioral responses high-earners		-0.065***
		(0.019)
High taxes high-incomes hurt economy		-0.158***
		(0.022)
No belief trickle down		-0.037
		(0.023)
Low trust in government		-0.063***
		(0.021)
Govt should do less		-0.215***
		(0.019)
Overestimate taxes		-0.001
		(0.009)
Overestimate low-income share		0.007
		(0.008)
Constant	0.399***	0.729***
	(0.019)	(0.034)
Observations	1996	1996

Notes: The table presents coefficient estimates from linear regressions. Both columns use support for redistribution at the margin—support for redistribution beyond that achieved by the current tax and transfer system; this ranges from -2 to $+2$, where positive (negative) values indicate a desire for redistribution from high-income (low/middle-income) to low/middle-income (high-income) individuals, and zero indicates no desired additional redistribution. The dependent variable takes a value of 1 if a Social Architect supports additional redistribution at the margin (values above 0) and 0 otherwise. $1(\nu < 0)$ is an indicator variable take a value of 1 for progressive weights and 0 otherwise. Definitions of the other explanatory variables can be found in Appendix Section A. HC3 standard errors in parentheses.
*p<0.1, **p<0.05, ***p<0.01

specification in Column (1) of Table A12 to the full specification in Column (2), divided by the Column (1) coefficient and expressed in percent; intuitively, it is the share of the overall welfare-weight–redistribution relationship that runs through each factor. These factor shares and the residual *Share unexplained* sum to 100%.

Consistent with our intended interpretation, welfare weights do capture fairness concerns, though the share is modest: fairness accounts for about 3% of the relationship. They also capture other concerns—the belief that higher taxes hurt the economy accounts for about 9%, and the broader view that the government should do less for about 15%. Other mechanisms and the two misperceptions account for a very small share of the relationship. The large majority of the relationship—about 73%—is left unexplained, indicating that the policy-relevant variation in welfare weights reflects considerations beyond the beliefs and misperceptions we are able to measure.

Table A13: Decomposition of Treatment Effects

	(1)
Fair distribution of income	3.08
Behavioral responses high-earners	0.78
High taxes high-incomes hurt economy	8.58
No belief trickle down	–0.20
Low trust in government	0.09
Govt should do less	14.86
Overestimate taxes	–0.01
Overestimate low-income share	0.24
Share unexplained	72.57

Notes: The table presents the Gelbach decomposition following Gelbach (2016). Each value reports the share (in percent) of the coefficient on $1(\nu < 0)$ in Column (1) of Table A12 that is explained by each factor. *Share unexplained* is the share of the coefficient that remains after accounting for all factors. Definitions of all variables can be found in Appendix Section A.

E.5 Comparison to Inverse-Optimum Weights

Inverse-Optimum Weights Implied by Tax Schedule: We obtain the inverse optimum welfare weights implied by the U.S. income tax schedule from Hendren (2020). Hendren (2020) computes MTRs using the universe of tax returns in 2012, incorporating ordinary income taxes, the alternative minimum tax (AMT), the earned income tax credits (EITC), and state, local, and Medicare taxes. The baseline compensated elasticities of taxable income (ETI) are 0.30 for top earners, 0.31 for earners in the EITC phase-in region, 0.14 for those in the phase-out region, and 0.30 for everyone else. We exclude negative incomes from our analysis. Figure A11 plots the inverse-optimum weights implied by the tax schedule against percentiles of the income distribution.

Inverse-Optimum Weights Implied by Tax and Transfer Policies: The inverse optimum weights implied by U.S. tax and transfer policies are derived using the framework outlined in Hendren and Sprung-Keyser (2020). In this framework, the welfare weight $g(z)$ assigned to beneficiaries of a policy earning z is the inverse of the policy's Marginal Value of Public Funds (MVPF)—the beneficiaries' welfare gain from \$1 of government spending. The MVPF of a policy affecting beneficiaries with incomes z^* is defined as the beneficiaries' willingness to pay for the policy (s^*) divided by the net cost (c) accrued from the policy to the government. To replicate this policy (with benefits s^*) through adjustments to the tax schedule, the policy would cost $s^*g(z^*)$, where $g(z^*)$ is the welfare weight. It would be cheaper to achieve s^* through the policy than through adjustments to the tax schedule if and only if $c \leq s^*g(z^*)$. Rewriting this expression yields the following equation: $MVPF = s^*/c \geq 1/g(z^*)$. We focus on tax and in-cash transfer policies, shown in Figure 6 of Hendren and Sprung-Keyser (2020), obtaining each policy's MVPF and the approximate income of its beneficiaries from their data. From this set, we drop policies with an MVPF of infinity or a negative MVPF, as our theoretical framework and experimental design cannot accommodate such values. Figure A12 plots the inverse-optimum weights implied by the tax and transfer policies against the approximate income of their beneficiaries.

Estimating and Plotting the Welfare Weights. We summarize the progressivity of each set of inverse-optimum weights with the constant-elasticity specification $g(c) \propto c^{\nu}$, estimated by PPML as in the main text. Figure 7 plots these fitted schedules together with the weights obtained from the experiment against the disposable income distribution, interpolating each with c^{ν} at its estimated ν and normalizing it to sum to one.

E.6 Calibration of Optimal Income Taxes

We explore the implications of the welfare weights elicited in the experiment for the optimal non-linear labor income taxes in the U.S. We use the optimal tax formula derived in Saez (2001), and follow the simulations presented in Støstad and Cowell (2024) and Mankiw et al. (2009).

The optimal tax formula solves a social planner's problem of maximizing social welfare subject to constraints. It expresses optimal marginal tax rates (MTRs) as a function of (i) the shape of the ability (wage) distribution, (ii) the elasticity of taxable income, and (iii) welfare weights. We estimate the ability distribution directly from the observed U.S. income distribution, apply elasticity estimates from existing literature, and incorporate welfare weights estimated in our paper.

Optimal Tax Formula

Individuals lie on a continuum of abilities (wages) w with density $f(w)$ and cumulative distribution function $F(w)$. An individual earns income $z = wl$, where l is labor supply, faces a tax schedule $T(z)$, and consumes $c = z - T(z) + G$, where G is a uniform lump-sum transfer (demogrant) financed by tax revenue. The planner observes only income z, not labor supply l or ability w, and so conditions taxes on income alone. Individual preferences are

$$U(c, l) = c - \frac{l^{1+\frac{1}{E_L}}}{1 + \frac{1}{E_L}}, \tag{A6}$$

which are quasi-linear in consumption and isoelastic in labor, with E_L denoting the elasticity of earnings with respect to the retention rate $1 - T'(z)$. Quasi-linearity implies that labor supply has no income effects, so the uncompensated and compensated elasticities coincide.

The planner's preference for redistribution is summarized by social marginal welfare weights $g(c) = c^{\nu}$, with $\nu < 0$, which decline in consumption and which we estimate in our experiment. Following Saez (2001), the planner's first-order condition is

$$\frac{T'(z(w))}{1 - T'(z(w))} = \left(\frac{1 + E_L(w)}{E_L(w)}\right) \frac{g(c(w))}{wf(w)} \left[\int_w^{\infty} \frac{f(\theta)}{g(c(\theta))}\, d\theta - \big(1 - F(w)\big)\frac{1}{p}\right], \tag{A7}$$

where $T'(z(w))$ is the optimal MTR at income $z(w)$ and θ indexes ability in the integral. The welfare weight $g(c(w)) = c(w)^{\nu}$ is the social value of a marginal dollar of consumption to an individual with wage w, so more progressive weights (more negative ν) imply more progressive optimal MTRs.

The term p is the marginal value of public funds, which measures the increase in social welfare from relaxing the planner's budget constraint by one unit. Because $g(c(w))$ is the social value of a dollar to an individual with wage w, its reciprocal $1/g(c(w))$ is the consumption cost of delivering one unit of social welfare to that individual, and $\int_0^{\infty} \frac{1}{g(c(w))} f(w)\, dw$ is the average cost across individuals. The value to the planner of a marginal unit of public funds is the inverse of this cost, which is given by

$$p = \left[\int_0^{\infty} \frac{f(w)}{g(c(w))}\, dw\right]^{-1}. \tag{A8}$$

Estimating the Wage-Ability Distribution

We estimate the ability (wage) distribution from the 2019 labor income distribution and use this exogenous ability distribution when we calibrate the tax formula in Equation (A7). Data on the labor income distribution is obtained from the Distributional National Accounts (DINA) micro-files of Piketty et al. (2018). Each observation in the data corresponds to a tax unit. Our analysis proceeds in several steps.

Step 1: Use the NBER TAXSIM model to find the federal marginal tax rate for each tax unit, based on the number of dependents, the age of the primary filer, and marital status. Add a 5% state tax rate, a 2.9% Medicare tax rate, and a 2.3% sales tax rate to obtain the total marginal rate.

Step 2: Assuming individuals optimize according to their utility in Equation (A6), invert the labor-supply condition to back out each tax unit's ability w:

$$w = \frac{z^{1/(1+E_L)}}{(1 - T'(z))^{E_L/(1+E_L)}},$$

evaluated at each unit's labor income z and marginal rate $T'(z)$. We set $E_L = 0.25$, a mid-range estimate of the elasticity of taxable income (Saez et al., 2012).

Step 3: Fit a smooth ability density with a Gaussian kernel density estimator (bandwidth $5,000), weighted by DINA sample weights, on a grid of 50,000 ability points. Because the kernel places some mass below zero, we truncate the density at zero and renormalize it to integrate to one over non-negative abilities.

Step 4: Replace the top 0.5% of the ability distribution with a Pareto tail, setting the Pareto parameter equal to the local value of the ability distribution, $\alpha(w) = \frac{w\, f(w)}{1-F(w)}$, evaluated just below the top 0.5%.

Calibrating Optimal Income Taxes

We solve for the optimal tax schedule by iterating to a fixed point, drawing from Mankiw et al. (2009) and Støstad and Cowell (2024). Starting from an initial schedule, we compute individuals' labor supply, incomes, and consumption; from consumption we obtain the welfare weights and the marginal value of public funds, and from these the marginal tax rates implied by Equation (A7). We iterate until the schedule converges, and at the converged optimum we verify the second-order condition that pre-tax income is non-decreasing in ability.

Step 1: Start with an initial flat tax rate of 35%.

Step 2: Compute labor choices under this schedule. Maximizing the utility in Equation (A6) with respect to l yields $l = (w\,(1 - T'(z)))^{E_L}$. We set $E_L = 0.25$, the mid-range estimate used above.

Step 3: Compute incomes $z = wl$ and consumption $c = z - T(z) + G$, where G is the lump-sum transfer (demogrant). Evaluate the welfare weights $g(c) = c^{\nu}$ at $\nu = -0.71$ and $\nu = -0.85$, the bounds obtained from our experiment, and compute the marginal value of public funds p from Equation (A8). A non-working group at the bottom, which consumes only the transfer G, is included when computing p.

Step 4: Compute the marginal tax schedule implied by Equation (A7).[2]

Step 5: Repeat the previous steps until the tax rates converge to a fixed point.

Results

The less progressive endpoint of the bounds ($\nu = -0.71$) yields an average optimal MTR of 56%, while the more progressive endpoint ($\nu = -0.85$) yields 60%. Thus, the bounds of progressivity estimates produce relatively tight bounds on optimal marginal tax rates.

[2]We update the schedule as the pointwise average of the new and previous schedules, which stabilizes the iteration.

F Additional Studies

We conducted two additional studies—Study 2 and Study 3—that preceded and informed the design of Study 1. Each tested a specific feature of the welfare-weight elicitation. Study 2 tests whether the elicited weights are sensitive to the gain-loss framing of the reforms; Study 3 tests whether they are sensitive to the presence of real stakes. We describe each study in turn and then summarize how their designs differ from Study 1.

F.1 Study 2: Gain-Loss Framing

Design. In Study 2, we tested whether Social Architects' welfare weights are sensitive to the framing of the reforms. Social Architects are randomly assigned to one of four treatments in a 2×2 design. The first dimension varies the framing. In Treatments Loss, reforms take money away from the higher-income Recipient and give money to the lower-income Recipient: each Recipient receives an initial endowment of \$1,500, and a Social Architect decides between the reforms (r_ℓ, $-r_h$) and (\$500, −\$500). In Treatments Gain, reforms instead give money to both Recipients: the endowment is incorporated into the reform amounts, so a Social Architect decides between (\$1,500+r_ℓ, \$1,500−r_h) and (\$2,000, \$1,000). The two framings induce identical final consumption for the Recipients, such that a Social Architect conditioning weights on consumption should assign the same weights across the two frames. The second dimension varies the income of the Recipient common across the comparisons: \$70,000 in Treatments 70K and \$500,000 in Treatments 500K. The complete set of instructions for Study 2 is in Appendix Sections H.

Data Collection and Sample. We recruited participants in the role of Social Architects from the U.S. general population using the data collection provider Lucid. We defined recruitment quotas based on gender, age, education, individual income, and region, designed to match the sample to the U.S. population. Participants had to pass one attention check and two comprehension checks to continue with the study. We implemented the survey using Qualtrics. Data collection began on 8 December 2021 and lasted approximately two weeks. Our final sample includes 1,965 participants.

Results. Column (1) of Table A14 presents estimates from a median regression of the progressivity of the welfare weights (ν_i) on treatment indicators, with Treatment Loss $\times$ 70K as the omitted category. We find mixed evidence of the role framing. Welfare weights are more progressive in Treatment Gain $\times$ 70K than in Treatment Loss $\times$ 70K (ν is 0.12 lower), and this difference is statistically significant. However, the median weights are identical across framings in Treatments 500K.

Table A14: Treatment Effects in Studies 2 and 3

	(1)	(2)
Gain x 70K	-0.125***	
	(0.038)	
Loss x 500K	0.203***	
	(0.038)	
Gain x 500K	0.203***	
	(0.038)	
Hypothetical		-0.307***
		(0.096)
Constant	-0.203***	-0.185***
	(0.027)	(0.068)
Observations	1965	997
Study:	Study 2	Study 3

Notes: The table presents coefficient estimates from a median regression.The regression in Column (1) uses data from Study 2 while the regression in Column (2) uses data from Study 3. The dependent variable is the progressivity of the welfare weights (ν_i). The explanatory variables include treatment indicators. The omitted category in Column (1) is Treatment Loss $\times$ 70K while the omitted category in Column (2) is Treatment Real. Standard errors in parentheses.
$^{*}p < 0.10,^{**} p < 0.05,^{***} p < 0.01.$

F.2 Study 3: Real vs. Hypothetical Stakes

Design. In Study 3, we test whether Social Architects' welfare weights are sensitive to the presence of real stakes. The elicitation procedure closely follows that of Study 2. Social Architects are randomly assigned to one of four treatments. The first two treatments vary the stakes: in Treatment Real, Social Architects make real decisions regarding real Recipients, whereas in Treatment Hypothetical, they make hypothetical decisions regarding hypothetical Recipients. The remaining two treatments examine the role of self-interest motives; which we do not discuss here. The complete set of instructions for Study 3 is in Appendix Sections I.

Data Collection and Sample. We recruited participants in the role of Social Architects from the data collection provider Prolific. Participants were recruited based on three quotas available on Prolific: age, sex, and ethnicity. We implemented the survey using oTree (Chen et al., 2016). Data collection began on 14 December 2022 and lasted eight days. Our final sample includes 1,992 participants.

Results. Column (2) of Table A14 presents estimates from a median regression of the progressivity of the welfare weights (ν_i) on an indicator equal to one for Treatment Hypothetical and zero for Treatment Real. Welfare weights are more progressive under hypothetical stakes (ν is 0.30 lower) than under real stakes. This pattern is consistent with Social Architects weighing the trade-offs more carefully when their choices carry real consequences.

F.3 Differences from Study 1

Studies 2 and 3 differ from Study 1 in three respects. First, the decision structure differs: in Study 1, Recipients have an endowment of \$1,000 and Social Architects choose between implementing a reform (r_ℓ, $-r_h$) and maintaining the status quo (0, 0); in the additional studies, Recipients' have an endowment of \$1,500 and Social Architects choose between a fixed reform (\$500, $-$\$500) and a varying reform (r_ℓ, $-r_h$). The decisions in Study 1 are simpler, as they require Social Architects to compare reforms against a status quo rather than comparing two distinct reforms. However, the decisions in Study 2 are better suited to testing the role of gain-loss framing, since the decision structure—and consequently their complexity—is held constant across frames. In Study 1, a gain-loss frame (comparing a reform to a status quo) is more complicated than a gain-gain frame (comparing two reforms), confounding the effect of framing with complexity.

Second, the income grid differs: Study 1 uses six Recipients with disposable incomes of \$15,000, \$30,000, \$60,000, \$120,000, \$240,000, and \$480,000, whereas the additional studies use seven Recipients with disposable incomes of \$8,000, \$35,000, \$70,000, \$100,000, \$170,000, \$250,000, and \$500,000. In Study 1, each Recipient's income is exactly double the previous one, so that the ratios of neighboring incomes are constant; this regular spacing is designed to reduce the cognitive burden of the task.

Third, Study 1 includes several additional robustness checks—including a consistency check, randomization of the starting decision, and prompts to verify the consequences of choices—that are absent from the additional studies.

G Instructions - Study 1

G.1 Wave 1

Bold text, underlining, tables, etc. appear as in the original screen.

Screen Break

[Consent screen]

Introduction
Welcome! This study is conducted by Unidistance Suisse, Switzerland and WZB, Germany. Our goal is to understand the views of U.S. residents on various topics. By carefully completing this survey, you are helping us to understand these views.

Requirements
To participate in this study, you must be a U.S. resident and at least 18 years old.

Time required
This study will take around **12 minutes**.

Compensation
You will receive **$2.5** for completing the study. The payment will be made in the next few days.

Checks
Our survey includes attention checks to test whether participants take our survey carefully. Additionally, we have implemented measures to ensure that participants do not use AI assistance during the surveys. Participants must complete the survey independently, without the help of AI tools. Participants who fail these checks cannot proceed with the survey, and will be asked to return the survey.

Follow-up study
You may be contacted for a follow-up study approximately four weeks from now. We will notify you of the study via Prolific. Your participation in the follow-up is very important to us.

Confidentiality

Your answers will remain anonymous and will be used for scientific purposes only. Strict confidentiality is guaranteed, and your identity can never be associated with your answers.

Voluntary participation
Participation in this study is voluntary. You may withdraw from the study at any time.

Questions about the survey
If you have questions about this study or your rights, please get in touch with us at krishna.srinivasan@unidistance.ch

Consent
I have received the above information and I am willing to participate in the study.

[Yes; No]

What is your Prolific ID?

——————————— Screen Break ———————————

[Screen shown if participant does not provide consent]

End of Survey

You did not give your consent to continue with the study.

Please close this survey and return your submission on Prolific by selecting the "Stop without completing" button.

——————————— Screen Break ———————————

[**Block 1: Background Questions**]

What is your gender?

[Female; Male; Non-binary; Prefer not to say]

How old are you?

[18 years old - 24 years old; 25 years old - 34 years old; 35 years old - 44 years old; 45 years old - 54 years old; 55 years old - 64 years old; 65 years old or above]

In which state do you currently reside?

[Alabama; ...; Wyoming; I do not reside in the U.S.]

What is the highest level of education you have completed?

[Primary education or less; Some high school; High school degree/GED; Some college; 2-year college degree; 4-year college degree; Master's degree; Doctoral degree; Professional degree (e.g., JD, MD, MBA)]

The next question is about your **total individual income**, before taxes, last year. This figure should include income from all sources, including salaries, wages, pensions, social security, dividends, interest, and all other income.

What was your total individual income, before taxes, last year?

[$29,999 and below; $30,000 to $59,999; $60,000 to $99,999; $100,000 to $149,999; $150,000 and above]

Screen Break

[Displayed if $29,999 and below is chosen]

You have reported that your total individual income, before taxes, last year was $29,999 and below.

[Displayed if $30,000 to $59,999 is chosen]

You have reported that your total individual income, before taxes, last year was $30,000 to $59,999.

[Displayed if $60,000 to $99,999 is chosen]

You have reported that your total individual income, before taxes, last year was $60,000 to $99,999.

[Displayed if $100,000 to $149,999 is chosen]

You have reported that your total individual income, before taxes, last year was $100,000 to $149,999.

[Displayed if $150,000 and above is chosen]

You have reported that your total individual income, before taxes, last year was $150,000 and above.

[Displayed in all cases]

Could you provide your best guess of what your total individual income, before taxes, last year was?

——————————— Screen Break ———————————

In politics, as of today, do you consider yourself a Republican, a Democrat or an independent?

[Republican; Democrat; Independent]

——————————— Screen Break ———————————

[Screen shown if participant does not reside in the U.S.]

End of survey

Unfortunately, you do not fulfill the requirements of this study since you do not reside in the U.S.

Thank you for your time.

Please close this survey and return your submission on Prolific by selecting the "Stop without completing" button.

——————————— Screen Break ———————————

[Attention check]

In surveys like ours, some participants do not carefully read the questions. This means that there are a lot of random answers that can compromise the results of research studies. To show that you read our questions carefully, please choose "Not at all interested" below:

[Extremely interested; Very interested; A little bit interested; Almost not interested; Not at all interested]

Screen Break

[Screen shown if participant failed the attention check]

End of survey

Unfortunately, you failed the attention check.

For this reason, you cannot continue the study and will not receive a payment.

Please close this survey and return your submission on Prolific by selecting the "Stop without completing" button.

Screen Break

[**Block 2: Eliciting Welfare Weights**]

Instructions

In this section, **you will be asked to decide whether you want to change the incomes of six real individuals in society.** These real individuals will be recruited from the U.S. general population. They are above the age of 18 and are U.S. citizens. As we will explain below, your decisions may have real consequences for two of these individuals.

You will be presented with several questions. In each question, you will be presented with two individuals and learn their disposable incomes. **Disposable income is defined as income after all taxes have been paid and transfers have been received (including federal and state taxes and transfers).**

In each question, you will be presented with a proposed change to the incomes of the individuals, and will be asked to indicate whether or not you prefer to implement this change. **If you prefer to implement the change, the income of the lower-income individual in the pair will increase, and the income of the higher-income individual in the pair will decrease by the amounts proposed in the change.** If you prefer not to implement the change, the incomes of the two individuals will remain unchanged. We will describe below how these changes are implemented.

Example

Here is an example of a question that you will see in the survey:

	Person #3	**Person #6**
Annual disposable income	\$60,000	\$120,000
Proposed change	\$500	-\$500

Please make your decision:

- ○ I prefer to implement the change
- ○ I prefer <u>not</u> to implement the change

In this example, there are two individuals: Person #3 with an income of \$60,000 and Person #6 with an income of \$120,000.

The proposed change involves increasing the income of Person #3 by $500 and decreasing the income of Person #6 by $500. If you prefer to implement the change, the final incomes of the individuals are Person #3: $60,500 and Person #6: $119,500. If you prefer not to implement the change, the incomes of the individuals remain unchanged.

Comparisons

You will face several questions like the one above, with the amounts in the proposed change varying across the questions. You will face several comparisons, with the income of the individuals varying across them. The following table indicates the comparisons.

Comparison	**Annual disposable income of individuals**
1	$15,000 vs. $60,000
2	$30,000 vs. $60,000
3	$60,000 vs. $120,000
4	$60,000 vs. $240,000
5	$60,000 vs. $480,000

Incentives and Implementation

At the end of the study, one participant will be randomly selected. If you are selected, one question will be randomly chosen from either this survey or the follow-up survey, and your choice on that question will be implemented. The two individuals involved in the selected question will then be affected by your decision. **Thus, if you are randomly selected, one of your choices may have real consequences for two other individuals.**

If you are selected, a $1,000 bonus will be transferred to the two individuals affected by your choice. The disposable incomes shown above for each person already include this $1,000 bonus, under the assumption that they receive it. The amounts specified in your proposed change will be added to or subtracted from that bonus.

[If Treatment Consequence Verification] For the first question in each comparison, you will be presented with two distributions of income for the pairs of individuals, which reflect the consequences of the two choices you were presented with. You will then be asked to decide which distribution you prefer. This is meant to help you consider the consequences of your choice before proceeding.

Please answer the following questions to show that you have understood the instructions. You can read the instructions above again if needed.

How many individuals will you make decisions regarding?

[Three; Six]

We will present you with the incomes of several individuals. What type of income will we present to you?

[Pre-tax income, Disposable income]

Please state True or False: "If you are randomly selected, one of your choices may have real consequences for two other individuals."

[True; False]

——————————— Screen Break ———————————

[*Note: If a participant fails any of the three comprehension checks, they are taken back to the instruction screen and informed which checks they failed. The error message reads: "In your first previous try, you answered [the first question/the second question/the third question/the first and second/the second and third/the first and third/all three] question[s] incorrectly. Please try again." Any check answered correctly is locked, while those failed require a new response. Participants must pass the remaining checks before proceeding. On subsequent attempts, if a participant fails a check, they receive the same error message.*]

——————————— Screen Break ———————————

[*We present the proposed change in each question as* ($\$r_\ell$, $-\$r_h$)*, where* $\$r_\ell$ *and* $-\$r_h$ *refer to the amounts accrued to the lower-income Recipient and high-income Recipients in the pair, respectively. Participants are randomized into one of two treatments that vary the amounts in the first question to be either (\$300, −\$700) or (\$700, −\$300).*]

[C1Q1: Comparison 1 Question]

Comparison 1/5, Question 1

[If first decision is ($300, −$700)]

Please consider each question carefully because if you are selected, one of your choices may have real consequences for two real individuals.

	Person #1	**Person #3**
Annual disposable income	$15,000	$60,000
Proposed change	$300	-$700

Please make your decision:

- ○ I prefer to implement the change
- ○ I prefer not to implement the change

If you prefer to implement the change, the final incomes of the individuals are Person #1: $15,300 and Person #3: $59,300. If you prefer not to implement the change, the incomes of individuals remain unchanged.

[If first decision is ($700, −$300)]

Please consider each question carefully because if you are selected, one of your choices may have real consequences for two real individuals.

	Person #1	**Person #3**
Annual disposable income	$15,000	$60,000
Proposed change	$700	-$300

Please make your decision:

- ○ I prefer to implement the change
- ○ I prefer not to implement the change

If you prefer to implement the change, the final incomes of the individuals are Person #1: $15,700 and Person #3: $59,700. If you prefer not to implement the change, the incomes of individuals remain unchanged.

Screen Break

[*Note: For brevity, the remaining questions in each comparison are presented in abbreviated form. Each bracketed line describes the question shown to participants given their path through the staircase.*]

[<u>If first decision is ($300, −$700)</u>]

[C1Q2.1: If change implemented in C1Q1, choose whether to implement ($200, −$800)]

[C1Q2.2: If change not implemented in C1Q1, choose whether to implement ($600, −$400)]

[<u>If first decision is ($700, −$300)</u>]

[C1Q2.1: If change implemented in C1Q1, choose whether to implement ($400, −$600)]

[C1Q2.2: If change not implemented in C1Q1, choose whether to implement ($800, −$200)]

Screen Break

[<u>If first decision is ($300, −$700)</u>]

[C1Q3.1: If change implemented in C1Q2.1, choose whether to implement ($100, −$900)]

[C1Q3.2: If change not implemented in C1Q2.1, choose whether to implement ($250, −$750)]

[C1Q3.3: If change implemented in C1Q2.2, choose whether to implement ($400, −$600)]

[C1Q3.4: If change not implemented in C1Q2.2, choose whether to implement ($800, −$200)]

[<u>If first decision is ($700, −$300)</u>]

[C1Q3.1: If change implemented in C1Q2.1, choose whether to implement ($200, −$800)]

[C1Q3.2: If change not implemented in C1Q2.1, choose whether to implement ($600, −$400)]

[C1Q3.3: If change implemented in C1Q2.2, choose whether to implement ($750, −$250)]

[C1Q3.4: If change not implemented in C1Q2.2, choose whether to implement ($900, −$100)]

Screen Break

[<u>If first decision is ($300, −$700)</u>]

[C1Q4.1: If change implemented in C1Q3.1, choose whether to implement ($50, −$950)]

[C1Q4.2: If change not implemented in C1Q3.1, choose whether to implement ($150, −$850)]

[If change implemented in C1Q3.2, participant indifferent between ($225, −$775)]

[If change not implemented in C1Q3.2, participant indifferent between ($275, −$725)]

[C1Q4.3: If change implemented in C1Q3.3, choose whether to implement ($350, −$650)]

[C1Q4.4: If change not implemented in C1Q3.3, choose whether to implement ($500, −$500)]

[C1Q4.5: If change implemented in C1Q3.4, choose whether to implement ($700, −$300)]

[C1Q4.6: If change not implemented in C1Q3.4, choose whether to implement ($900, −$100)]

[<u>If first decision is ($700, −$300)</u>]

[C1Q4.1: If change implemented in C1Q3.1, choose whether to implement ($100, −$900)]

[C1Q4.2: If change not implemented in C1Q3.1, choose whether to implement ($300, −$700)]

[C1Q4.3: If change implemented in C1Q3.2, choose whether to implement ($500, −$500)]

[C1Q4.4: If change not implemented in C1Q3.2, choose whether to implement ($650, −$350)]

[If change implemented in C1Q3.3, participant indifferent between ($725, −$275)]

[If change not implemented in C1Q3.3, participant indifferent between ($775, −$225)]

[C1Q4.5: If change implemented in C1Q3.4, choose whether to implement ($850, −$150)]

[C1Q4.6: If change not implemented in C1Q3.4, choose whether to implement ($950, −$50)]

Screen Break

[<u>If first decision is ($300, −$700)</u>]

[If change implemented in C1Q4.1, participant indifferent between ($25, −$975)]

[If change not implemented in C1Q4.1, participant indifferent between ($75, −$925)]

[If change implemented in C1Q4.2, participant indifferent between ($125, −$875)]

[If change not implemented in C1Q4.2, participant indifferent between ($175, −$825)]

[If change implemented in C1Q4.3, participant indifferent between ($325, −$675)]

[If change not implemented in C1Q4.3, participant indifferent between ($375, −$625)]

[C1Q5.1: If change implemented in C1Q4.4, choose whether to implement ($450, −$550)]

[C1Q5.2: If change not implemented in C1Q4.4, choose whether to implement ($550, −$450)]

[C1Q5.3: If change implemented in C1Q4.5, choose whether to implement ($650, −$350)]

[C1Q5.4: If change not implemented in C1Q4.5, choose whether to implement ($750, −$250)]

[C1Q5.5: If change implemented in C1Q4.6, choose whether to implement ($850, −$150)]

[C1Q5.6: If change not implemented in C1Q4.6, choose whether to implement ($950, −$50)]

[If first decision is ($700, −$300)]

[C1Q5.1: If change implemented in C1Q4.1, choose whether to implement ($50, −$950)]

[C1Q5.2: If change not implemented in C1Q4.1, choose whether to implement ($150, −$850)]

[C1Q5.3: If change implemented in C1Q4.2, choose whether to implement ($250, −$750)]

[C1Q5.4: If change not implemented in C1Q4.2, choose whether to implement ($350, −$650)]

[C1Q5.5: If change implemented in C1Q4.3, choose whether to implement ($450, −$550)]

[C1Q5.6: If change not implemented in C1Q4.3, choose whether to implement ($550, −$450)]

[If change implemented in C1Q4.4, participant indifferent between ($625, −$375)]

[If change not implemented in C1Q4.4, participant indifferent between ($675, −$325)]

[If change implemented in C1Q4.5, participant indifferent between ($825, −$175)]

[If change not implemented in C1Q4.5, participant indifferent between ($875, −$125)]

[If change implemented in C1Q4.6, participant indifferent between ($925, −$75)]

[If change not implemented in C1Q4.6, participant indifferent between ($975, −$25)]

Screen Break

[If first decision is ($300, −$700)]

[If change implemented in C1Q5.1, participant indifferent between ($425, −$575)]

[If change not implemented in C1Q5.1, participant indifferent between ($475, −$525)]

[If change implemented in C1Q5.2, participant indifferent between ($525, −$475)]

[If change not implemented in C1Q5.2, participant indifferent between ($575, −$425)]

[If change implemented in C1Q5.3, participant indifferent between ($625, −$375)]

[If change not implemented in C1Q5.3, participant indifferent between ($675, −$325)]

[If change implemented in C1Q5.4, participant indifferent between ($725, −$275)]

[If change not implemented in C1Q5.4, participant indifferent between ($775, −$225)]

[If change implemented in C1Q5.5, participant indifferent between ($825, −$175)]

[If change not implemented in C1Q5.5, participant indifferent between ($875, −$125)]

[If change implemented in C1Q5.6, participant indifferent between ($925, −$75)]

[If change not implemented in C1Q5.6, participant indifferent between ($975, −$25)]

[If first decision is ($700, −$300)]

[If change implemented in C1Q5.1, participant indifferent between ($25, −$975)]

[If change not implemented in C1Q5.1, participant indifferent between ($75, −$925)]

[If change implemented in C1Q5.2, participant indifferent between ($125, −$875)]

[If change not implemented in C1Q5.2, participant indifferent between ($175, −$825)]

[If change implemented in C1Q5.3, participant indifferent between ($225, −$775)]

[If change not implemented in C1Q5.3, participant indifferent between ($275, −$725)]

[If change implemented in C1Q5.4, participant indifferent between ($325, −$675)]

[If change not implemented in C1Q5.4, participant indifferent between ($375, −$625)]

[If change implemented in C1Q5.5, participant indifferent between ($425, −$575)]

[If change not implemented in C1Q5.5, participant indifferent between ($475, −$525)]

[If change implemented in C1Q5.6, participant indifferent between ($525, −$475)]

[If change not implemented in C1Q5.6, participant indifferent between ($575, −$425)]

——— Screen Break ———

[*Note: Comparisons 2-5 are identical to Comparison 1, with the following exceptions:
In Comparison 2, the Recipients are Person #2: $30,000 and Person #3: $60,000;
In Comparison 3, the Recipients are Person #3: $60,000 and Person #4: $120,000;
In Comparison 4, the Recipients are Person #3: $60,000 and Person #5: $240,000;
In Comparison 5, the Recipients are Person #3: $60,000 and Person #6: $480,000*]

[*Note: For half the participants, the order of the comparisons is reversed*]

——— Screen Break ———

[<u>Consistency check</u>]

In this final comparison, all participants will see a pair of individuals that they have encountered before. It is very important for us that you once again consider each question carefully. Thank you very much.

[In this comparison, the Recipients are Person #3: $60,000 and Person #4: $120,000]

——— Screen Break ———

How confident are you that the choices you made in the previous screens reflect what you really think?

[Not at all; Very little; Little; Somewhat; Very much]

Screen Break

[*Note. Participants are randomized to various treatments. The first treatment dimension varies the order of the comparisons. For half the participants, the order is as presented above, while for the other half, the order of the first five comparisons is reversed. The sixth comparison is identical across the two treatments. The second treatment dimension varies whether the Recipient common across the comparisons earns an income of $60,000 (as shown above) or $120,000. The third treatment dimension varies whether the low-income Recipient is presented on the left (as shown above) or the right of the screen. The fourth treatment dimension varies whether participants face a consequence verification on implied outcomes prior to proceeding. Participants with a such verifications are shown the two final-income distributions implied by their decision and must select the distribution consistent with their choices before advancing. If their preference on whether to implement the reform and their preferred final income distribution are consistent, they can proceed with the survey. If there is an inconsistency, they see the following error message: "Your choice on whether to implement the change (above) and your preferred distribution of final incomes (below) contradict each other. Please make your choice and indicate your preferred distribution of final incomes once again." This verification is only presented in the first decision of each comparison. Figure A1 presents a screenshot. The fifth treatment varies the questions in the staircase (explained above), with the reform in the first decision being either ($300, −$700) or ($700, −$300). The first four treatment dimensions have implications for the instructions and decisions, while the fifth dimension only has implications for the decisions.*]

Screen Break

[**Block 3: Support for Redistribution**]

[*Note: The order of the two questions is counterbalanced across participants.*]

In the following screens, we would like to ask you some general questions about your views on society. Your opinion and thoughts are important to us.

Some people think that the government in Washington ought to reduce the income differences between the rich and the poor, perhaps by raising the taxes of wealthy families or by giving income assistance to the poor. Others think that the government should not concern itself with reducing this income difference between the rich and the poor.

Here is a scale from 1 to 7. Think of a score of 1 as meaning that the government ought to reduce the income differences between rich and poor, and a score of 7 meaning that the government should not concern itself with reducing income differences. What score between 1 and 7 comes closest to the way you feel?

[1: Government should do something to reduce income differences between rich and poor; 2; 3; 4; 5:; 6:; 7: Government should not concern itself with income differences]

——————————— Screen Break ———————————

Consider the current **disposable incomes** of individuals in society. Disposable income is defined as **income after all taxes have been paid and transfers have been received.**

Do you think that, given the current disposable incomes of individuals in society, incomes should be further redistributed or should they remain as they are?

Please provide your answer on a scale from -2 to $+2$.

- A -2 means that income should be further redistributed by taking from lower/middle-income individuals and giving to higher-income individuals.
- A $+2$ means that income should be further redistributed by taking from higher-income individuals and giving to lower/middle-income individuals.

Given the current disposable incomes of individuals in society ...

[-2: Incomes should be further redistributed by taking from lower/middle-income individuals and giving to higher-income individuals; -1:; +0: Incomes should remain as they are; +1:; +2: Income should be further redistributed by taking from higher-income individuals and giving to lower/middle-income individuals].

——————————— Screen Break ———————————

[**Block 4: Knowledge**]

The next set of questions is about the income tax system in the United States. In order for your answers to be most helpful to us, it is really important that you provide your best guesses to these questions. Although you may find some questions difficult, it is very important for our research that you try your best. Thank you very much!

In 2024, individuals (single filers) with income over $609,350 were in the top federal personal income tax bracket.

Out of every 100 households in the U.S., how many are in the top federal personal income tax bracket?

[slider 0-100]

What is the marginal income tax rate applied to incomes at the top federal personal income tax bracket?

[slider 0%-100%]

What share of their total income do people in the top federal personal income tax bracket pay in taxes?

[slider 0-100]

Out of every 100 U.S. households, how many pay no federal income taxes?

[slider 0-100]

Imagine a middle class household that is right at the middle of the income distribution, such that half of all households in the U.S. earn more than this household and half earn less.

What share of their income do you think such a household pays in federal income taxes?

[slider 0-100]

Out of every 100 individuals in the U.S., how many earn a disposable income below $15,000?

[slider 0-100]

——————————————— Screen Break ———————————————

[**Block 5: Mechanisms**]

[*Note: The order of the questions is randomized across participants*]

Please answer the following last set of questions.

If the federal personal income tax rate were to increase for the richest people in the economy, to what extent would it encourage them to work less?

[A great deal; A lot; A moderate amount; A little; None at all]

Do you think that increasing income taxes on high-income households would hurt economic activity, not have an effect on economic activity, or help economic activity in the U.S.?

[Hurt economic activity in the U.S.; Not have an effect on economic activity in the U.S.; Help economic activity in the U.S.]

Typically, when the top federal income tax rate on high earners is cut, do you think that the lower class and working class mostly win or mostly lose from this change?

[Mostly lose; Neither lose nor win; Mostly win]

Consider the current disposable incomes of individuals in society, defined as income after all taxes have been paid and transfers have been received. Do you think that the current distribution of disposable incomes in society is unfair or fair?

[Very unfair; Somewhat unfair; Neither unfair nor fair; Somewhat fair; Very fair]

How much of the time do you think you can trust the federal government to do what is right?

[Never; Only some of the time; Most of the time; Just about always]

Some people think the government is trying to do too many things that should be left to individuals and businesses. Others think that the government should do more to solve our country's problems. Which comes closer to your own view?

[Government is doing too much; Government is doing just the right amount; Government should do more]

——————————————— Screen Break ———————————————

End of survey

Thank you for your time!

We will pay you your $2.5 participation payment in the following days.

Please click the following link to finish the survey:

[link]

G.2 Wave 2

[Consent screen]

Introduction
Welcome! This study is conducted by Unidistance, Switzerland and WZB, Germany. You previously participated in our survey. We invite you to participate in our follow-up survey.

Time required
This study will take around **5 minutes**.

Compensation
You will receive **$1** for completing the study. The payment will be made in the next few days.

Checks
Our survey includes attention checks to test whether participants take our survey carefully. Additionally, we have implemented measures to ensure that participants do not use AI assistance during the surveys. Participants must complete the survey independently, without the help of AI tools. Participants who fail these checks cannot proceed with the survey, and will be asked to return the survey.

Confidentiality
Your answers will remain anonymous and will be used for scientific purposes only. Strict confidentiality is guaranteed, and your identity can never be associated with your answers.

Voluntary participation
Participation in this study is voluntary. You may withdraw from the study at any time.

Questions about the survey
If you have questions about this study or your rights, please get in touch with us at krishna.srinivasan@unidistance.ch

Consent
I have received the above information and I am willing to participate in the study.

[Yes; No]

What is your Prolific ID?

Screen Break

[Screen shown if participant does not provide consent]

End of survey

You did not give your consent to continue with the study.

Please close this survey and return your submission on Prolific by selecting the "Stop without completing" button.

Screen Break

[Attention check]

In surveys like ours, some participants do not carefully read the questions. This means that there are a lot of random answers that can compromise the results of research studies. To show that you read our questions carefully, please choose "Not at all interested" below:

[Extremely interested; Very interested; A little bit interested; Almost not interested; Not at all interested]

Screen Break

[Screen shown if participant failed the attention check]

End of survey

Unfortunately, you failed the attention check.

For this reason, you cannot continue the study and will not receive a payment.

Please close this survey and return your submission on Prolific by selecting the "Stop without completing" button.

Screen Break

In the next part, we will present you with several decisions. **These decisions may look very similar to the ones you faced in our previous survey.**

However, it is very important for us that you consider these questions carefully.

——————————————— Screen Break ———————————————

[**Block 1: Eliciting Welfare Weights**]

[*Participants are presented with questions designed to elicit their welfare weights. The questions are identical to those in Wave 1. Each participant is assigned to the same treatment group in both waves. Only one sentence in the instructions for Wave 2 is different from those in Wave 1. The sentence in Wave 1 "If you are selected, we will randomly select one question (from this survey or the follow-up survey) and implement your choice on this question." is replaced in Wave 2 with "If you are selected, we will randomly select one question (from this survey or the previous survey) and implement your choice on this question."*]

——————————————— Screen Break ———————————————

[**Block 2: Support for Redistribution**]

[*Note: Participants are presented with both questions on support for redistribution shown to participants in Block 3 of Wave 1*]

H Instructions - Study 2

Bold text, underlining, tables, etc., appear as on the original screens.

H.1 Treatment Loss × 70K

[Consent screen]

Introduction
Welcome to this research study. We appreciate your participation. We are a non-partisan group of researchers from University of Zurich and Erasmus University Rotterdam. This study contains real choices and questions regarding your demographic characteristics. No matter what your political views are, by completing this survey you are contributing to our knowledge as a society.

Time required
Approximately **10 minutes.** You will have a maximum of one hour to finish the survey after starting it.

Requirements
You must be a **U.S. resident** to participate in this study. You must also be above the age of 18. The survey contains attention checks. You must pass these check in order to proceed with the survey.

Confidentiality
All data obtained from you will be used for research purposes only. Data will be anonymized immediately after collection. Researchers will at no point have access to any information that could be used to personally identify you.

Voluntary participation
It is voluntary to participate in the project, and you can at any time choose to withdraw your consent without stating any reason.

Questions about the Survey
If you have questions about this study or your rights, please get in touch with us at Krishna.srinivasan@econ.uzh.ch

Consent

I have received the above information about the project and am willing to participate.

[Yes; No]

——————————— Screen Break ———————————

[Screen shown if participant does not provide consent]

End of survey

You did not give your consent to continue with the study.

Thank you for your time.

You will be automatically redirected in 5 seconds.

——————————— Screen Break ———————————

[Block 1: Background Questions]

What is your sex?

[Male; Female]

How old are you?

[18 years old - 34 years old; 35 years old - 44 years old; 45 years old - 54 years old; 55 years old - 64 years old; Above 65 years old]

In which state do you currently reside?

[Northeast (ME, NH, VT, MA, CT, RI, NY, PA, NJ); Midwest (OH, MI, IN, WI, IL, MN, IA, MO, ND, SD, NE, KS); South (DE, MD, DC, VA, WV, NC, SC, GA, FL, KY, TN, AL, MS, AR, LA, OK, TX); Pacific (MT, WY, CO, NM, ID, UT, AZ, NV, WA, OR, CA, AK, HI); I do not reside in the US]

What is the highest level of education you have completed?

[Less than High School; High School/GED; Some College; Associate's Degree; Bachelor's degree; Master's degree; Doctoral or Profession Degree (PhD, ED.D, JD, DVM, DO, MD, DDS, or similar)]

As of today, do you consider yourself a Republican, a Democrat, or an Independent?

[Republican; Democrat; Independent]

The next question is about your **total individual income in 2020 before taxes**. This figure should include income from all sources, including salaries, wages, pensions, Social Security, dividends, interest, and all other income. What was your total individual income (USD) in 2020?

[$29,999 and below; $30,000 to $59,999; $60,000 to $99,999; $100,000 to $149,999; $150,00 and above

[Displayed if $29,999 and below is chosen]

You have reported that your total individual income in 2020 before taxes was $29,999 and below.

[Displayed if $30,000 to $59,999 is chosen]

You have reported that your total individual income in 2020 before taxes was between $30,000 and $59,999.

[Displayed if $60,000 to $99,999 is chosen]

You have reported that your total individual income in 2020 before taxes was between $60,000 and $99,999.

[Displayed if $100,000 to $149,999 is chosen]

You have reported that your total individual income in 2020 before taxes was between $100,000 and $149,999.

[Displayed if $150,000 and above is chosen]

You have reported that your total individual income in 2020 before taxes was above $150,000.

[Displayed in all cases]

Could you provide your best guess of what your **total individual income** was?

Screen Break

[If quotas are full]

End of survey

Unfortunately, we already have the number of participants needed for this study.

Thank you for your time.

You will be automatically redirected in 5 seconds.

——— Screen Break ———

[If a participant does not reside in the U.S.]

End of survey

Unfortunately, you do not fulfil the requirements of this study since you do not reside in the U.S.

Thank you for your time.

You will be automatically redirected in 5 seconds.

——— Screen Break ———

[Attention check]

In surveys like ours, some participants do not carefully read the questions. This means that there are a lot of random answers that can compromise the results of research studies. To show that you read our questions carefully, please choose both "Extremely interested" and "Not at all interested" below:

[Extremely interested; Very interested; A little bit interested; Almost not interested; Not at all interested]

——— Screen Break ———

[Screen shown if participant failed the attention check]

End of survey

Sorry, you failed the attention check. You were supposed to select both "Extremely interested" and "Not at all interested."

You cannot continue with the study.

Thank you for your time.

You will be automatically redirected in 5 seconds.

——————————————— Screen Break ———————————————

[**Block 2: Eliciting Welfare Weights**]

[Instructions screen]

Instructions

In this study, you will make several choices involving **seven real people**. These people will be selected at random from a survey panel and will not participate in the same survey as you. These people are above the age of 18 and are U.S. citizens. The incomes of the seven people are as follows:

Person	**After-tax annual income**
Person A	\$8000
Person B	\$35,000
Person C	\$70,000
Person D	\$100,000
Person E	\$170,000
Person F	\$250,000
Person G	\$500,000

Here is an example of a question that you will see in the survey:

	Person C	Person G
After-tax annual income	**$70,000**	**$500,000**

Question 2/4: Please choose your preferred alternative

Person C: +$750 Person G: -$1250	Person C: +$500 Person G: -$500

In this question, if you choose the option on the left, then $1250 will be taken away from Person G and $750 will be given to Person C. If you choose the option on the right, then $500 will be taken away from Person G and $500 will be given to Person C.

If you choose the option on the left, the final incomes of the two people **(including an initial $1500 bonus)** will be Person C: $72,250 and Person G: $500,250. If you choose the option on the right, the final incomes of the two people (including an initial $1500 bonus) will be Person C: $72,000 and Person G: $501,000.

You will face four questions like the one you saw above in each "decision screen." **Overall, you will face six decision screens with four questions in each.** In each question, you will see a different amount in the option on the left. In each decision screen, you will see a different pair of people.

There is a chance that you may be randomly selected in this study. If you are randomly selected, your choice on one randomly selected question on one randomly selected decision screen will be implemented. **This means that if you are randomly selected, one of your choices will have real consequences for two other people.** The final bonus of these two people will be transferred to them at the end of the study.

Please answer the following questions to demonstrate that you have understood the instructions. You can read the instructions above again if you feel the need to.

Please state True or False: "In this study, you will make several choices involving seven real people."

[True; False]

Please state True or False: "If you are randomly selected, one of your choices will have real consequences for two other people."

[True; False]

(You will be allowed to move to the next screen in 30 seconds)

——————————— Screen Break ———————————

[If a participant fails the comprehension check]

End of survey

The correct answers were "True" and "True". You answered incorrectly.

You cannot continue with the study.

Thank you for your time.

You will be automatically redirected in 5 seconds.

——————————— Screen Break ———————————

[C1Q1: Comparison 1 Question 1]

Decision Screen 1/6

Please consider each question carefully because if you are selected, one of your choices will have real consequences for two other persons.

	Person A	Person C
After-tax annual income	**$8,000**	**$70,000**

Question 1/4: Please choose your preferred alternative

Person A: +$1000 Person C: -$1000	Person A: +$500 Person C: -$500

Screen Break

[*Note: For brevity, the remaining questions in each comparison are presented in abbreviated form. Each bracketed line describes the question shown to participants given their path through the staircase.*]

[C1Q2.1: If $(500, -500)$ chosen in C1Q1, choose between $(1250, -750)$ and $(500, -500)$]

[C1Q2.2: If $(1000, -1000)$ chosen in C1Q1, choose between $(750, -1250)$ and $(500, -500)$]

Screen Break

[C1Q3.1: If $(500, -500)$ chosen in C1Q2.1, choose between $(1375, -625)$ and $(500, -500)$]

[C1Q3.2: If $(1250, -750)$ chosen in C1Q2.1, choose between $(1125, -875)$ and $(500, -500)$]

[C1Q3.3: If $(500, -500)$ chosen in C1Q2.2, choose between $(875, -1125)$ and $(500, -500)$]

[C1Q3.4: If $(750, -1250)$ chosen in C1Q2.2, choose between $(625, -1375)$ and $(500, -500)$]

Screen Break

[C1Q4.1: If $(500, -500)$ chosen in C1Q3.1, choose between $(1450, -550)$ and $(500, -500)$]

[C1Q4.2: If $(1375, -625)$ chosen in C1Q3.1, choose between $(1300, -700)$ and $(500, -500)$]

[C1Q4.3: If $(500, -500)$ chosen in C1Q3.2, choose between $(1200, -800)$ and $(500, -500)$]

[C1Q4.4: If $(1125, -875)$ chosen in C1Q3.2, choose between $(1050, -950)$ and $(500, -500)$]

[C1Q4.5: If $(500, -500)$ chosen in C1Q3.3, choose between $(950, -1050)$ and $(500, -500)$]

[C1Q4.6: If $(875, -1125)$ chosen in C1Q3.3, choose between $(800, -1200)$ and $(500, -500)$]

[C1Q4.7: If $(500, -500)$ chosen in C1Q3.4, choose between $(700, -1300)$ and $(500, -500)$]

[C1Q4.8: If $(625, -1375)$ chosen in C1Q3.4, choose between $(550, -1450)$ and $(500, -500)$]

Screen Break

[*Note: Comparisons 2-6 are identical to Comparison 1, with the following exceptions:*
In Comparison 2, the Recipients are: Person B: $35,000 and C: $70,000;
In Comparison 3, the Recipients are Person C: $70,000 and D: $100,000;
In Comparison 4, the Recipients are Person C: $70,000 and E: $170,000;
In Comparison 5, the Recipients are Person C: $70,000 and F: $250,000;
In Comparison 6, the Recipients are Person C: $70,000 and G: $500,000]

[*Note: For half the participants, the order of the Comparisons is reversed*]

Screen Break

[**Block 3: Support for Redistribution**]

[*Note: The order of the two questions is counterbalanced across participants.*]

We have some final questions. It is important for us that you answer them carefully.

The top income tax category in 2020 includes those with an annual individual income of over $518,400. Do you think that income taxes levied on these people in the top income category should be increased, stay the same, or decreased?

[1: Increased a lot; 2:; 3:; 4: Stay the same; 5; 6; 7: Decreased a lot]

Some people think that the government in Washington ought to reduce the income differences between the rich and the poor, perhaps by raising the taxes of wealthy families or by giving income assistance to the poor. Others think that the government should not concern itself with reducing this income difference between the rich and the poor.

Here is a scale from 1 to 7. Think of a score of 1 as meaning that the government ought to reduce the income differences between rich and poor, and a score of 7 meaning that the government should not concern itself with reducing income differences. What score between 1 and 7 comes closest to the way you feel?

[1: Government should do something to reduce income differences between rich and poor; 2; 3; 4; 5:; 6:; 7: Government should not concern itself with income differences]

———————— Screen Break ————————

End of survey

Thank you for your time!

You will be automatically redirected in 5 seconds.

H.2 Treatment Loss $\times$ 500K

[*Note: All questions are identical to those in Treatment Loss $\times$ 70K. The incomes of the Recipients in the comparisons are as follows:*
In Comparison 1, the Recipients are: Person A: \$8,000 vs. G: \$500,000;
In Comparison 2, the Recipients are: Person B: \$35,000 vs. G: \$500,000;
In Comparison 3, the Recipients are Person C: \$70,000 vs. G: \$500,000;
In Comparison 4, the Recipients are Person D: \$100,000 vs. G: \$500,000;
In Comparison 5, the Recipients are Person E: \$170,000 vs. G: \$500,000;
In Comparison 6, the Recipients are Person F: \$250,000 vs. G: \$500,000]

[*Note: For half the participants, the order of the Comparisons is reversed*]

H.3 Treatment Gain × 70K

[All questions, with the exceptions of those listed below, are identical to those in Treatment Loss × 70K]

[Instructions screen]

Instructions

In this study, you will make several choices involving **seven real people**. These people will be selected at random from a survey panel and will not participate in the same survey as you. These people are above the age of 18 and are U.S. citizens. The incomes of the seven people are as follows:

Person	**After-tax annual income**
Person A	$8000
Person B	$35,000
Person C	$70,000
Person D	$100,000
Person E	$170,000
Person F	$250,000
Person G	$500,000

Here is an example of a question that you will see in the survey:

	Person C	**Person G**
After-tax annual income	**$70,000**	**$500,000**

Question 2/4: Please choose your preferred alternative

Person C: +$2250 Person G: +$250	Person C: +$2000 Person G: +$1000

In this question, if you choose the option on the left, then $250 will be given to Person G and $2250 will be given to Person C. If you choose the option on the right, then $1000 will be given to Person G and $2000 will be given to person C.

If you choose the option on the left, the final incomes of the two people will be Person C: $72,250 and Person G: $500,250. If you choose the option on the right, the final incomes of the two people will be Person C: $72,000 and Person G: $501,000.

You will face four questions like the one you saw above in each "decision screen." **Overall, you will face six decision screens with four questions in each.** In each question, you will see a different amount in the option on the left. In each decision screen, you will see a different pair of people.

There is a chance that you may be randomly selected in this study. If you are randomly selected, your choice on one randomly selected question on one randomly selected decision screen will be implemented. **This means that if you are randomly selected, one of your choices will have real consequences for two other people.** The final bonus of these two people will be transferred to them at the end of the study.

Please answer the following questions to demonstrate that you have understood the instructions. You can read the instructions above again if you feel the need to.

Please state True or False: "In this study, you will make several choices involving seven real people."

[True; False]

Please state True or False: "If you are randomly selected, one of your choices will have real consequences for two other people."

[True; False]

(You will be allowed to move to the next screen in 30 seconds)

——— Screen Break ———

[*Note: The incomes of the Recipients in the six Comparisons are identical to the incomes of the Recipients in Treatment Loss × 70K. All comparisons have the following questions.*]

[C1Q1: Choose between $(2500, 500)$ and $(2000, 1000)$]

——— Screen Break ———

[C1Q2.1: If $(2000, 1000)$ chosen in C1Q1, choose between $(2750, 750)$ and $(2000, 1000)$]

[C1Q2.2: If $(2500, 500)$ chosen in C1Q1, choose between $(2250, 250)$ and $(2000, 1000)$]

——————————— Screen Break ———————————

[C1Q3.1: If $(2000, 1000)$ chosen in C1Q2.1, choose between $(2875, 875)$ and $(2000, 1000)$]

[C1Q3.2: If $(2750, 750)$ chosen in C1Q2.1, choose between $(2625, 625)$ and $(2000, 1000)$]

[C1Q3.3: If $(2000, 1000)$ chosen in C1Q2.2, choose between $(2375, 375)$ and $(2000, 1000)$]

[C1Q3.4: If $(2250, 250)$ chosen in C1Q2.2, choose between $(2125, 125)$ and $(2000, 1000)$]

——————————— Screen Break ———————————

[C1Q4.1: If $(2000, 1000)$ chosen in C1Q3.1, choose between $(2950, 950)$ and $(2000, 1000)$]

[C1Q4.2: If $(2875, 875)$ chosen in C1Q3.1, choose between $(2800, 800)$ and $(2000, 1000)$]

[C1Q4.3: If $(2000, 1000)$ chosen in C1Q3.2, choose between $(2700, 700)$ and $(2000, 1000)$]

[C1Q4.4: If $(2625, 625)$ chosen in C1Q3.2, choose between $(2550, 550)$ and $(2000, 1000)$]

[C1Q4.5: If $(2000, 1000)$ chosen in C1Q3.3, choose between $(2450, 450)$ and $(2000, 1000)$]

[C1Q4.6: If $(2375, 375)$ chosen in C1Q3.3, choose between $(2300, 300)$ and $(2000, 1000)$]

[C1Q4.7: If $(2000, 1000)$ chosen in C1Q3.4, choose between $(2200, 200)$ and $(2000, 1000)$]

[C1Q4.8: If $(2125, 125)$ chosen in C1Q3.4, choose between $(2050, 50)$ and $(2000, 1000)$]

H.4 Treatment Gain × 500K

[<u>All questions are identical to those in Treatment Gain × 70K, while the incomes of the Recipients in the six Comparisons are identical to those in Treatment Loss × 500K.</u>]

I Instructions - Study 3

Bold text, underlining, tables, etc., appear as on the original screen.

I.1 Treatment Real

[Consent screen]

Introduction
Welcome to this research study. We appreciate your participation. We are a non-partisan group of researchers from University of Zurich and Erasmus University Rotterdam. This study contains real choices and questions regarding your demographic characteristics. No matter what your political views are, by completing this survey you are contributing to our knowledge as a society.

Time required
Approximately **12 minutes.**

Requirements
You must be a U.S. resident to participate in this study. You must also be above the age of 18. The survey contains attention checks. You must pass these check in order to proceed with the survey.

Confidentiality
All data obtained from you will be used for research purposes only. Data will be anonymized immediately after collection. Researchers will at no point have access to any information that could be used to personally identify you.

Voluntary participation
It is voluntary to participate in the project, and you can at any time choose to withdraw your consent without stating any reason.

Questions about the Survey
If you have questions about this study or your rights, please get in touch with us at Krishna.srinivasan@econ.uzh.ch

Consent
I have received the above information about the project and am willing to participate.

[Yes; No]

What is your prolific ID?

——————————— Screen Break ———————————

[Screen shown if participant does not provide consent]

You did not give your consent to continue with the study.

Thank you for your time.

Please return your submission on Prolific by selecting the 'Stop without completing' button.

——————————— Screen Break ———————————

[Block 1: Background Questions]

What is your sex?

[Male; Female]

How old are you?

[18 years old - 34 years old; 35 years old - 44 years old; 45 years old - 54 years old; 55 years old - 64 years old; 65 years old or above]

In which state do you currently reside?

[Alabama; ...; Wyoming; I do not reside in the U.S.]

In which ZIP code do you live? (5 digits)

What is the highest level of education you have completed?

[Less than High School; High School/GED; Some College; Associate's Degree; Bachelor's degree; Master's degree; Doctoral or Profession Degree (PhD, ED.D, JD, DVM, DO, MD, DDS, or similar)]

As of today, do you consider yourself a Republican, a Democrat, or an Independent?

[Republican; Democrat; Independent]

The next question is about your **total individual income in 2021 before taxes**. This figure should include income from all sources, including salaries, wages, pensions, social security, dividends, interest, and all other income. What was your total individual income (USD) in 2021?

[$29,999 and below; $30,000 to $59,999; $60,000 to $99,999; $100,000 to $149,999; $150,000 and above]

——————————— Screen Break ———————————

[Displayed if $29,999 and below is chosen]

You have reported that your total individual income in 2021 before taxes was $29,999 and below.

[Displayed if $30,000 to $59,999 is chosen]

You have reported that your total individual income in 2021 before taxes was $30,000 to $59,999.

[Displayed if $60,000 to $99,999 is chosen]

You have reported that your total individual income in 2021 before taxes was $60,000 to $99,999.

[Displayed if $100,000 to $149,999 is chosen]

You have reported that your total individual income in 2021 before taxes was $100,000 to $149,999.

[Displayed if $150,000 and above is chosen]

You have reported that your total individual income in 2021 before taxes was $150,000 and above.

[Displayed in all cases]

Could you provide your best guess of what your **total individual income** was?

——————————————— Screen Break ———————————————

[If a participant does not reside in the U.S.]

End of survey

Unfortunately, you do not fulfil the requirements of this study since you do not reside in the U.S.

Thank you for your time.

Please return your submission on Prolific by selecting the 'Stop without completing' button.

——————————————— Screen Break ———————————————

[Attention check]

In surveys like ours, some participants do not carefully read the questions. This means that there are a lot of random answers that can compromise the results of research studies. To show that you read our questions carefully, please choose both "Extremely interested" and "Not at all interested" below:

[Extremely interested; Very interested; A little bit interested; Almost not interested; Not at all interested]

——————————————— Screen Break ———————————————

[**Block 2: Eliciting Welfare Weights**]

[Instructions screen]

Instructions

In this study, you will make several choices involving **seven real people**. These people will be selected at random from a survey panel and will not participate in the same survey as

Person	**After-tax annual income**
Person A	$8,000
Person B	$35,000
Person C	$70,000
Person D	$100,000
Person E	$170,000
Person F	$250,000
Person G	$500,000

you. These people are above the age of 18 and are U.S. citizens. The incomes of the seven people **after all taxes paid and transfers received** are as follows:

Here is an example of a question that you will see in the survey:

	Person C	**Person G**
After-tax annual income	**$70,000**	**$500,000**

Question 2/4: Please choose your preferred alternative

Person C: +$750 Person G: -$1250	Person C: +$500 Person G: -$500

In this question, if you choose the option on the left, then $1250 will be taken away from Person G and $750 will be given to Person C. If you choose the option on the right, then $500 will be taken away from Person G and $500 will be given to Person C.

If you choose the option on the left, the final incomes of the two people **(including an initial $1500 bonus)** will be Person C: $72,250 and Person G: $500,250. If you choose the option on the right, the final incomes of the two people (including an initial $1500 bonus) will be Person C: $72,000 and Person G: $501,000.

You will face four questions like the one you saw above in each "decision screen." **Overall, you will face six decision screens with four questions in each.** In each question, you will see a different amount in the option on the left. In each decision screen, you will see a different pair of people.

One participant in this study will be randomly selected. If you are randomly selected, your choice on one randomly selected question on one randomly selected decision screen will be implemented. **This means that if you are randomly selected, one of your choices will have real consequences for two other people.** The final bonus of these two people will be transferred to them at the end of the study.

Please answer the following questions to demonstrate that you have understood the instructions. You can read the instructions above again if you feel the need to.

Please state True or False: "In this study, you will make several choices involving seven real people."

[True; False]

Please state True or False: "If you are randomly selected, one of your choices will have real consequences for two other people."

[True; False]

(You will be allowed to proceed to the next screen in 30 seconds)

[The timer updates dynamically. When the time elapses, the text disappears.]

——— Screen Break ———

[Shown if a participant fails at least two out of three checks (one attention check and two comprehension checks)]

End of survey

Sorry, you answered at least two out of three comprehension/attention checks incorrectly.

You cannot continue with the study.

Thank you for your time.

Please return your submission on Prolific by selecting the 'Stop without completing' button.

[If a participant fails only one out of three checks (one attention check and two comprehension checks)]

End of survey

Thank you for your time.

We will pay you your £2 participation fee in the following days.

Please click the following link to finish the survey.

——— Screen Break ———

[C1Q1: Comparison 1 Question]

Decision Screen 1/6

Please consider each question carefully because if you are selected, one of your choices will have real consequences for two other persons.

	Person A	**Person C**
After-tax annual income	**\$8,000**	**\$70,000**

Question 1/4: Please choose your preferred alternative:

Person A: +\$1000 Person C: -\$1000 ○	Person A: +\$500 Person C: -\$500 ○

——— Screen Break ———

[*Note: For brevity, the remaining questions in each comparison are presented in abbreviated form. Each bracketed line describes the question shown to participants given their path through the staircase.*]

[C1Q2.1: If $(500, -500)$ chosen in C1Q1, choose between $(1250, -750)$ and $(500, -500)$]

[C1Q2.2: If $(1000, -1000)$ chosen in C1Q1, choose between $(750, -1250)$ and $(500, -500)$]

——— Screen Break ———

[C1Q3.1: If $(500, -500)$ chosen in C1Q2.1, choose between $(1375, -625)$ and $(500, -500)$]

[C1Q3.2: If $(1250, -750)$ chosen in C1Q2.1, choose between $(1125, -875)$ and $(500, -500)$]

[C1Q3.3: If $(500, -500)$ chosen in C1Q2.2, choose between $(875, -1125)$ and $(500, -500)$]

[C1Q3.4: If $(750, -1250)$ chosen in C1Q2.2, choose between $(625, -1375)$ and $(500, -500)$]

———————————— Screen Break ————————————

[C1Q4.1: If $(500, -500)$ chosen in C1Q3.1, choose between $(1450, -550)$ and $(500, -500)$]

[C1Q4.2: If $(1375, -625)$ chosen in C1Q3.1, choose between $(1300, -700)$ and $(500, -500)$]

[C1Q4.3: If $(500, -500)$ chosen in C1Q3.2, choose between $(1200, -800)$ and $(500, -500)$]

[C1Q4.4: If $(1125, -875)$ chosen in C1Q3.2, choose between $(1050, -950)$ and $(500, -500)$]

[C1Q4.5: If $(500, -500)$ chosen in C1Q3.3, choose between $(950, -1050)$ and $(500, -500)$]

[C1Q4.6: If $(875, -1125)$ chosen in C1Q3.3, choose between $(800, -1200)$ and $(500, -500)$]

[C1Q4.7: If $(500, -500)$ chosen in C1Q3.4, choose between $(700, -1300)$ and $(500, -500)$]

[C1Q4.8: If $(625, -1375)$ chosen in C1Q3.4, choose between $(550, -1450)$ and $(500, -500)$]

———————————— Screen Break ————————————

[*Comparisons 2-6 are identical to Comparison 1, with the following exceptions:*
In Comparison 2, the Recipients are: Person B: $35,000 and C: $70,000;
In Comparison 3, the Recipients are Person C: $70,000 and D: $100,000;
In Comparison 4, the Recipients are Person C: $70,000 and E: $170,000;
In Comparison 5, the Recipients are Person C: $70,000 and F: $250,000;
In Comparison 6, the Recipients are Person C: $70,000 vs. G: $500,000]

[*For half the participants, the order of the Comparisons is reversed*]

———————————— Screen Break ————————————

How confident are you that the choices you made in the previous screens reflect what you really think?

Please provide your answer on a scale of 1 to 5. A 1 indicates "Not confident all," and a 5 indicates "Completely confident."

[5: Completely confident; 4:; 3:; 2:; 1: Not confident at all]

——————————— Screen Break ———————————

[**<u>Block 3: Welfarist and Non-Welfarist Concerns</u>**]

In the following screens, we would like to ask you some general questions about your views on society. Your opinion and thoughts are important to us.

Consider the current incomes of individuals in society obtained after all taxes are paid and transfers received.

Which of the following statements comes closest to how you feel?

High-income individuals ...

[do <u>not</u> deserve their current income and do <u>not</u> need their current income; deserve their current income but do <u>not</u> need their current income; do <u>not</u> deserve their current income but need their current income; deserve their current income and need their current income]

Which of the following statements comes closest to how you feel?

Low-income individuals ...

[do <u>not</u> deserve their current income and do <u>not</u> need their current income; deserve their current income but do <u>not</u> need their current income; do <u>not</u> deserve their current income but need their current income; deserve their current income and need their current income]

——————————— Screen Break ———————————

Consider the current incomes of individuals in society obtained after all taxes are paid and transfers received.

Do you think that, given the current incomes of individuals in society, incomes should be further redistributed or should not be further redistributed?

Please provide your answer on a scale from -2 to +2 where a +2 means that income should be further redistributed by taking from the higher-income individuals and giving to the lower/middle-income individuals while a -2 means that income should be further redistributed by taking from the lower/middle-income individuals and giving to the higher-income individuals.

[-2: Incomes should be further redistributed by taking from the lower/middle-income individuals and giving to the higher-income individuals; -1:; +0: Incomes should **not** be further redistributed; +1:; +2: Incomes should be further redistributed by taking from the higher-income individuals and giving to the lower/middle-income individuals]

——————————— Screen Break ———————————

[**Block 4: Knowledge**]

The next set of questions is about the income tax system in the United States. These are questions for which there are right or wrong answers.

In order for your answers to be most helpful to us, it is really important that you answer these questions as accurately as you can. Although you may find some questions difficult, it is very important for our research that you try your best. Thank you very much!

Out of 100 households in the U.S., how many are in the top federal personal income tax bracket?

[slider 0-100]

What share of their total income do people in the top federal personal income tax bracket pay in taxes?

[slider 0-100]

Out of 100 U.S. households, how many pay no federal income taxes?

[slider 0-100]

Imagine a middle class household that is right at the middle of the income distribution, such that half of all households in the U.S. earn more than this household and half earn less. What share of their income do you think such a household pays in federal income taxes?

[slider 0-100]

Out of every 100 individuals in the U.S., how many earn an income (after all taxes paid and transfers received) below $35,000?

[slider 0-100]

We would now like to ask you what you think about the life opportunities of children from very poor families.

For the following question, we focus on 500 families that represent the U.S. population. We divide them into five groups on the basis of their income, with each group containing 100 families. These groups are:

- The poorest 100 families
- The second poorest 100 families
- The middle 100 families
- The second richest 100 families
- The richest 100 families

How many out of 100 children coming from the poorest 100 families will grow up to be among the richest 100 families?

——— Screen Break ———

[**Block 4: Mechanisms**]

We would like to ask you what you think the distribution of after-tax income in the U.S. should be.

There are **7 tax groups** (tax brackets) in the U.S. Group 1 includes households with the lowest incomes and Group 7 includes households with the highest incomes. Groups 2 through 6 include households with incomes in the middle.

Column 2 of the table below lists the **CURRENT** average annual after-tax income of all households in each group. The after-tax income is obtained by subtracting all federal income taxes (e.g., ordinary income taxes, alternative minimum taxes) from the pre-tax income and adding all federal transfers (e.g., tax credits) to the pre-tax income.

In Column 3 of the table below, we list the average federal income tax rate of each group. This rate was determined based on the ordinary income taxes that households paid. As an example, if a household with a pre-tax income of $80,000 has an average tax rate of 15%, they would pay 80000*0.15 = $12,000 in taxes.

We would like you to indicate what you think the average tax rate for each tax group in the U.S. should be. This can be done as follows. **You can increase or decrease the average tax rates of the first six groups. The average tax rate of group 7 adjusts automatically so that all seven groups together pay as much taxes as they currently do.**

Column 4 of the table below and the figure below indicate your **DESIRED** average annual after-tax incomes. The numbers in the table as well as the figure update automatically as you change the average tax rates.

Your choices will sometimes be limited for a variety of reasons. For example, you cannot set the tax rate for a group such that their average after-tax income becomes lower than the average after-tax income of the group below them or higher than the average after-tax income of the group above them.

Note also that there may be rounding-off errors in various calculations.

You can go back to the initial situation by refreshing the page.

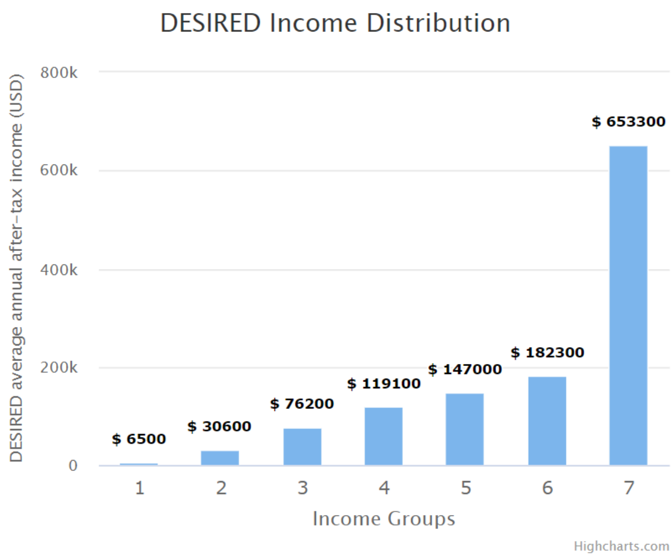

Income group	Annual after-tax income **(CURRENT)**	Average tax rate	Annual after-tax income **(DESIRED)**
1	$6,500	9%	$6,500
2	$30,600	11%	$30,600
3	$76,200	15%	$76,200
4	$119,100	19%	$119,100
5	$147,000	21%	$147,000
6	$182,300	25%	$182,300
7	$653,300	31%	$653,300

Screen Break

Please answer the following last set of questions.

Which has more to do with why a person is rich?

[Because she or he worked harder than others; Because she or he had more advantages than others]

If the federal personal income tax rate were to increase for the richest people in the economy, to what extent would it encourage them to work less?

[A great deal; A lot; A moderate amount; A little; None at all]

Do you think that increasing income taxes on high-income households would hurt economic activity, not have an effect on economic activity, or help economic activity in the U.S.?

[Hurt economic activity in the U.S.; Not have an effect on economic activity in the U.S.; Help economic activity in the U.S.]

Typically, when the top federal income tax rate on high earners is cut, do you think that the lower class and working class mostly win or mostly lose from this change?

[Mostly lose; Neither lose nor win; Mostly win]

Some people think that income inequality in society can affect the level of crime, trust, corruption, and social unrest in society.

How big of an issue do you think income inequality is in America?

[Not an issue at all; A small issue; An issue; A serious issue; A very serious issue]

How much of the time do you think you can trust the federal government to do what is right?

[Always; Most of the time; Only some times; Never]

Screen Break

End of survey

Thank you for your time!

We will pay you your £2 participation fee in the following days.

Please click the following link to finish the survey.

I.2 Treatment Hypothetical

[*Note: All questions, with the exceptions of those listed below, are identical to those in Treatment Real*]

[**Block 2: Eliciting Welfare Weights**]

[Instructions screen]

Instructions

In this study, you will make several choices involving **seven hypothetical people**. These people are not real but you should imagine them as above the age of 18 and U.S. citizens. The incomes of the seven people **after all taxes paid and transfers received** are as follows:

Person	**After-tax annual income**
Person A	$8,000
Person B	$35,000
Person C	$70,000
Person D	$100,000
Person E	$170,000
Person F	$250,000
Person G	$500,000

Here is an example of a question that you will see in the survey:

	Person C	**Person G**
After-tax annual income	**$70,000**	**$500,000**

Question 2/4: Please choose your preferred alternative

Person C: +$750 Person G: -$1250	Person C: +$500 Person G: -$500

In this question, if you choose the option on the left, then $1250 will be taken away from Person G and $750 will be given to Person C. If you choose the option on the right, then $500 will be taken away from Person G and $500 will be given to Person C.

If you choose the option on the left, the final incomes of the two people **(including an initial $1500 bonus)** will be Person C: $72,250 and Person G: $500,250. If you choose the option on the right, the final incomes of the two people (including an initial $1500 bonus) will be Person C: $72,000 and Person G: $501,000.

You will face four questions like the one you saw above in each "decision screen." **Overall, you will face six decision screens with four questions in each.** In each question, you will see a different amount in the option on the left. In each decision screen, you will see a different pair of people.

The choices you make in the survey will not have real consequences.

Please answer the following questions to demonstrate that you have understood the instructions. You can read the instructions above again if you feel the need to.

Please state True or False: "In this study, you will make several choices involving seven hypothetical people."

[True; False]

Please state True or False: "Your choices will **not** have real consequences."

[True; False]

(You will be allowed to proceed to the next screen in 30 seconds)

[The timer updates dynamically. When the time elapses, the text disappears.]

——————————— Screen Break ———————————

[<u>C1Q1: shown to all participants</u>]

Decision Screen 1/6

Please consider each question carefully.

	Person A	Person C
After-tax annual income	**$8,000**	**$70,000**

Question 1/4: Please choose your preferred alternative:

Person A: +$1000 Person C: -$1000 ○	Person A: +$500 Person C: -$500 ○

[*Note: All Comparisons and questions and identical to those in Treatment Real. Treatment Hypothetical omits the sentences about sentence about stakes shown in Treatment Real.*]

I.3 Treatment No Self-Interest

[*All questions, with the exceptions of those listed below, are identical those in Treatment Real*]

[**Block 1: Background Questions**]

[*Note: In the Demographics screen, all questions, with the exception of the question on own income, is the same as in Treatment Real*]

The next question is about your **total individual income in 2021 before taxes**. This figure should include income from all sources, including salaries, wages, pensions, Social Security, dividends, interest, and all other income. What was your total individual income (USD) in 2021?

[$22,000 and below; $22,000 to $53,000; $53,000 to $85,000; $85,000 to $135,000; $135,000 to $210,000; $210,000 to $375,000; $375,000 and above]

Screen Break

[Displayed if $22,000 and below is chosen]

You have reported that your total individual income in 2021 before taxes was $22,000 and below.

[Displayed if $22,000 to $53,000 is chosen]

You have reported that your total individual income in 2021 before taxes was $22,000 to $53,000.

[Displayed if $53,000 to $85,000 is chosen]

You have reported that your total individual income in 2021 before taxes was $53,000 to $85,000.

[Displayed if $85,000 to $135,000 is chosen]

You have reported that your total individual income in 2021 before taxes was $85,000 to $135,000.

[Displayed if $135,000 to $210,000 is chosen]

You have reported that your total individual income in 2021 before taxes was $135,000 to $210,000.

[Displayed if $210,000 to $375,000 is chosen]

You have reported that your total individual income in 2021 before taxes was $210,000 to $375,000.

[Displayed if $375,000 and above is chosen]

You have reported that your total individual income in 2021 before taxes was $375,000 and above.

[Displayed in all cases]

Could you provide your best guess of what your **total individual income** was?

——————————— Screen Break ———————————

[**Block 2: Eliciting Welfare Weights**]

[Instructions screen]

Instructions

In this study, you will make several choices involving **seven real people**. These people will be selected at random from a survey panel and will not participate in the same survey as you. These people are above the age of 18 and are U.S. citizens. The incomes of the seven people **after all taxes paid and transfers received** put them in the following income brackets:

Person	**After-tax annual income**
Person A	$22,000 and below
Person B	$22,000 to $53,000
Person C	$53,000 to $85,000
Person D	$85,000 to $135,000
Person E	$135,000 to $210,000
Person F	$210,000 to $375,000
Person G	$375,000 and above

Here is an example of a question that you will see in the survey:

	Person C	**Person G**
After-tax annual income	**\$53,000 to \$85,000**	**\$375,000 and above**

Question 2/4: Please choose your preferred alternative

Person C: +\$750 Person G: -\$1250	Person C: +\$500 Person G: -\$500

In this question, if you choose the option on the left, then \$1250 will be taken away from Person G and \$750 will be given to Person C. If you choose the option on the right, then \$500 will be taken away from Person G and \$500 will be given to Person C.

If you choose the option on the left, the final income brackets of the two people **(including an initial \$1500 bonus)** will be Person C: \$55,250 to \$87,250 and Person G: \$375,250 and above. If you choose the option on the right, the final incomes of the two people (including an initial \$1500 bonus) will be Person C: \$55,000 to \$87,000 and Person G: \$376,000 and above.

You will face four questions like the one you saw above in each "decision screen." **Overall, you will face six decision screens with four questions in each.** In each question, you will see a different amount in the option on the left. In each decision screen, you will see a different pair of people.

One participant in this study will be randomly selected. If you are randomly selected, your choice on one randomly selected question on one randomly selected decision screen will be implemented. **This means that if you are randomly selected, one of your choices will have real consequences for two other people.** The final bonus of these two people will be transferred to them at the end of the study.

Please answer the following questions to demonstrate that you have understood the instructions. You can read the instructions above again if you feel the need to.

Please state True or False: "In this study, you will make several choices involving seven real people."

[True; False]

Please state True or False: "If you are randomly selected, one of your choices will have real consequences for two other people."

[True; False]

(You will be allowed to proceed to the next screen in 30 seconds)

[The timer updates dynamically. When the time elapses, the text disappears.]

——— Screen Break ———

[C1Q1: shown to all participants]

Decision Screen 1/6

Please consider each question carefully because if you are selected, one of your choices will have real consequences for two other persons.

	Person A	**Person C**
After-tax annual income	**$22,000 and below**	**$53,000 to $85,000**

Question 1/4: Please choose your preferred alternative:

Person A: +$1000 Person C: -$1000 ○	Person A: +$500 Person C: -$500 ○

[*Note: All questions are identical to those in Treatment Real. Comparisons 1 to 6 are identical to the corresponding Comparisons in Treatment Real, with the exception that the incomes of the Recipients are different. The pair of Recipients they view is as follows:*
In Comparison 2, the Recipients are Person B: $22,000 to $53,000 and C: $53,000 to $85,000;
In Comparison 3, the Recipients are Person C: $53,000 to $85,000 and D: $85,000 to $135,000;
In Comparison 4, the Recipients are Person C: $53,000 to $85,000 and E: $135,000 to

$210,000;
In Comparison 5, the Recipients are Person C: $53,000 to $85,000 and F: $210,000 to $375,000;
In Comparison 6, the Recipients are Person C: $53,000 to $85,000 and G: $375,000 and above]

[*For half the participants, the order of the Comparisons is reversed*]

I.4 Treatment Self-Interest

[*All questions, with the exception of those listed below, are identical to those in Treatment No Self-Interest*]

[**Block 2: Eliciting Welfare Weights**]

[Instructions screen]

Instructions

In this study, you will make several choices involving **six real people** and you. These six people will be selected at random from a survey panel and will not participate in the same survey as you. These people are above the age of 18 and are U.S. citizens. The incomes of the six people **after all taxes paid and transfers received** put them in the following income brackets:

Note that in this study, you are Person [A/B/C/D/E/F/G] earning [income].

Person	**After-tax annual income**
Person A	$22,000 and below
Person B	$22,000 to $53,000
Person C	$53,000 to $85,000
Person D	$85,000 to $135,000
Person E	$135,000 to $210,000
Person F	$210,000 to $375,000
Person G	$375,000 and above

Here is an example of a question that you will see in the survey:

	Person C	**Person G**
After-tax annual income	**$53,000 to $85,000**	**$375,000 and above**

Question 2/4: Please choose your preferred alternative

Person C: +$750 Person G: -$1250	Person C: +$500 Person G: -$500

In this question, if you choose the option on the left, then $1250 will be taken away from Person G and $750 will be given to Person C. If you choose the option on the right, then $500 will be taken away from Person G and $500 will be given to Person C.

If you choose the option on the left, the final income brackets of the two people **(including an initial $1500 bonus)** will be Person C: $55,250 to $87,250 and Person G: $375,250 and above. If you choose the option on the right, the final incomes of the two people (including an initial $1500 bonus) will be Person C: $55,000 to $87,000 and Person G: $376,000 and above.

You will face four questions like the one you saw above in each "decision screen." **Overall, you will face six decision screens with four questions in each.** In each question, you will see a different amount in the option on the left. In each decision screen, you will see a different pair of people.

Remember that in this study, you are Person [A/B/C/D/E/F/G] earning [income].

One participant in this study will be randomly selected. If you are randomly selected, your choice on one randomly selected question on one randomly selected decision screen will be implemented. **This means that if you are randomly selected, one of your choices will have real consequences. If the selected question involves a payment to you, then we will pay out the bonus to you and to the other person. If the selected question involves a payment to two other persons, then we will pay out the bonus to these two other persons.** The final bonus will be transferred at the end of the study. If you are among the winners, we will contact you in a few months and pay out your bonus via prolific.

Please answer the following questions to demonstrate that you have understood the instructions. You can read the instructions above again if you feel the need to.

Please state True or False: "In this study, you will make several choices involving six real people and you."

[True; False]

Please state True or False: "If you are randomly selected, one of your choices will have real consequences for two other people or for you and one other person."

[True; False]

(You will be allowed to proceed to the next screen in 30 seconds)

[The timer updates dynamically. When the time elapses, the text disappears.]

———————————————— Screen Break ————————————————

[C1Q1: Comparison 1 Question 1]

Decision Screen 1/6

Please consider each question carefully because if you are selected, one of your choices will have real consequences.

	Person A	**Person C**
After-tax annual income	**$22,000 and below**	**$53,000 to $85,000**

Question 1/4: Please choose your preferred alternative:

Person A: +$1000 Person C: -$1000 ○	Person A: +$500 Person C: -$500 ○

[*All questions and Comparisons are identical to those in Treatment No Self-Interest except that in the relevant Comparisons, we replace "Person [A/B/C/D/E/F/G]" with "You." Furthermore, the first sentence in all Comparisons is different.*]

References

Bursztyn, L., Handel, B., Jiménez-Durán, R. and Roth, C. (2025), 'When product markets become collective traps: The case of social media', *American Economic Review* **115**(12), 4105–4136.

Chen, D. L., Schonger, M. and Wickens, C. (2016), 'oTree—An open-source platform for laboratory, online, and field experiments', *Journal of Behavioral and Experimental Finance* **9**, 88–97.

Feenberg, D. and Coutts, E. (1993), 'An introduction to the taxsim model', *Journal of Policy Analysis and Management* **12**(1), 189–194.

Gelbach, J. B. (2016), 'When do covariates matter? And which ones, and how much?', *Journal of Labor Economics* **34**(2), 509–543.

Hendren, N. (2020), 'Measuring economic efficiency using inverse-optimum weights', *Journal of Public Economics* **187**, 104198.

Hendren, N. and Sprung-Keyser, B. (2020), 'A unified welfare analysis of government policies', *The Quarterly Journal of Economics* **135**(3), 1209–1318.

Luttmer, E. F. P. and Samwick, A. A. (2018), 'The welfare cost of perceived policy uncertainty: Evidence from social security', *American Economic Review* **108**(2), 275–307.

Mankiw, N. G., Weinzierl, M. and Yagan, D. (2009), 'Optimal taxation in theory and practice', *Journal of Economic Perspectives* **23**(4), 147–174.

Piketty, T., Saez, E. and Zucman, G. (2018), 'Distributional national accounts: methods and estimates for the united states', *The Quarterly Journal of Economics* **133**(2), 553–609.

Saez, E. (2001), 'Using elasticities to derive optimal income tax rates', *The Review of Economic Studies* **68**(1), 205–229.

Saez, E., Slemrod, J. and Giertz, S. H. (2012), 'The elasticity of taxable income with respect to marginal tax rates: A critical review', *Journal of Economic Literature* **50**(1), 3–50.

Stantcheva, S. (2021), 'Understanding tax policy: How do people reason?', *The Quarterly Journal of Economics* **136**(4), 2309–2369.

Støstad, M. N. and Cowell, F. (2024), 'Inequality as an externality: Consequences for tax design', *Journal of Public Economics* **235**, 105139.